

Architecting a Data Lake

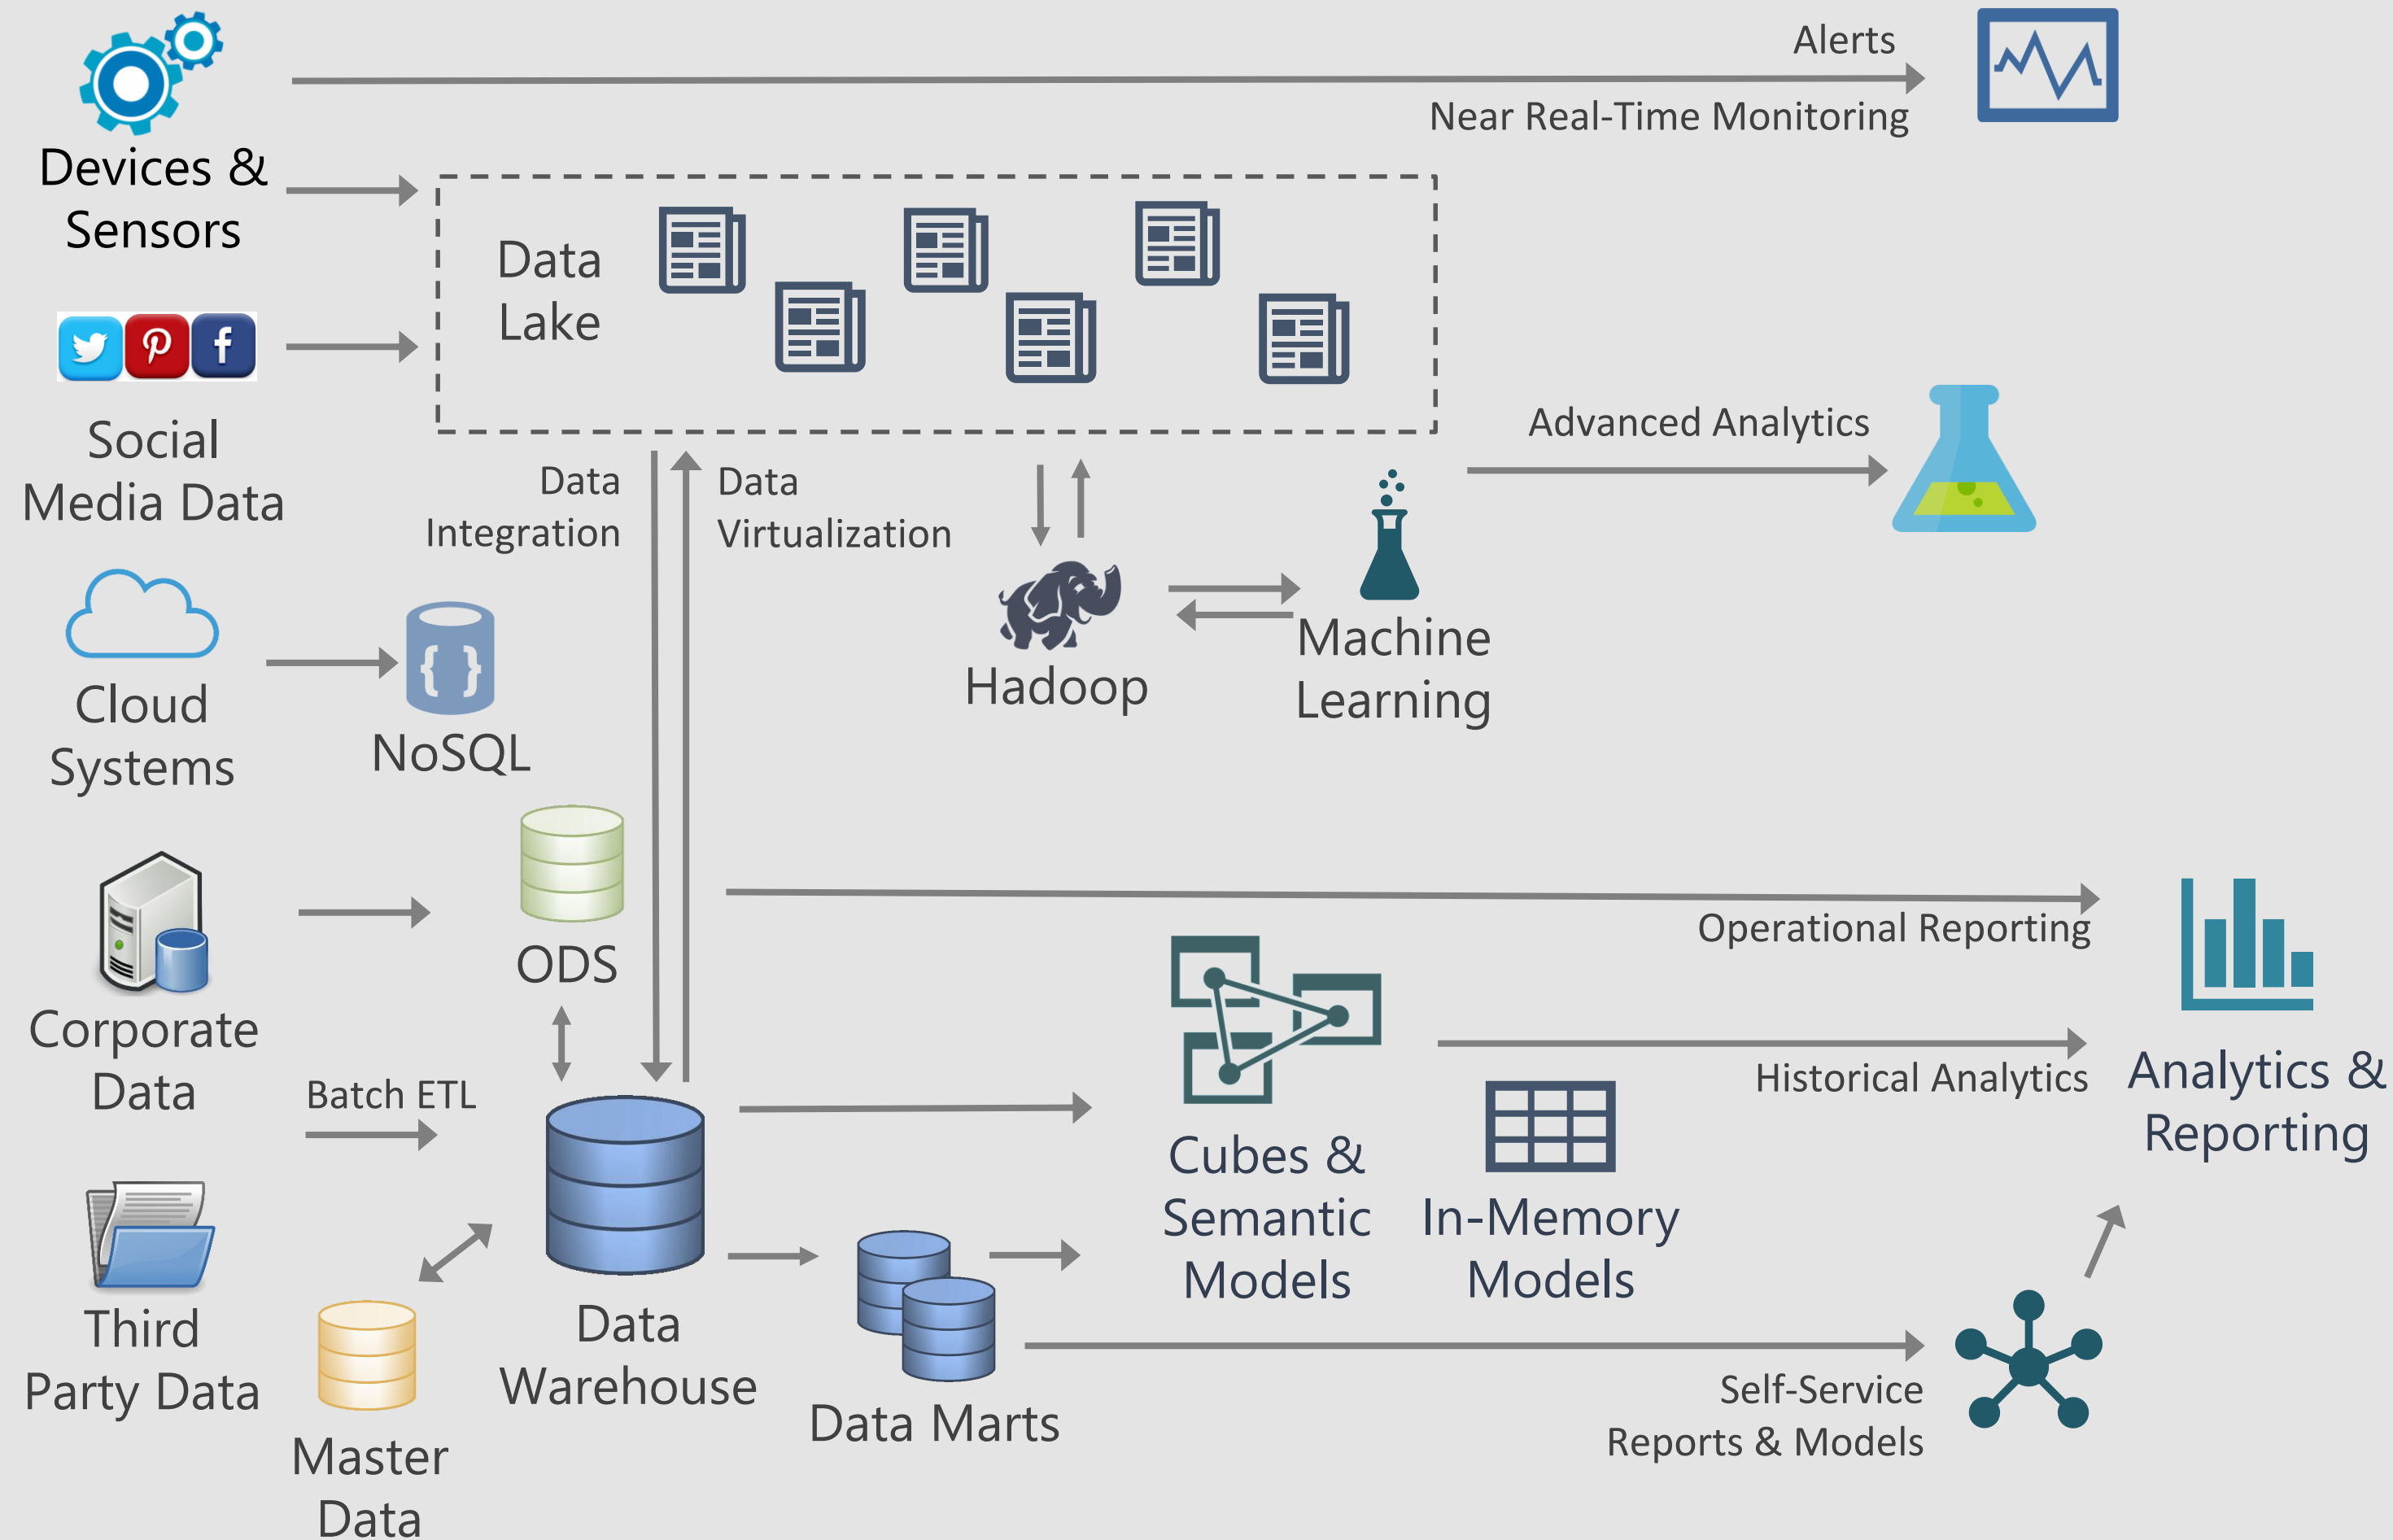
Chad Gronbach
Chief Technology Architect
Microsoft Technology Center - Boston

Content Credit: [JamesSerra.com](https://www.jameserra.com)



Modern Multi-Platform Architectures

Modern Data Warehousing & Analytics



Multi-platform architecture

- ✓ Handle a variety of data types & sources
- ✓ Larger data volumes at lower latency
- ✓ Bimodal: self-service + corporate BI to support all types of users
- ✓ Newer cloud services
- ✓ Advanced analytics scenarios
- ✓ Balance data integration & data virtualization

Definitions

(1/3)

Data Warehouse

Repository of data from multiple sources, cleansed & enriched for reporting; generally 'schema on write'

Data Lake

Repository of data for multi-structured data; generally 'schema on read'

Hadoop

(1) Data storage via HDFS (Hadoop Distributed File System), and
(2) Set of Apache projects for data processing and analytics

Lambda Architecture

Data processing & storage with batch, speed, and serving layers

ETL

Extract > Transform > Load: traditional paradigm associated with data warehousing and 'schema on write'

ELT

Extract > Load > Transform: newer paradigm associated with data lakes & 'schema on read'

Semantic Model

User-friendly interface for users on top of a data warehouse and/or data lake

Definitions

(2/3)

Data Integration

Physically moving data to integrate multiple sources together

Data Virtualization

Access to one or more distributed data sources without requiring the data to be physically materialized in another data structure

Federated Query

A type of data virtualization: access & consolidate data from multiple distributed data sources

Polyglot Persistence

A multi-platform strategy which values using the most effective technology based on the data itself ("best fit engineering")

Schema on Write

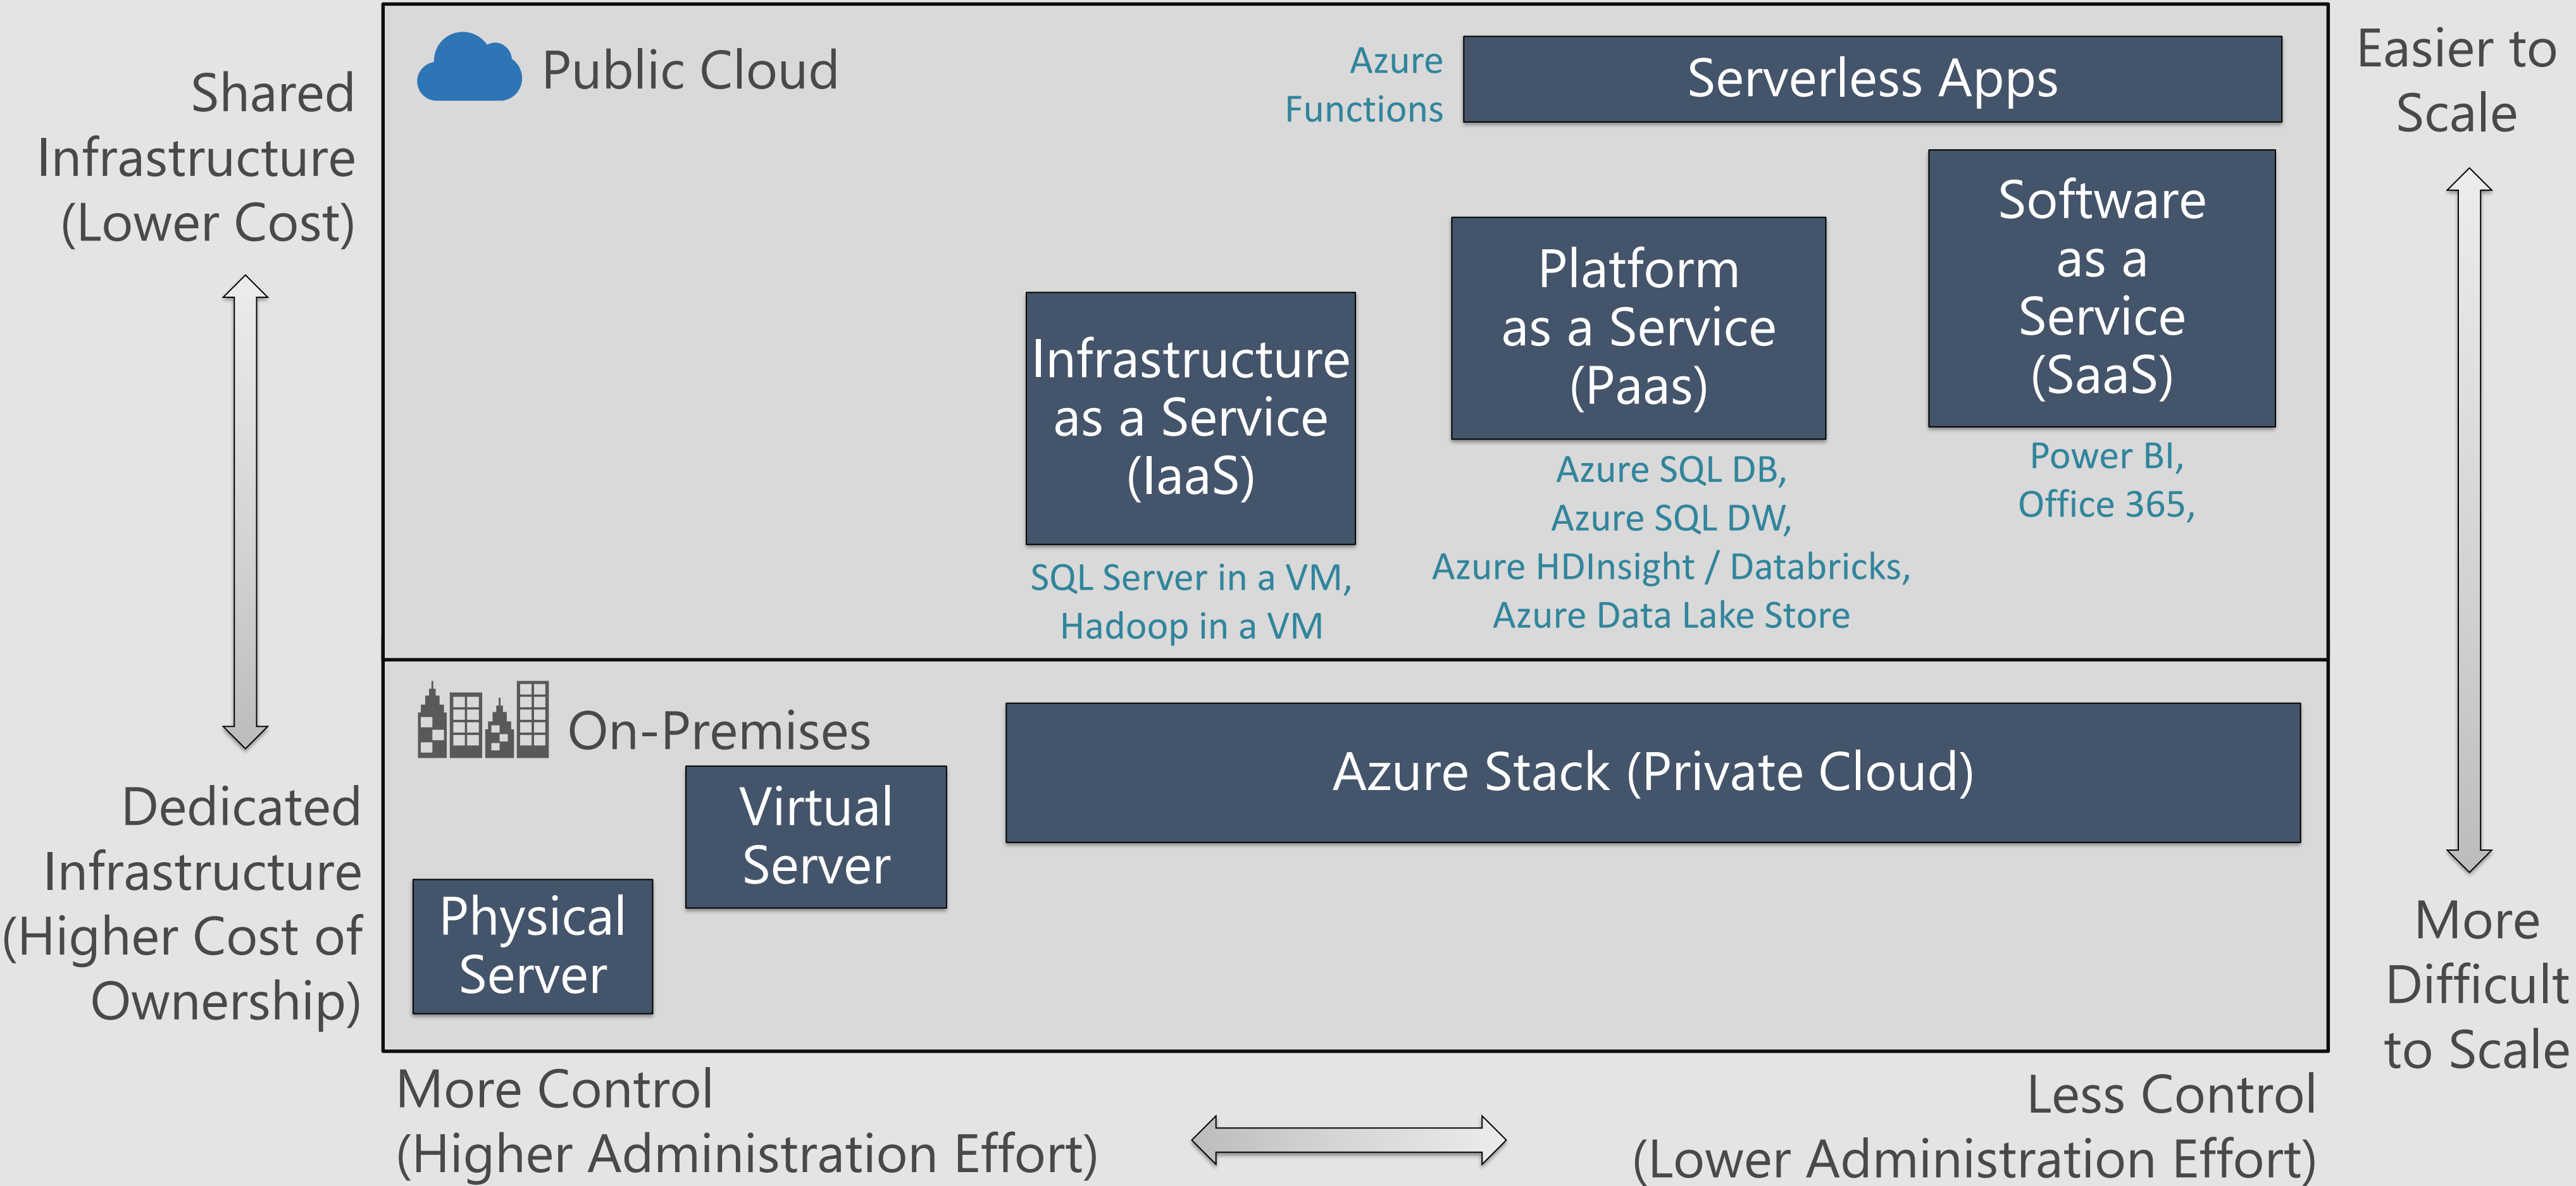
Data structure is applied at design time, requiring additional up-front effort to formulate a data model (relational DBs)


Schema on Read

Data structure is applied at query time rather than when the data is initially stored (data lakes, NoSQL)

Definitions

(3/3)





What are some common challenges
of analytical environments?

Challenges of Analytical Environments

Agility

- ✓ Reducing time to value
- ✓ Minimizing chaos with self-service
- ✓ Evolving & maturing technology
- ✓ Balancing schema-on-read with schema-on-write
- ✓ How strict to be with dimensional design?

Complexity

- ✓ Hybrid scenarios
- ✓ Multi-platform architecture
- ✓ Ever-increasing data volumes
- ✓ Diversity of file types & formats
- ✓ Effort & cost of data integration
- ✓ Many skillsets needed

Balance

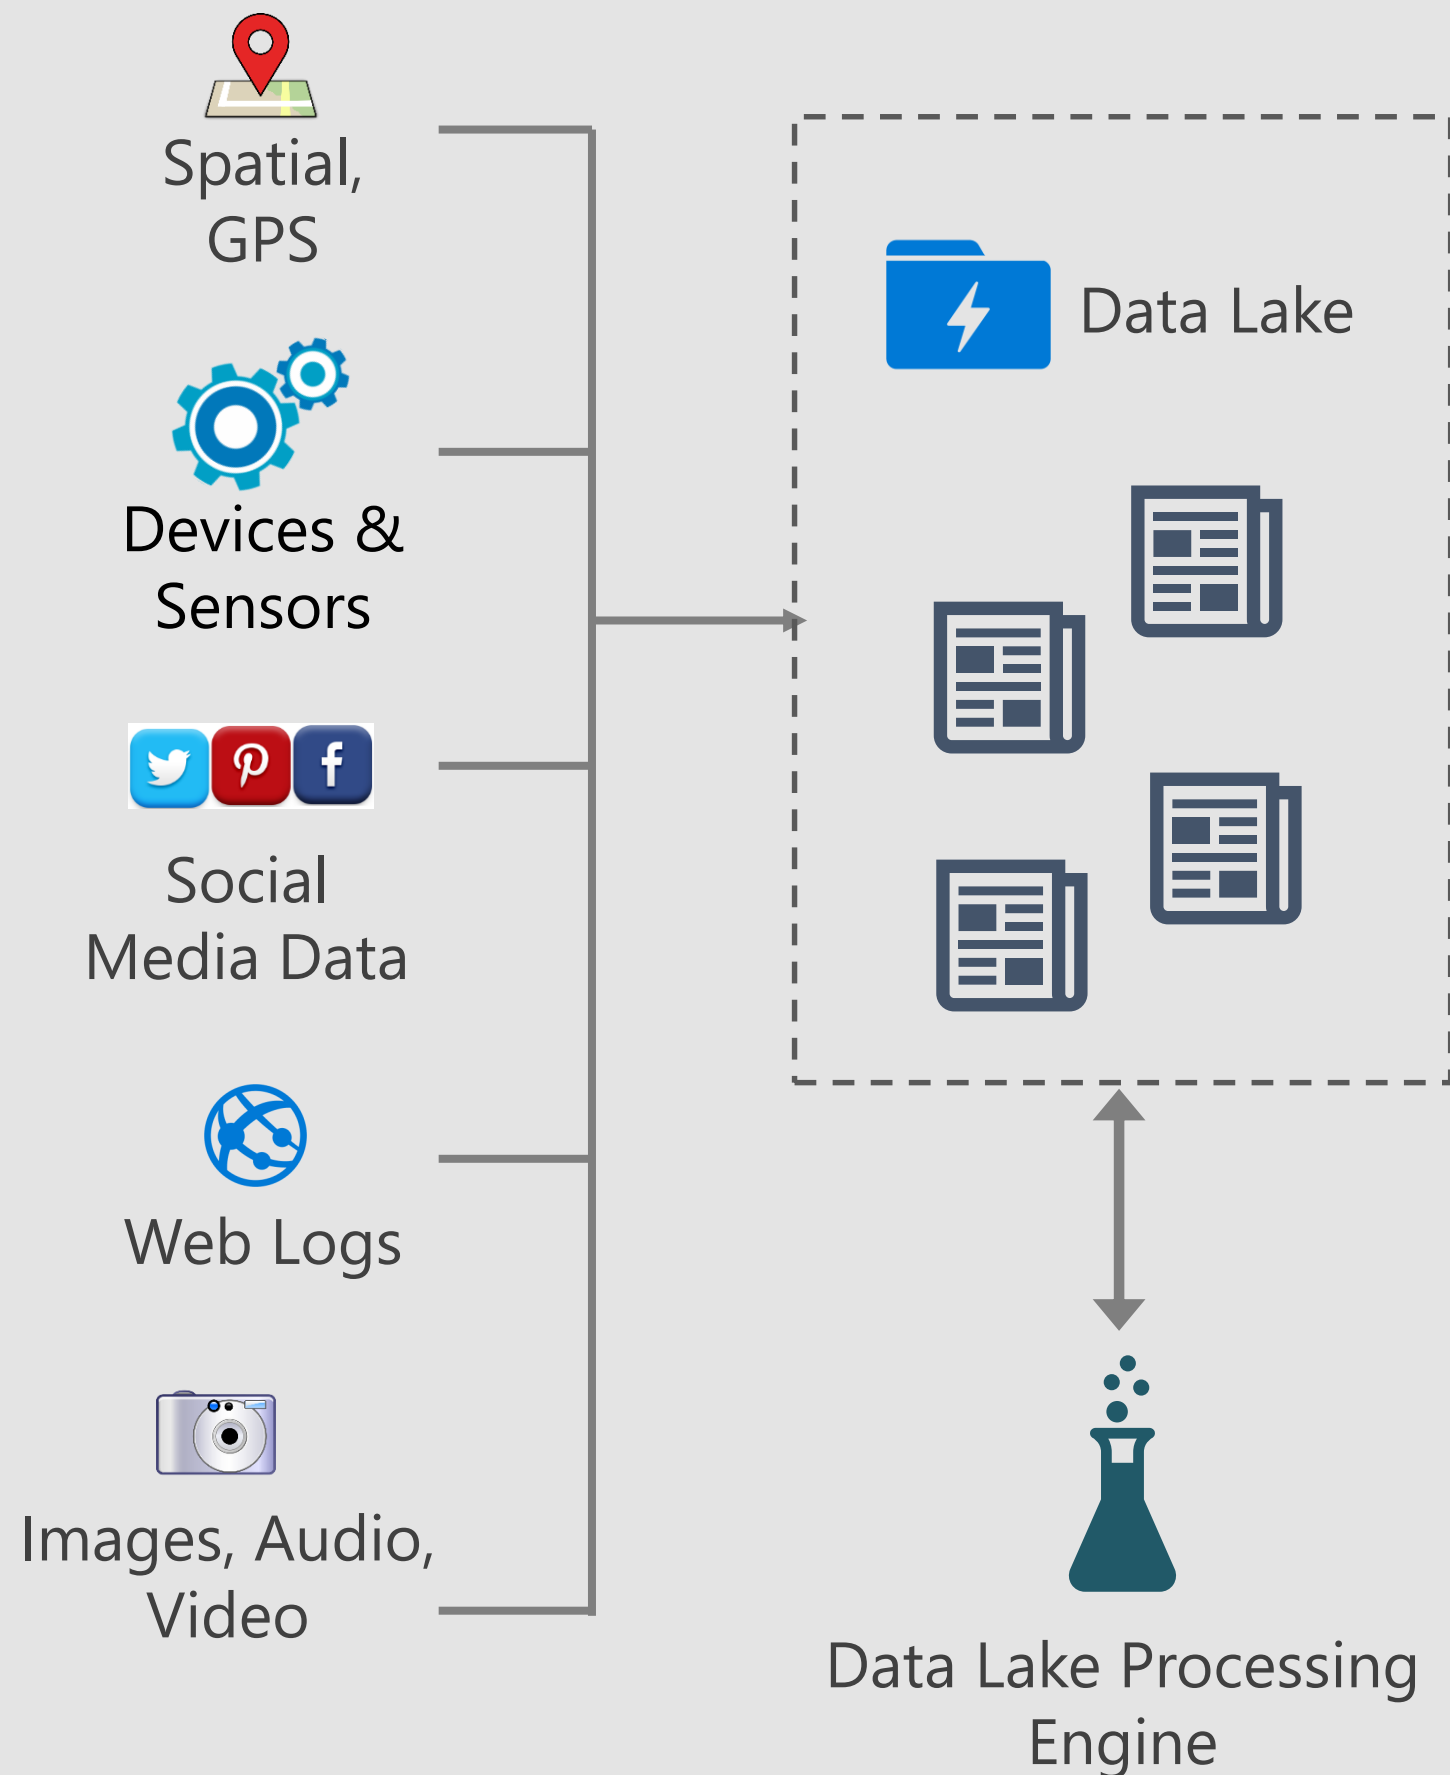
- ✓ Self-service solutions challenge corporate DW solutions
- ✓ Operationalizing valuable user-created solutions (including data science)
- ✓ Handling ownership changes of a productionized solution

Never-Ending

- ✓ Data quality
- ✓ User trust
- ✓ Master data
- ✓ Security
- ✓ Governance
- ✓ Performance

Data Lake Overview & Use Cases

Data Lake



A **repository** for storing large quantities of disparate sources of data in its native format

One **architectural platform** to house all types of data:

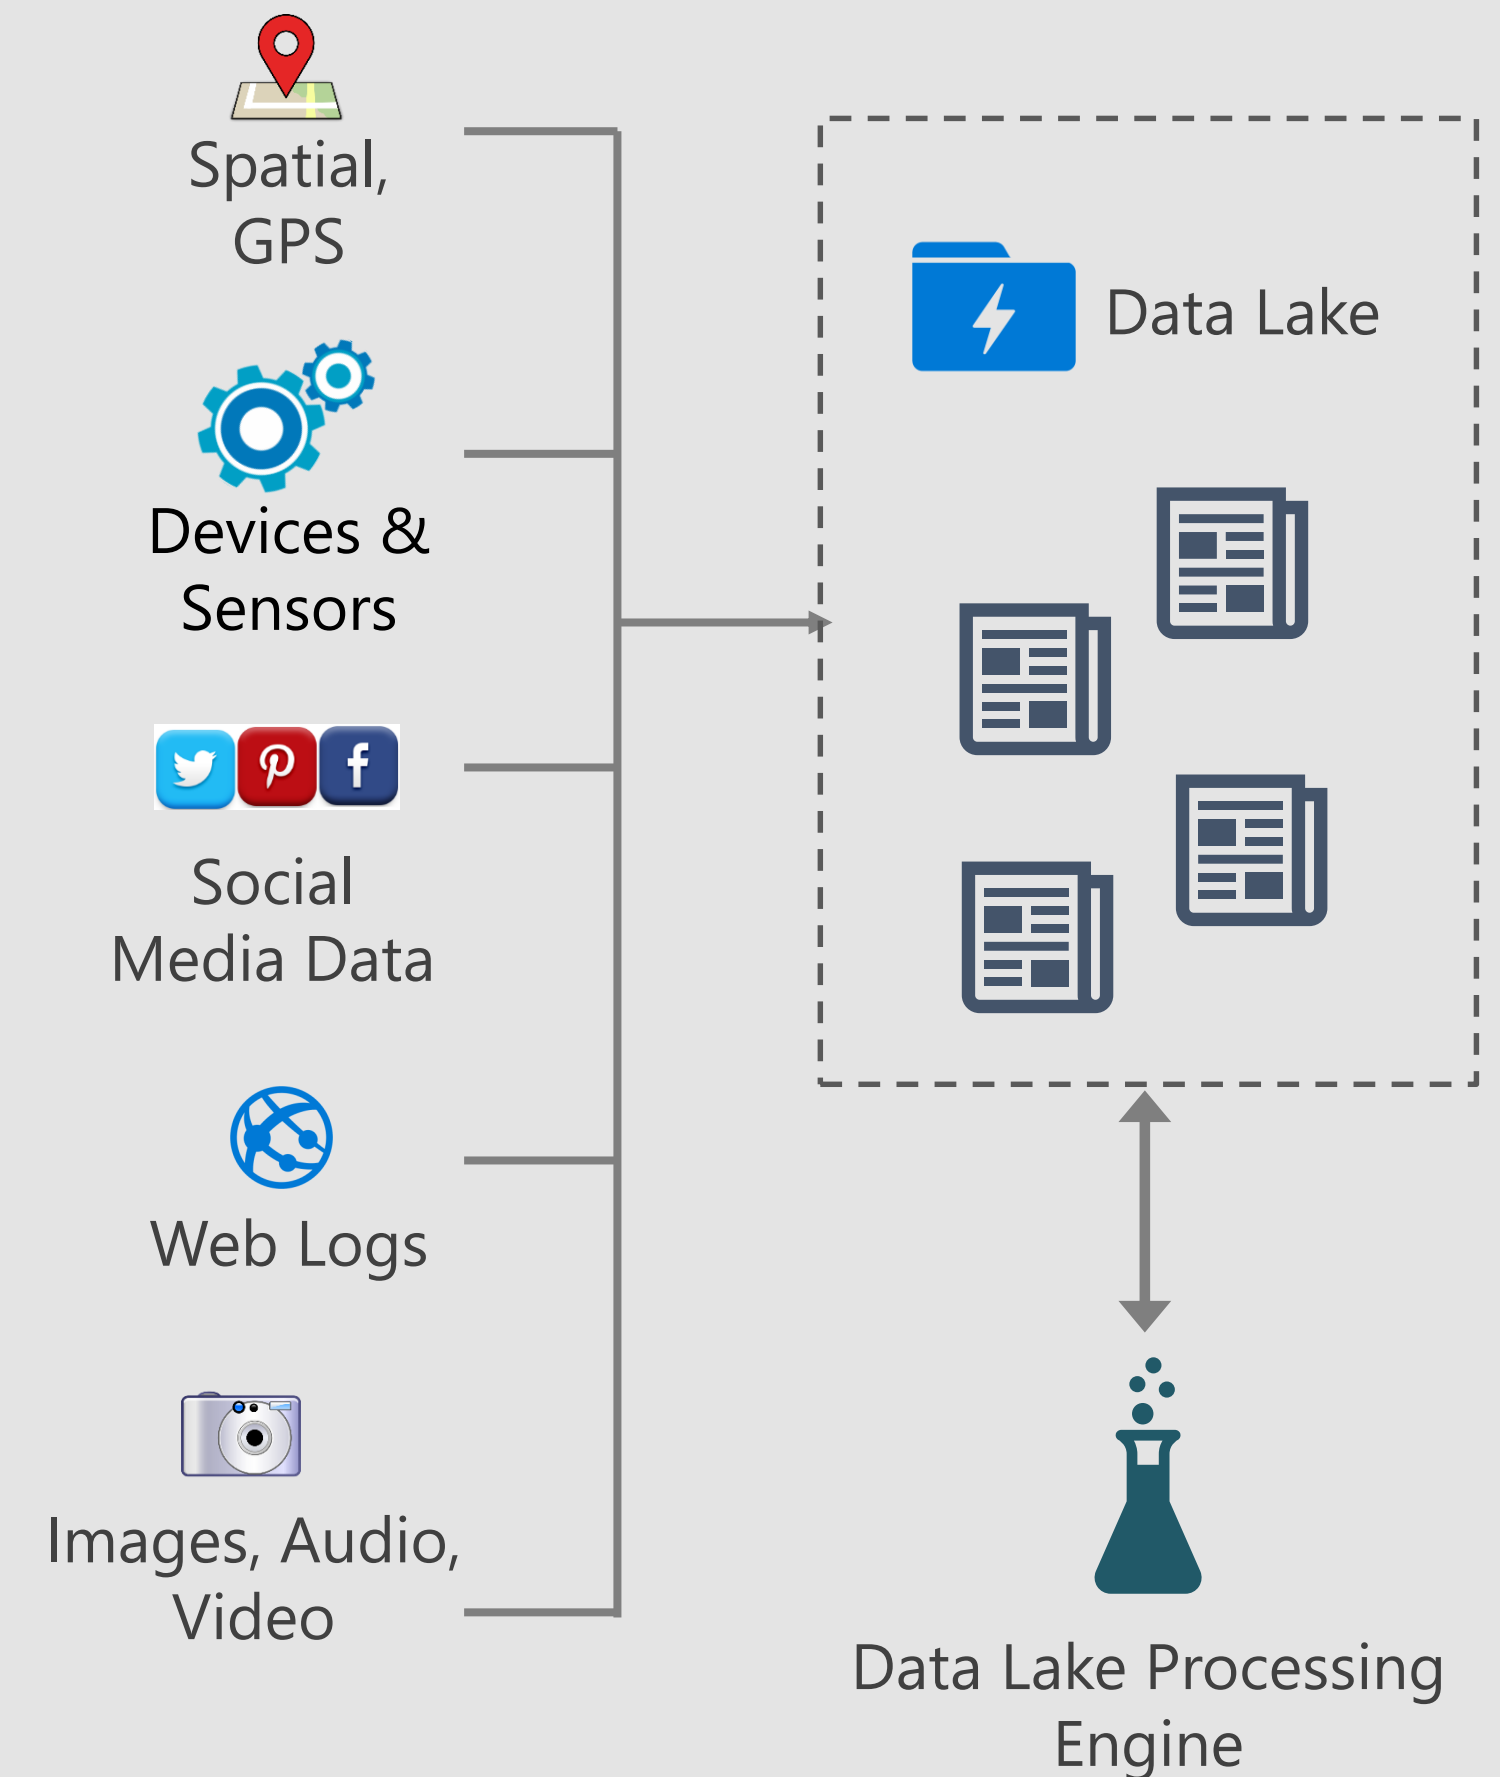
- ✓ Machine-generated data (ex: IoT, logs)
- ✓ Human-generated data (ex: tweets, e-mail)
- ✓ Traditional operational data (ex: sales, inventory)

B A **processing engine** for analyzing data

Data Lake Objectives

(1/2)


- ✓ Reduce up-front effort by ingesting data in any format, any size, without requiring a schema initially
- ✓ Make acquiring new data easy, so it can be available for data science & analysis quickly
- ✓ Store large volume of multi-structured data in its native format
- ✓ Storage for additional types of data which were historically difficult to obtain or store
- ✓ Reduce the long-term ownership cost of data management & storage



Data Lake Objectives

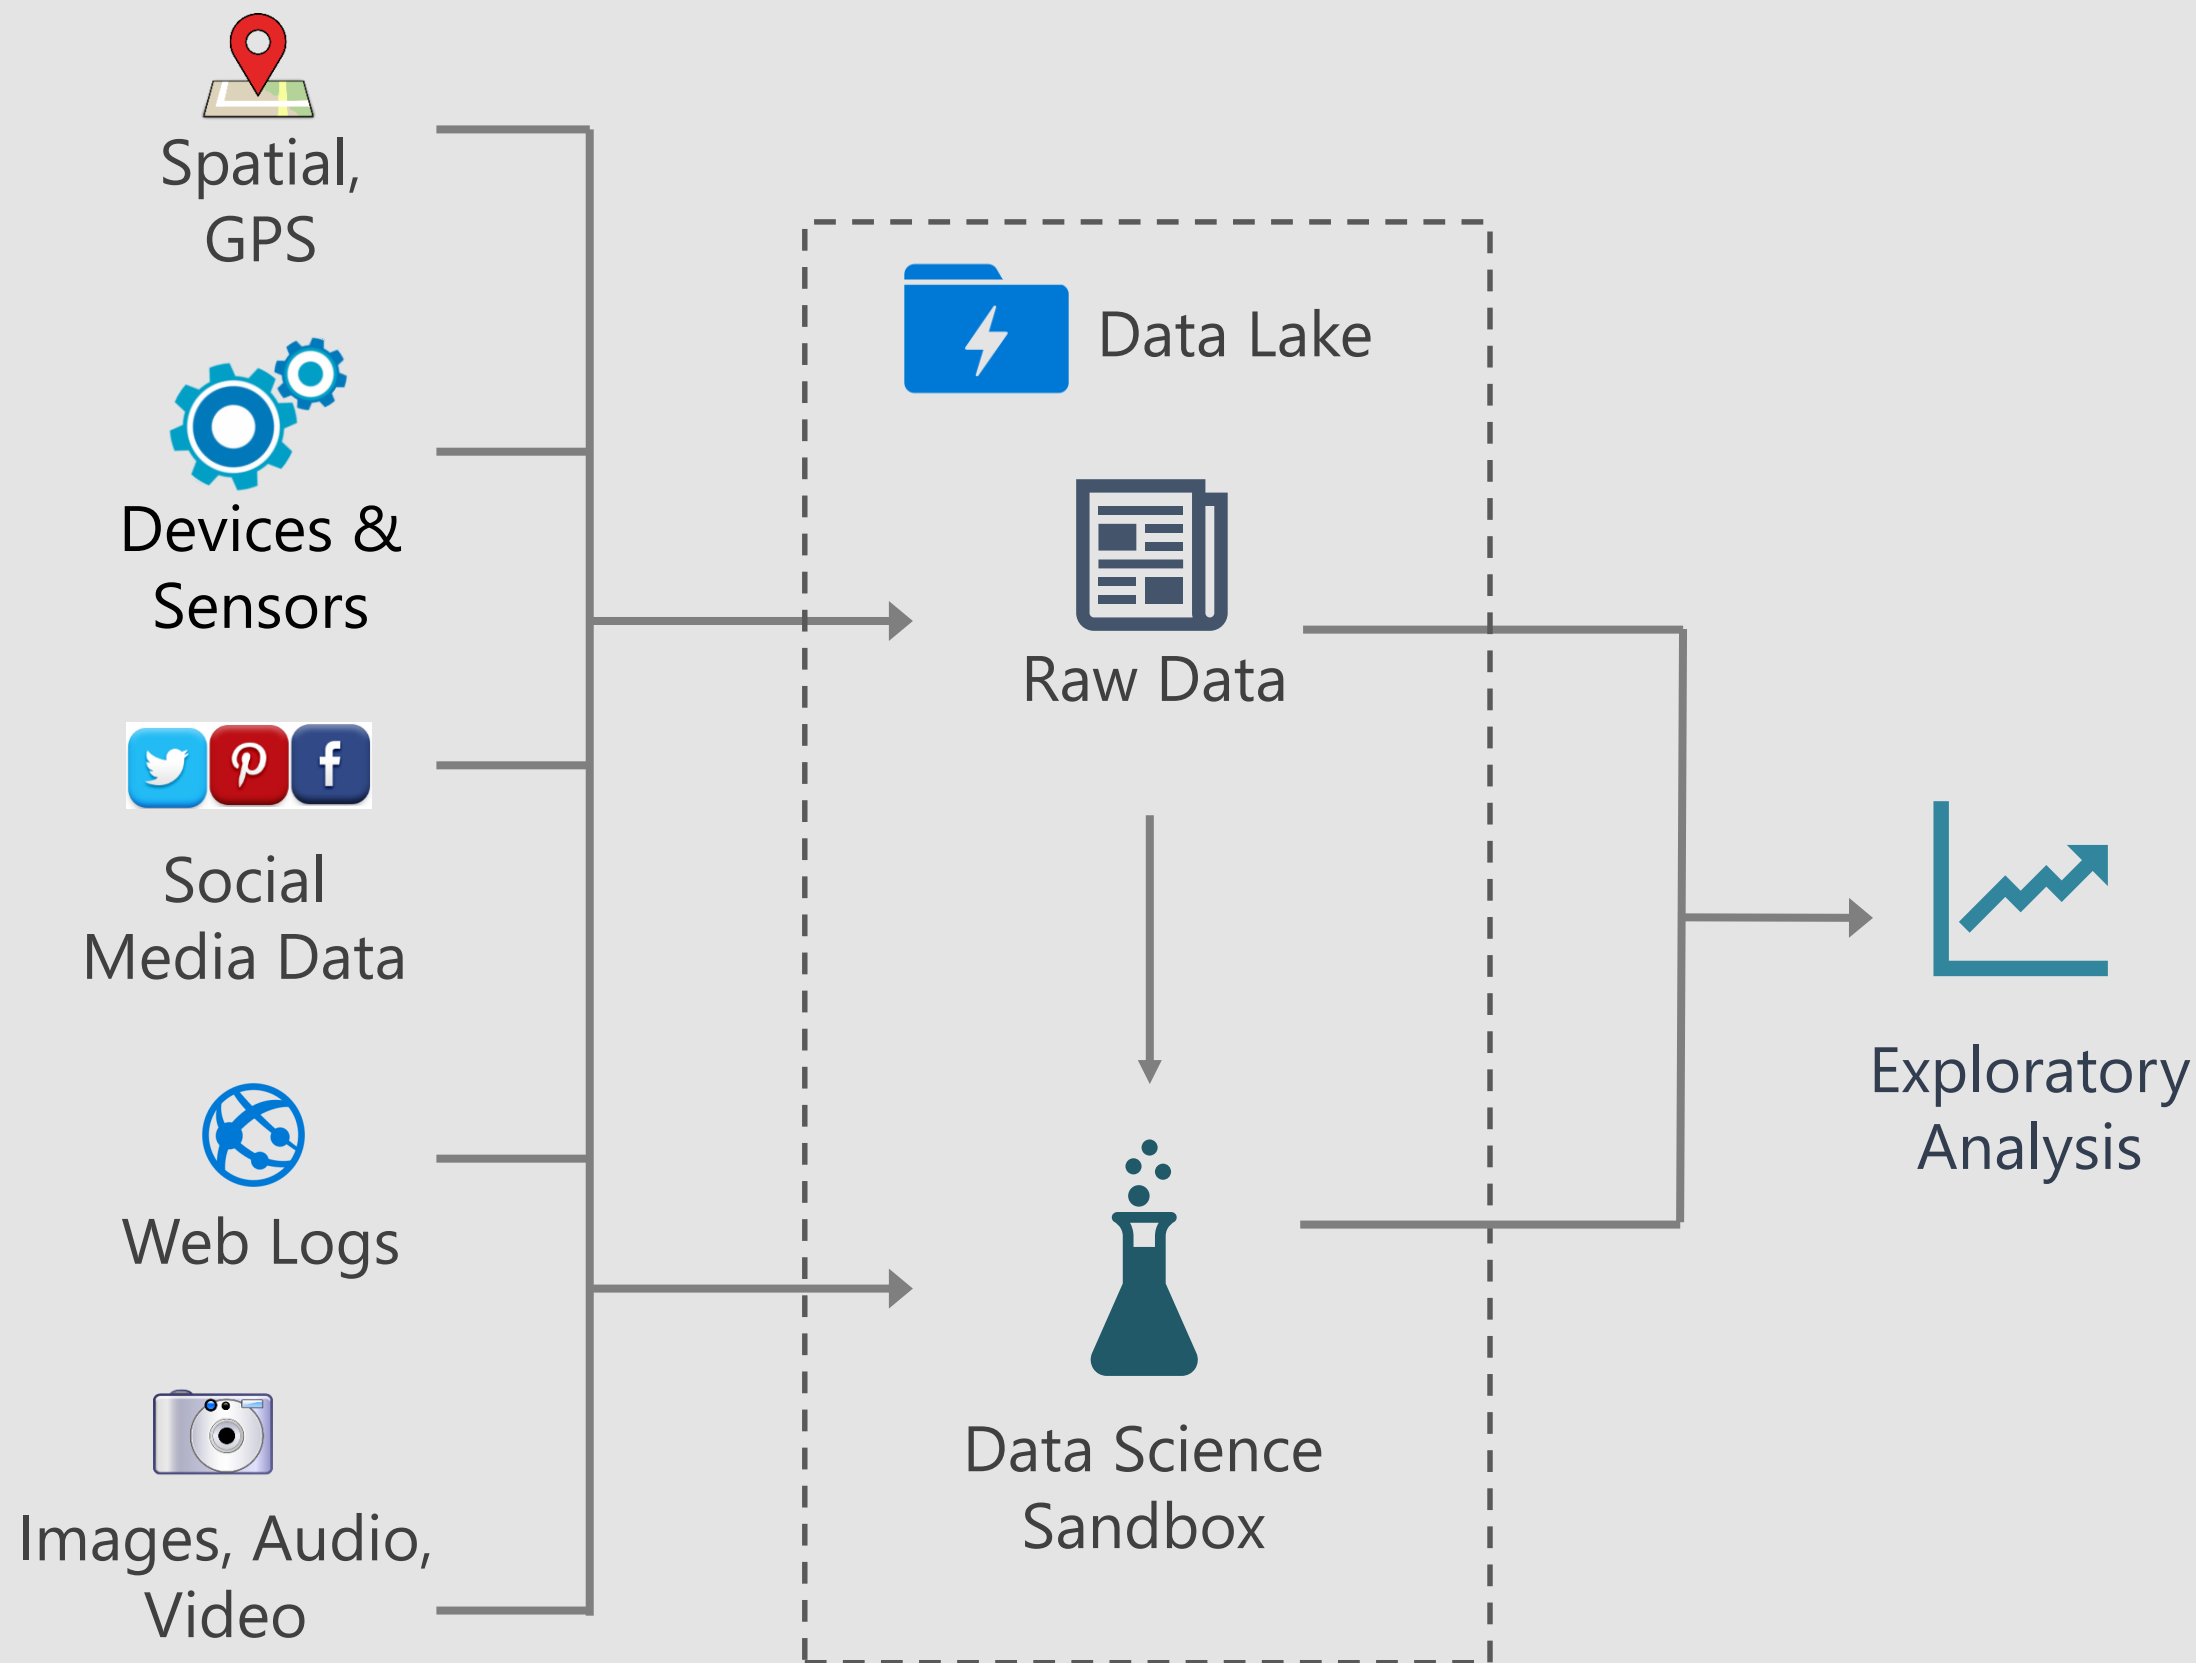
(2/2)

- ✓ Schema-on-read: **Defer work to 'schematize'** after value & requirements are known
- ✓ Achieve **agility** faster than a traditional data warehouse can to speed up **decision-making** ability
- ✓ Access to **low-latency data**
- ✓ Different / new **value** proposition vs. traditional data warehousing
- ✓ Facilitate **advanced analytics** scenarios



Data Lake Use Cases

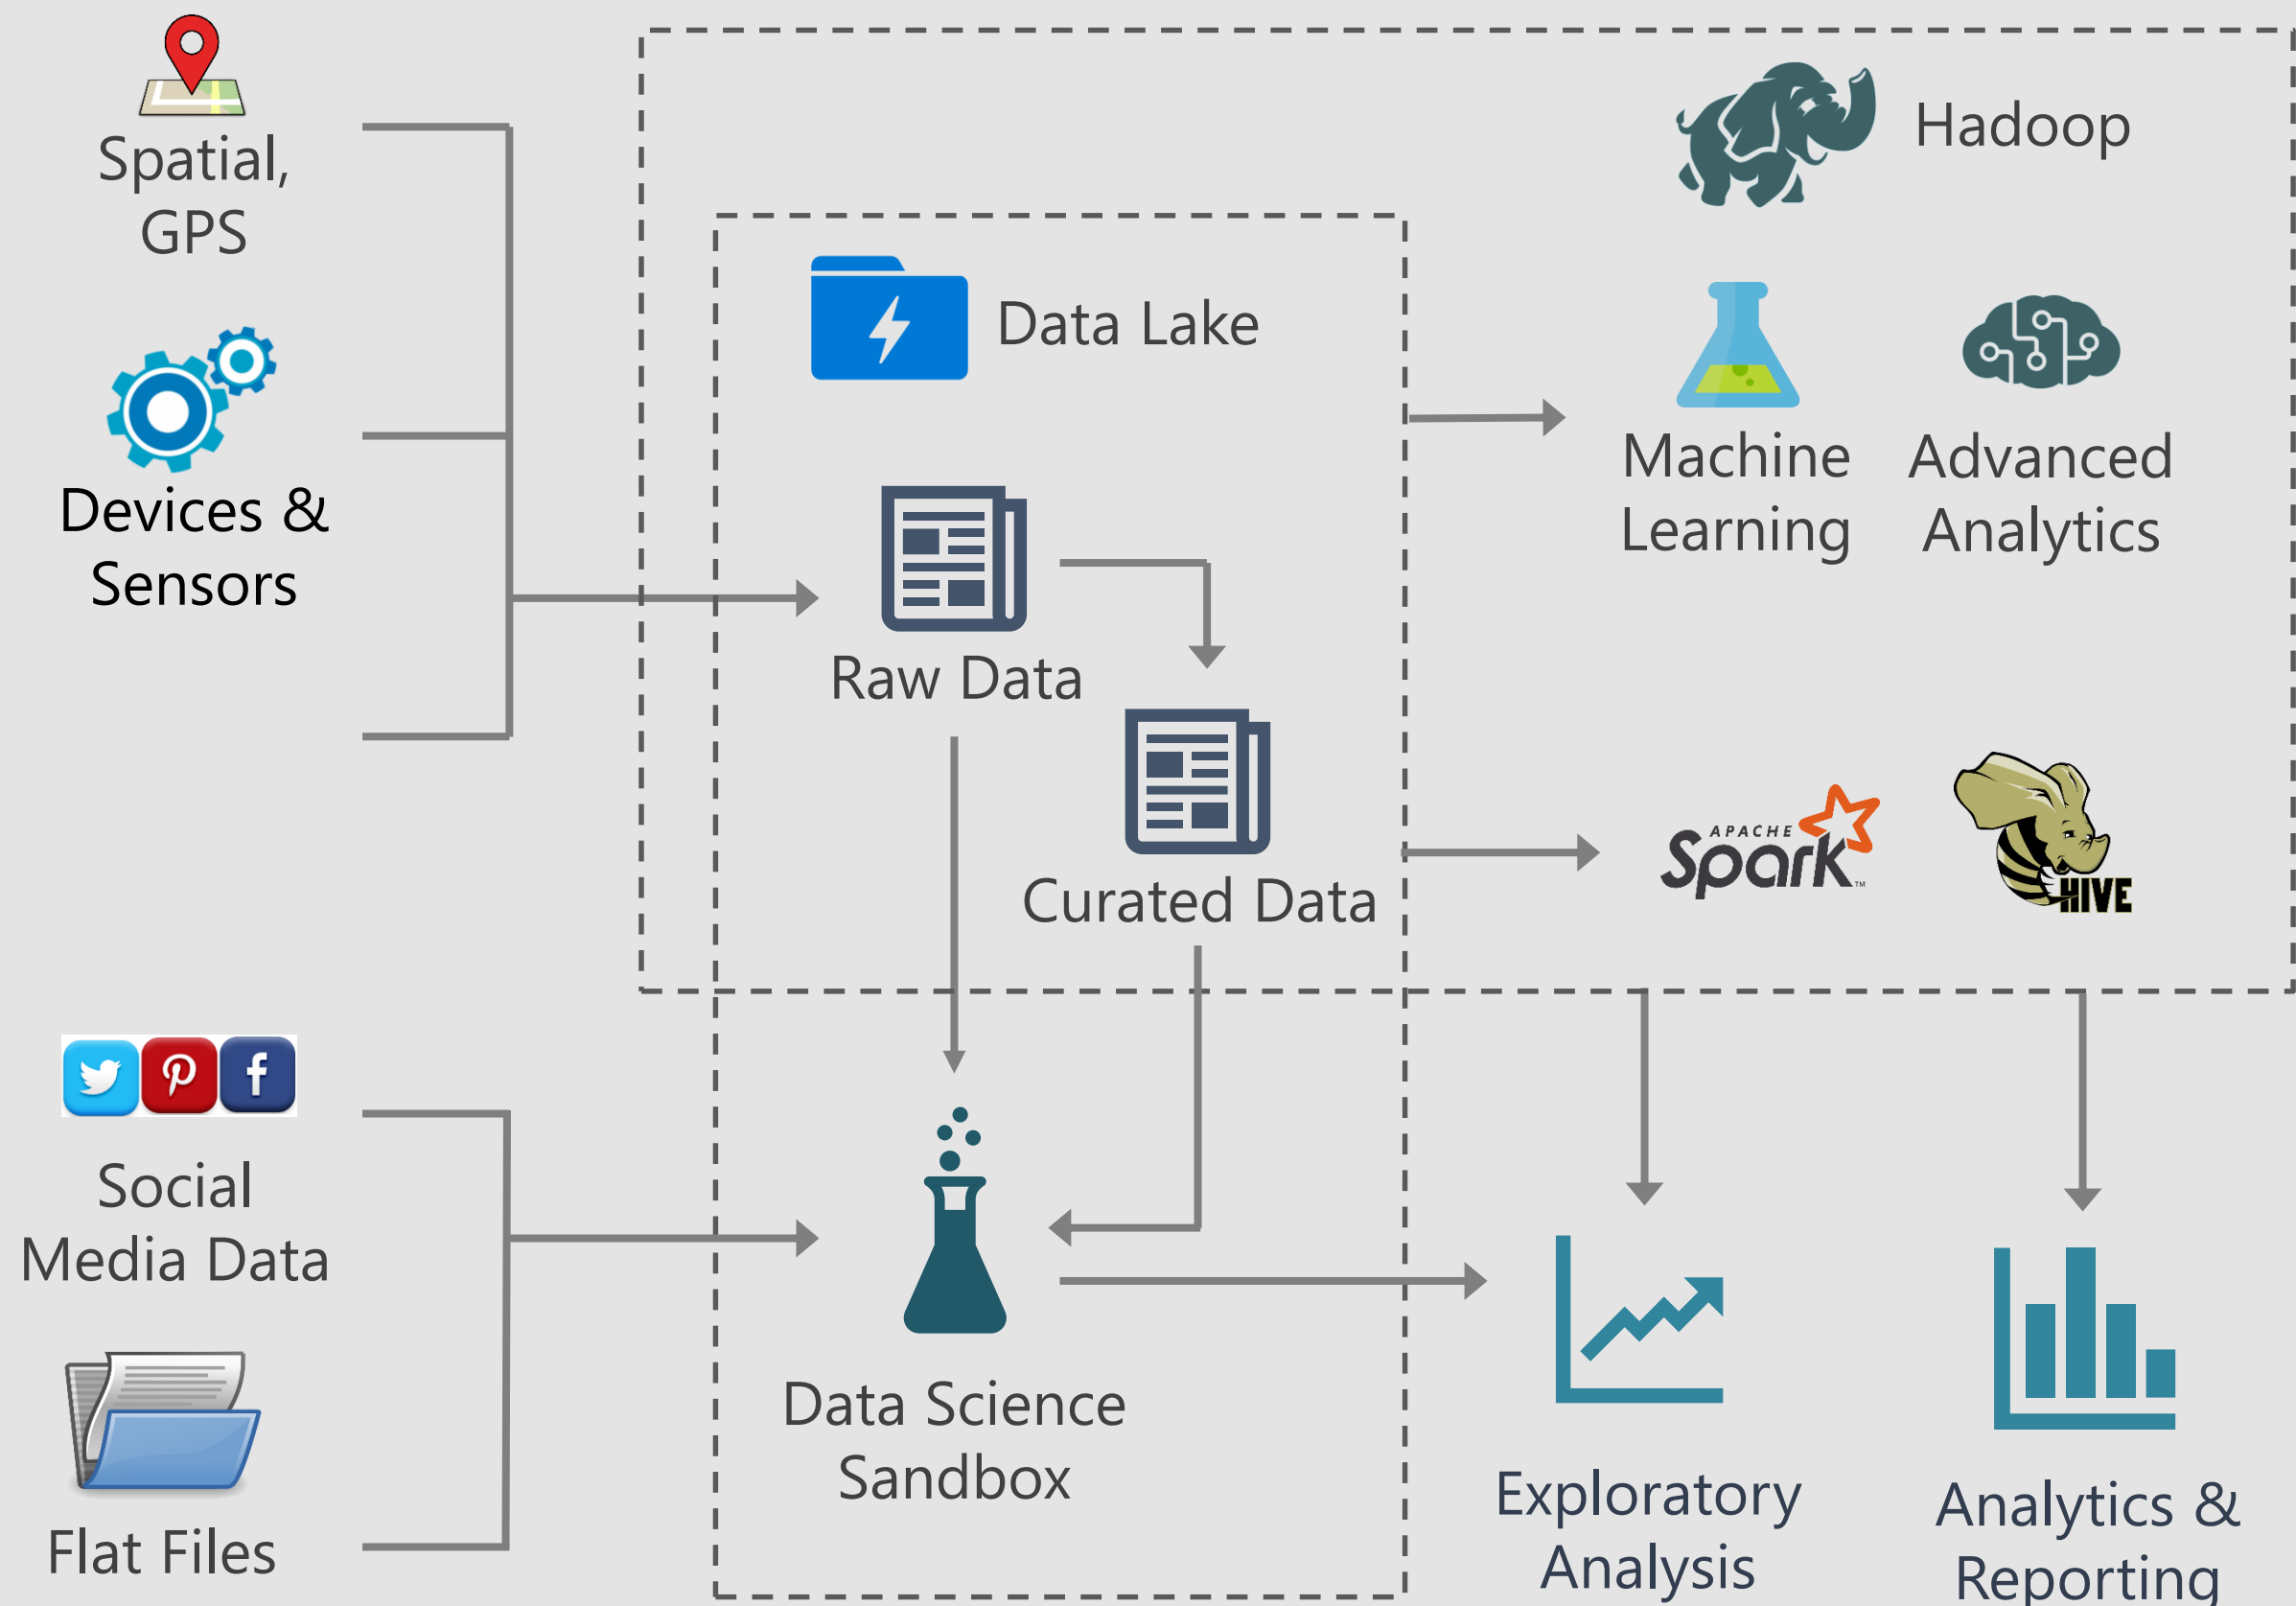
Ingestion of New File Types



- ✓ Preparatory file storage for multi-structured data
- ✓ Exploratory analysis + POCs to determine value of new data types & sources
- ✓ Affords additional time for longer-term planning while accumulating data or handling an influx of data

Data Lake Use Cases

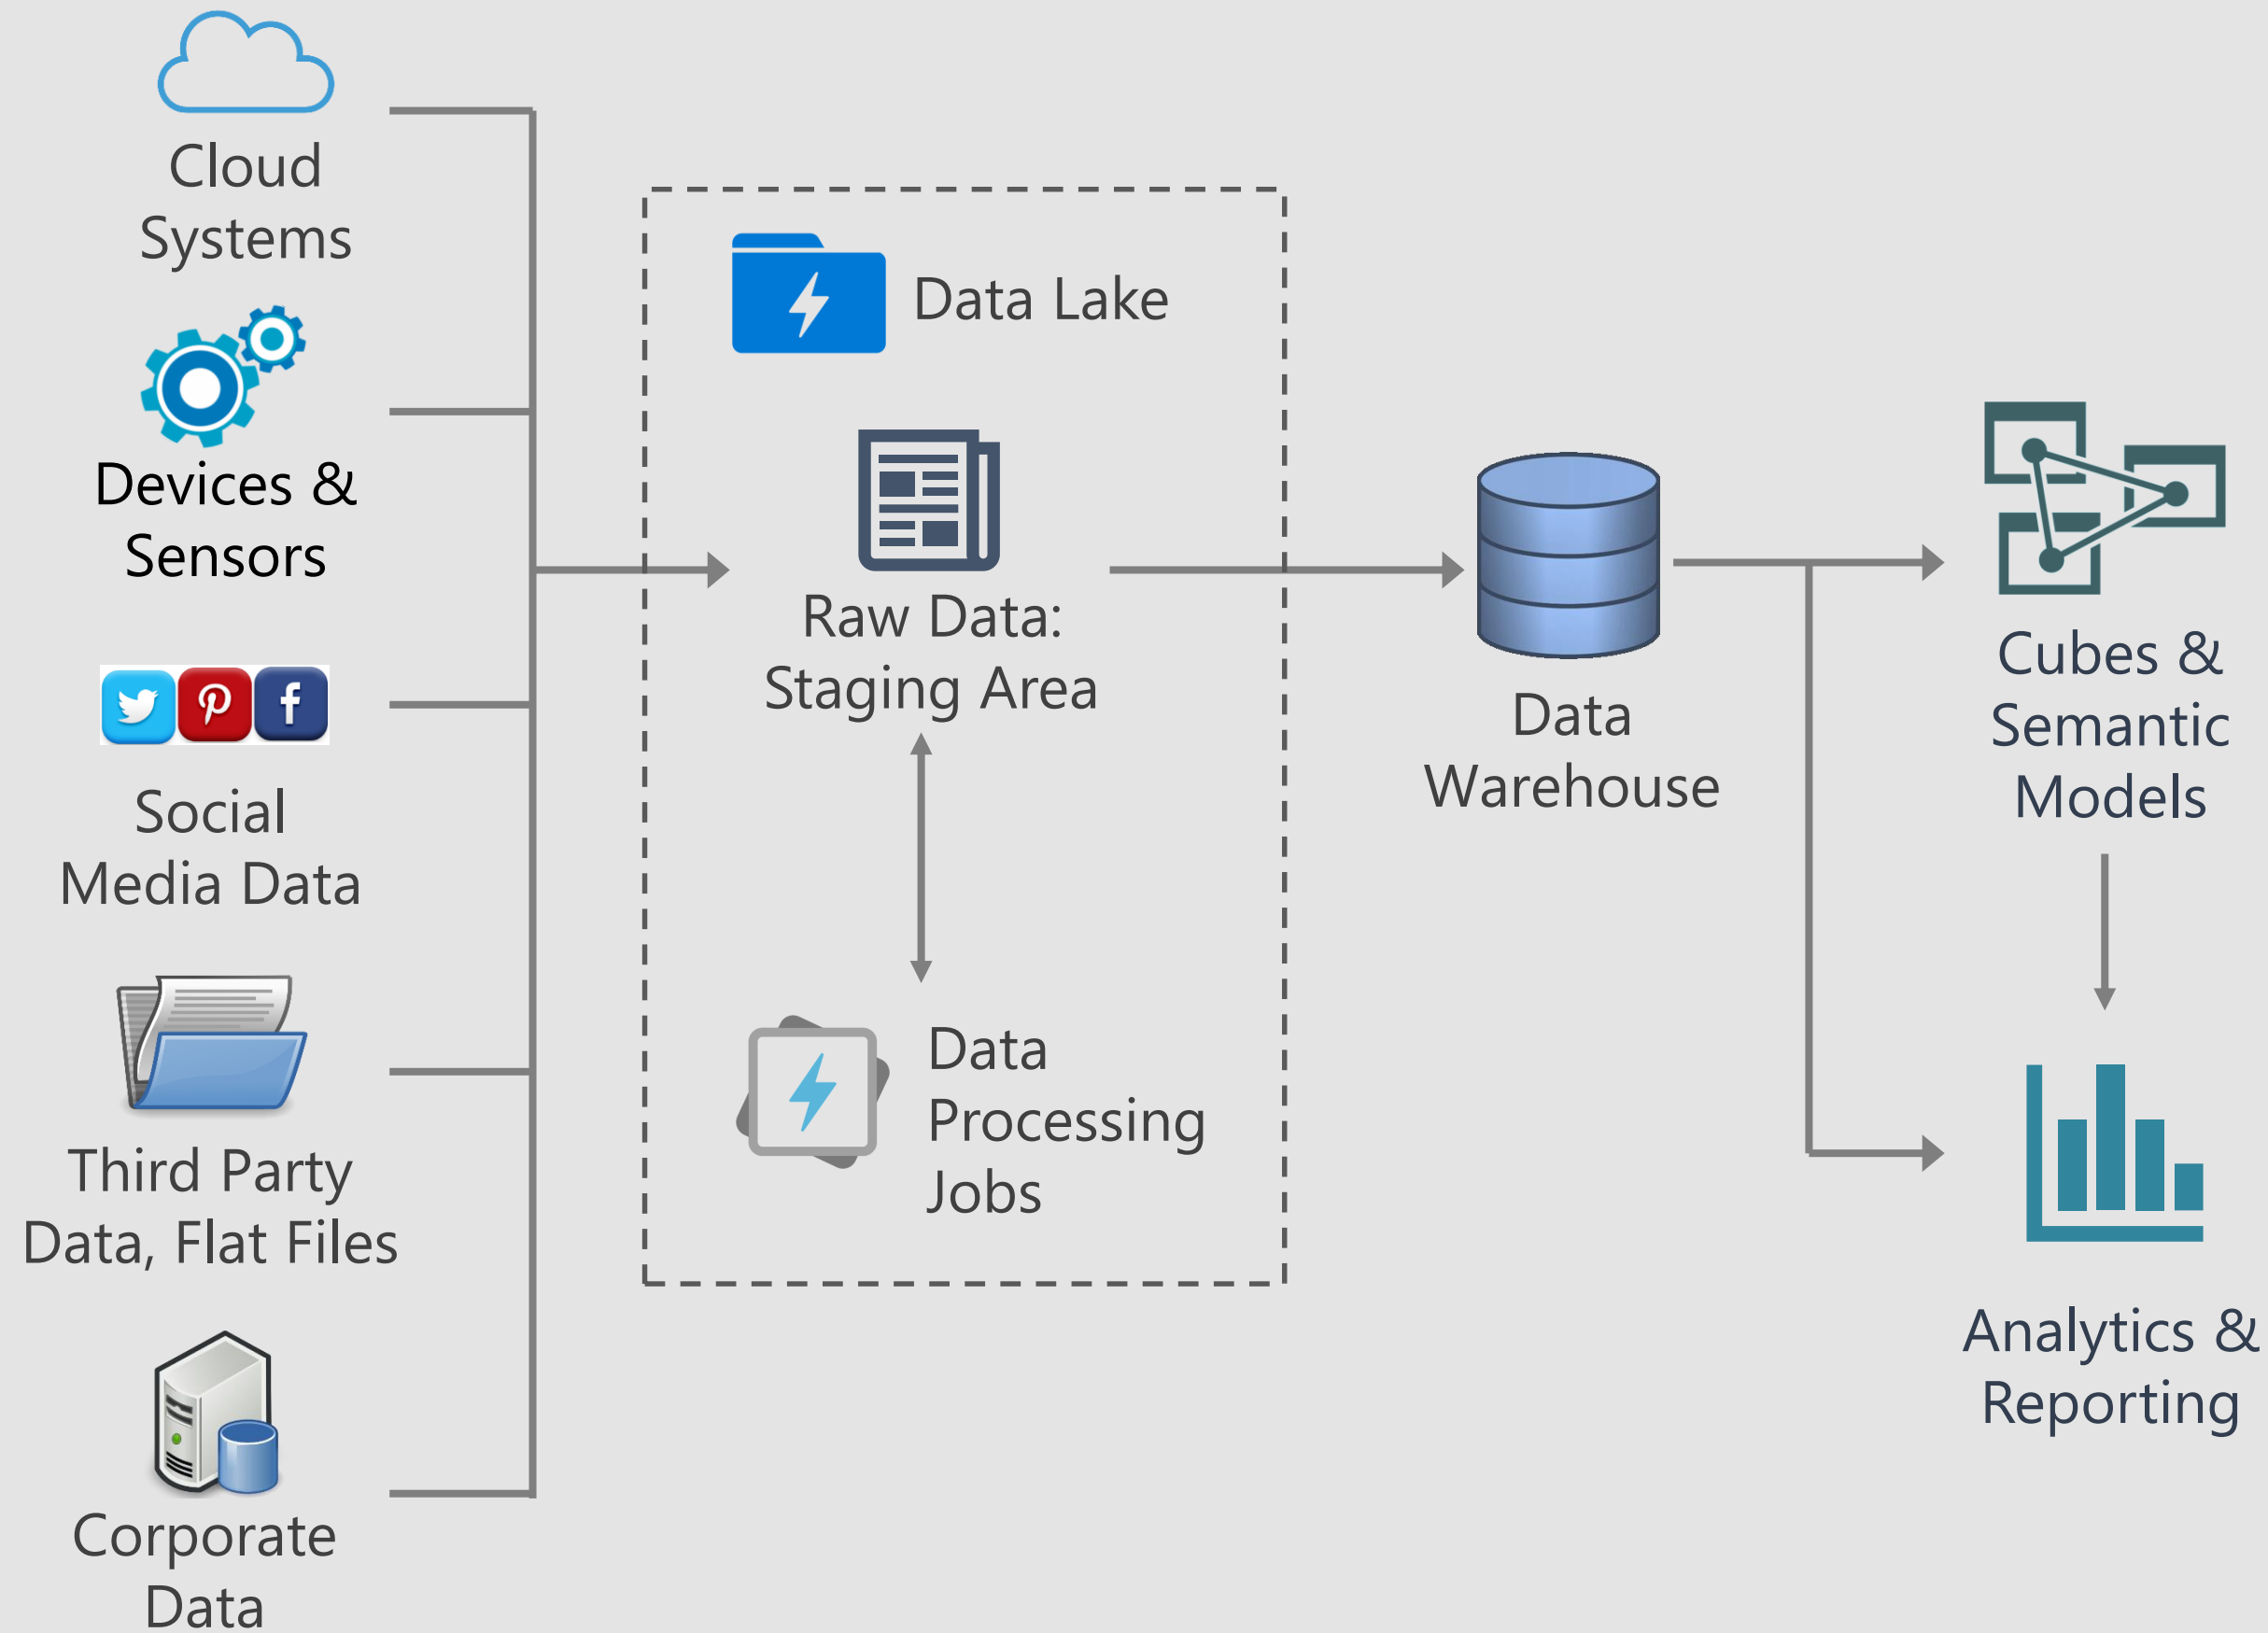
Data Science Experimentation | Hadoop Integration



- ✓ Sandbox solutions for initial data prep, experimentation, and analysis
- ✓ Migrate from proof of concept to operationalized solution
- ✓ Integrate with open source projects such as Hive, Pig, Spark, Storm, etc.
- ✓ Big data clusters
- ✓ SQL-on-Hadoop solutions

Data Lake Use Cases

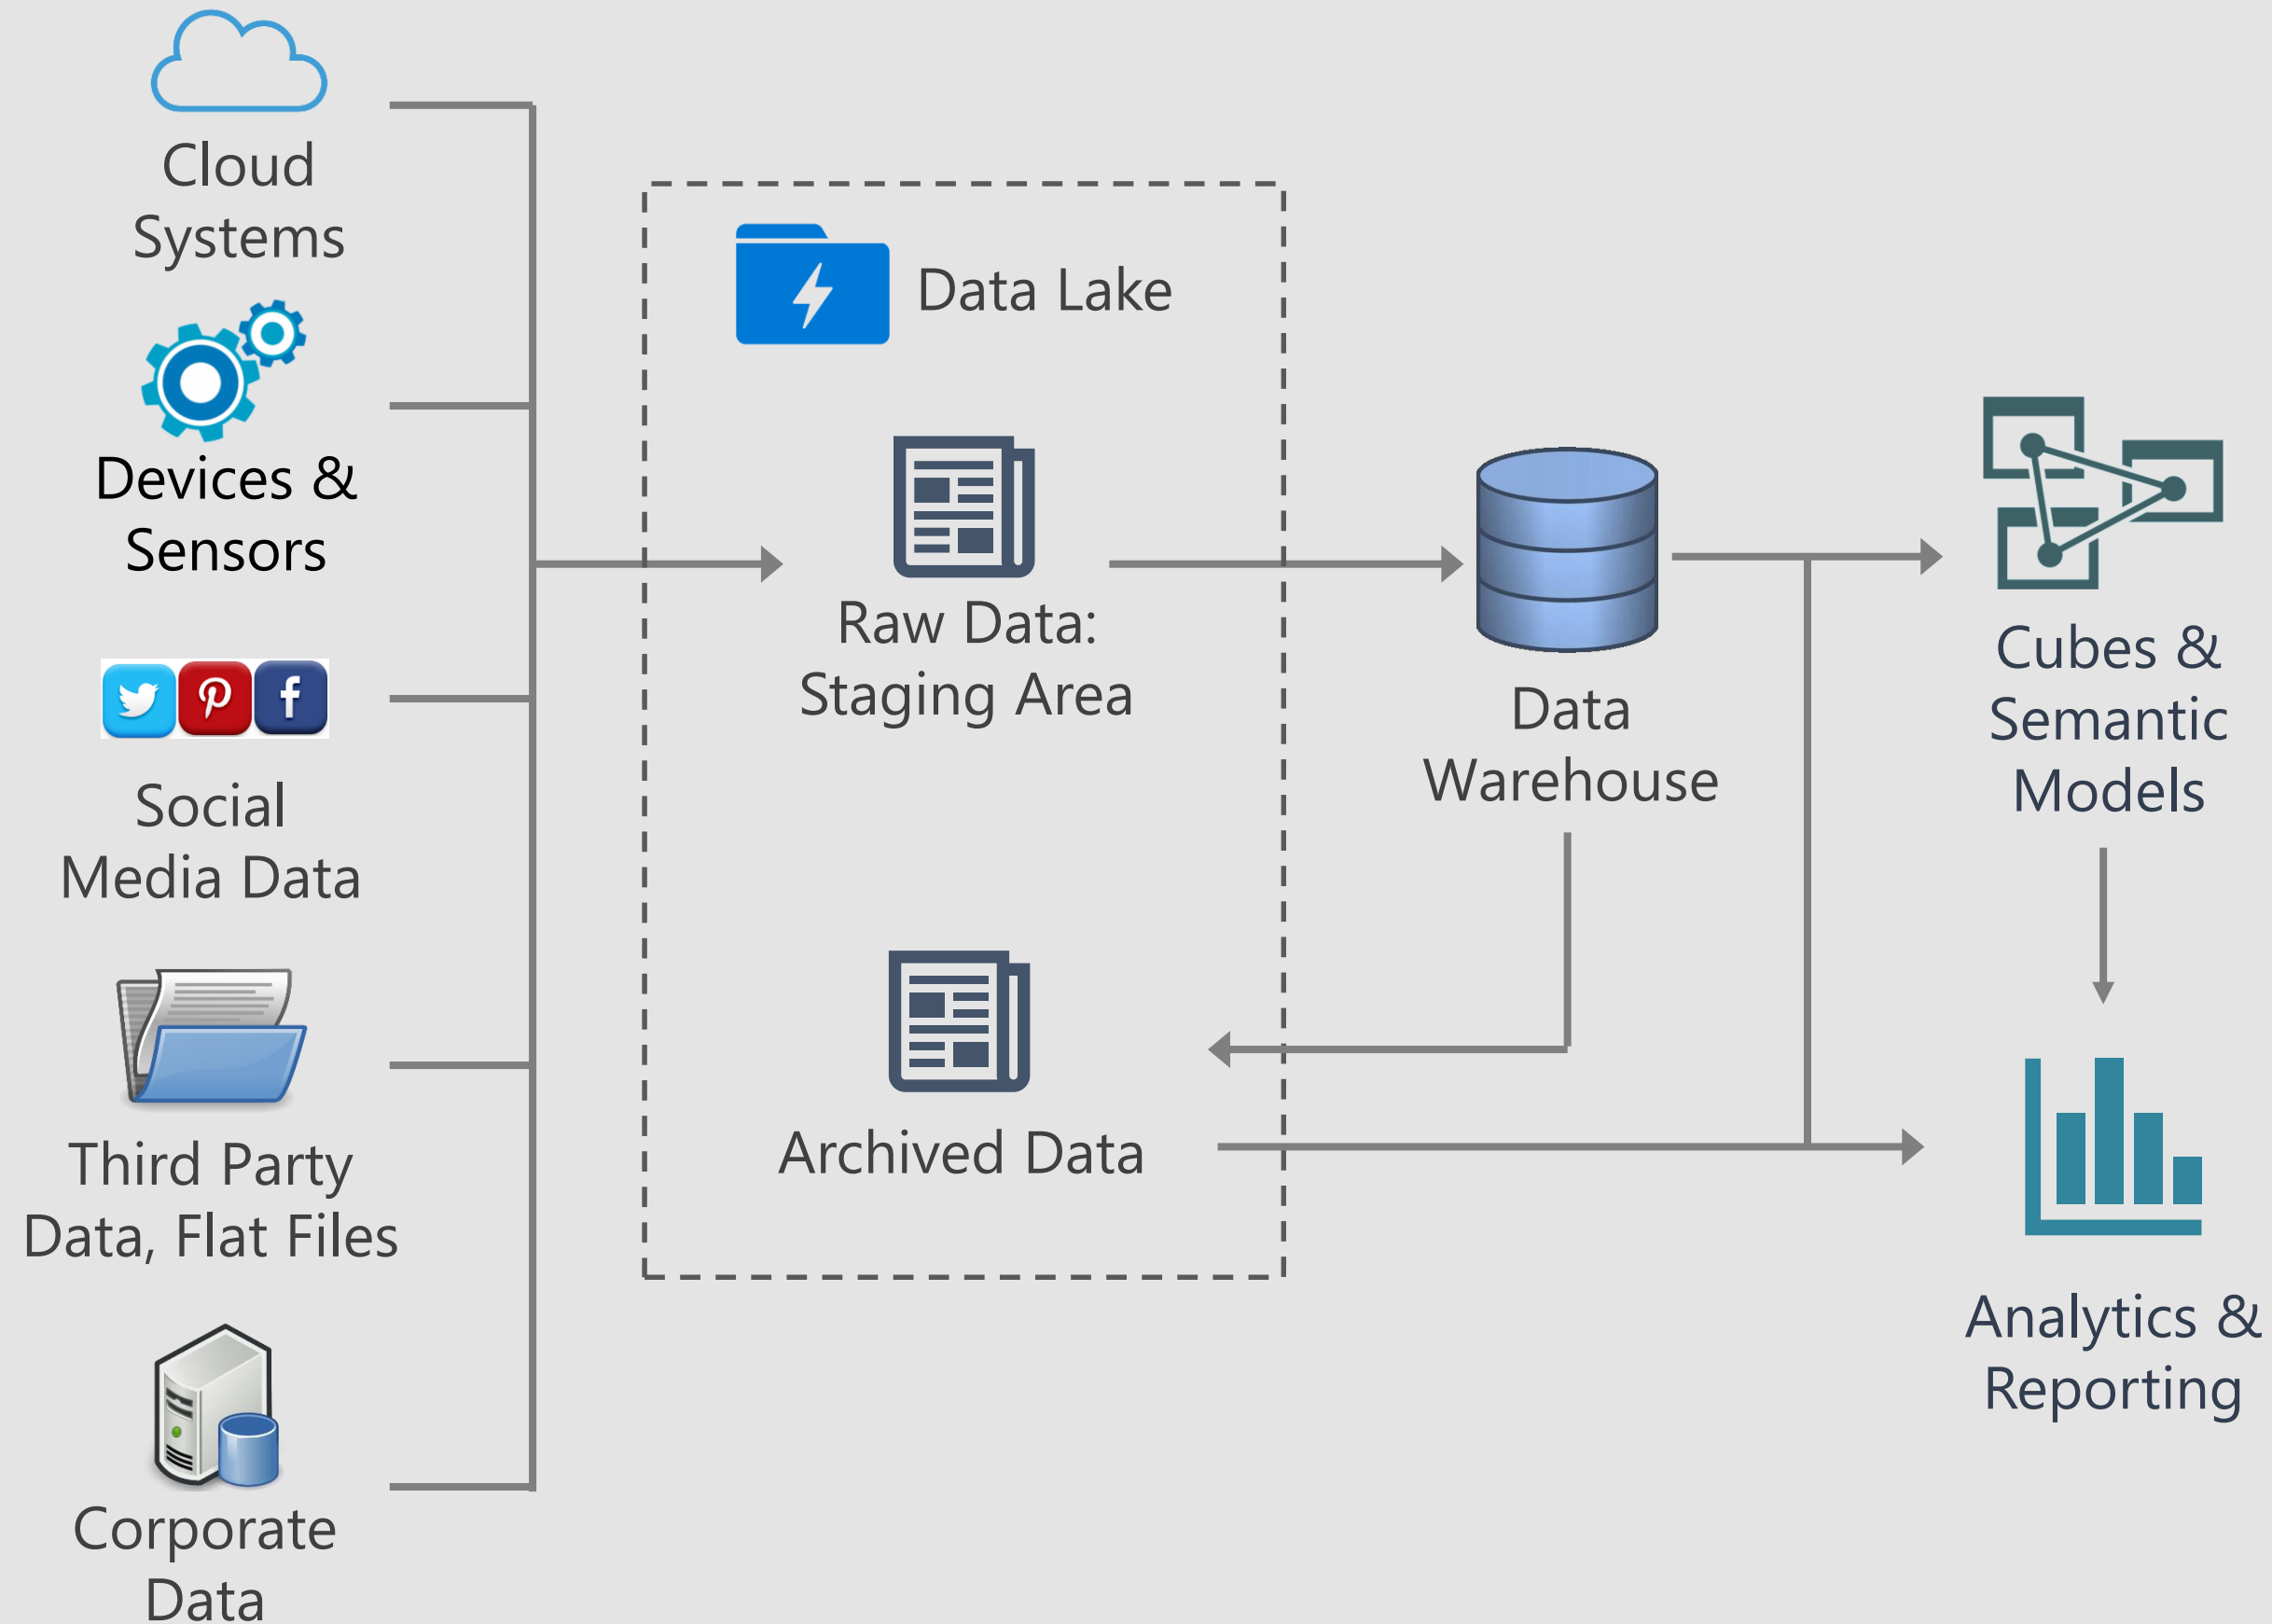
Data Warehouse Staging Area



- ✓ ELT strategy
- ✓ Reduce storage needs in relational platform by using the data lake as landing area
- ✓ Practical use for data stored in the data lake
- ✓ Potentially also handle transformations in the data lake

Data Lake Use Cases

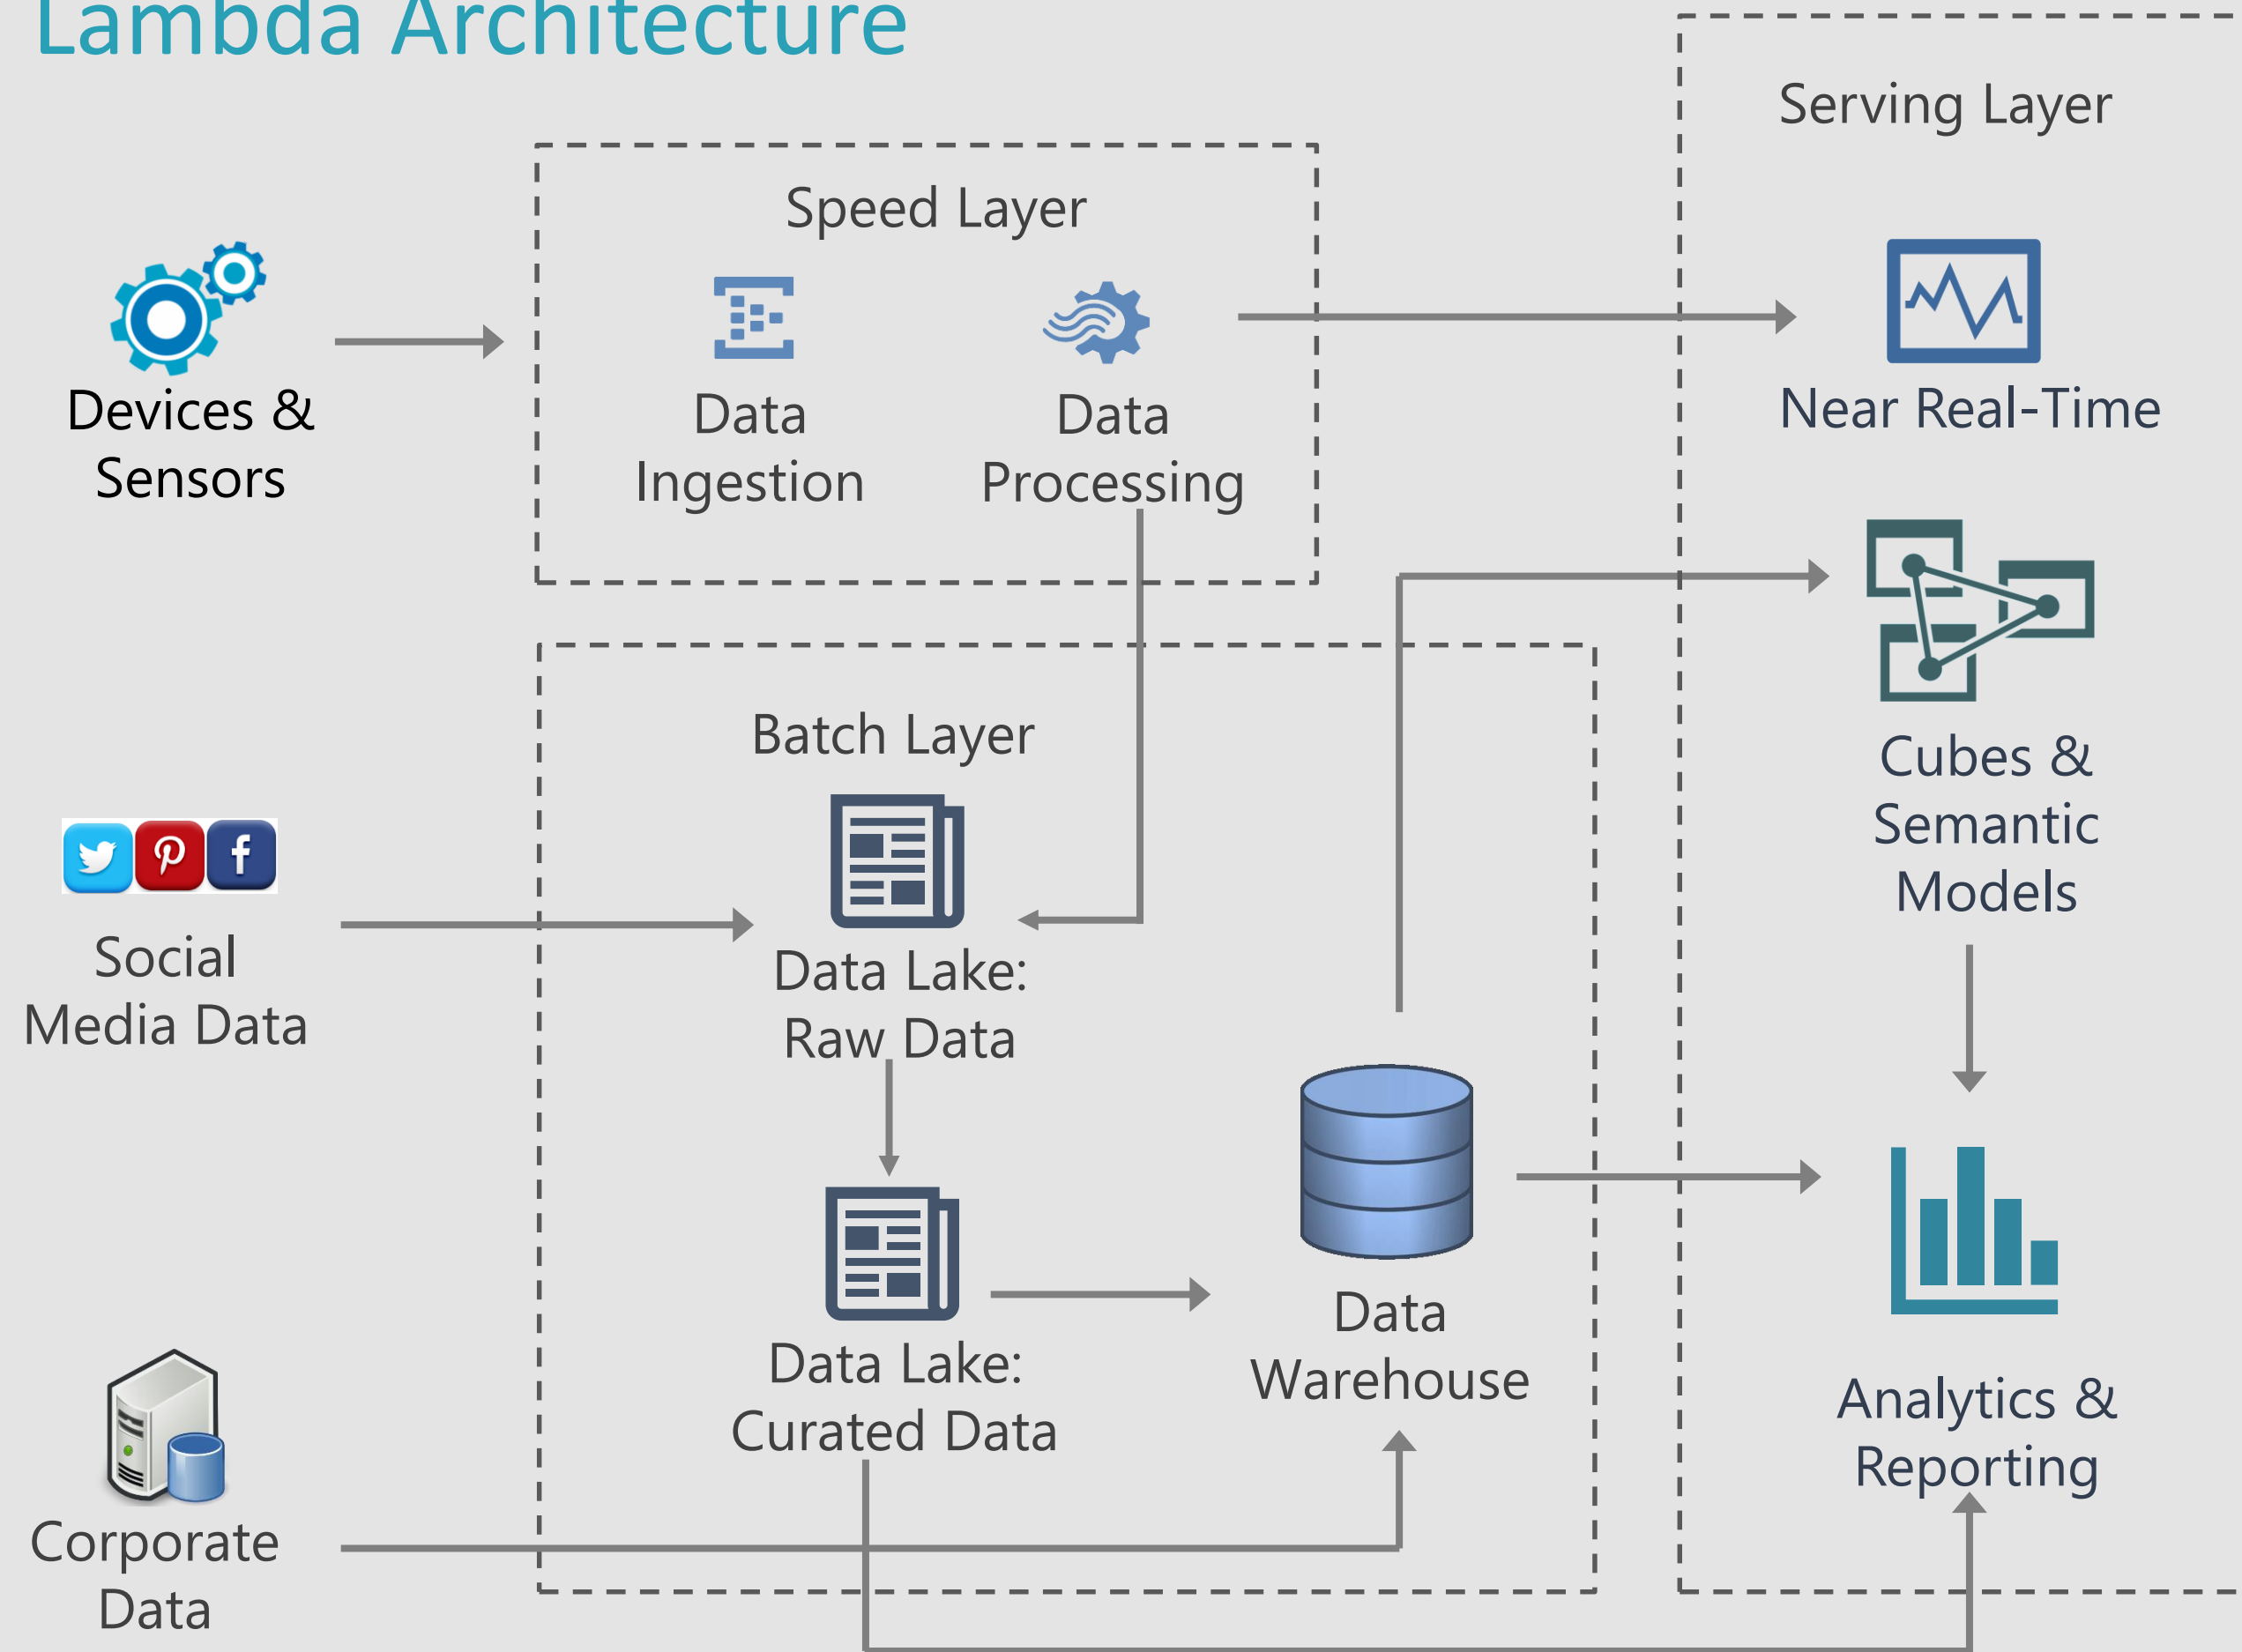
Integration with DW | Data Archival | Centralization




- ✓ Grow around existing DW
- ✓ Aged data available for querying when needed
- ✓ Complement to the DW via data virtualization
- ✓ Federated queries to access current data (relational DB) + archive (data lake)

Data Lake Use Cases

Lambda Architecture



- ✓ Support for low-latency, high-velocity data in near real time
- ✓ Support for batch-oriented operations



What are some initial considerations for deciding if a data lake is right for you?

Is a Data Lake Right For You?

Initial Considerations:

Do you have *non-relational* data?

Do you have *IoT* type of data?


Do you have *advanced analytics scenarios* on unusual datasets?

Do you need to *offload ETL processing (ELT)* and/or *archival data* from a data warehouse?

Readiness:


Are you ready willing to learn *different development patterns* and/or *new technologies*?

Are you ready to handle the *trade-offs of 'schema on read'* vs 'schema on write'?




What are some key differences between
a data warehouse & a data lake?

Data Lake + Data Warehouse: Inverse Relationship



Data Lake focuses on:

- ✓ Agility
- ✓ Flexibility
- ✓ Easy data acquisition
- ✓ Early exploration activities



Data warehouse focuses on:

- ✓ Cleansed, user-friendly data
- ✓ Reliability
- ✓ Standardization
- ✓ Process-oriented operationalization

Schema on Read

↓ Less effort

↑ More effort

Data acquisition

Data retrieval

Schema on Write

↑ More effort

↓ Less effort

Data Lake Challenges

Technology

- ✓ Addt'l component(s) in a multi-layered architecture

- ✓ Unknown storage & scalability

- ✓ Data retrieval

- ✓ Working with uncurated data

- ✓ Performance

- ✓ Change management

Process

- ✓ Right balance of deferred work vs. up-front work

- ✓ Ignoring established best practices for data management

- ✓ Data quality

- ✓ Governance

- ✓ Security

- ✓ Disaster recovery for large solutions

People

- ✓ Expectations & trust

- ✓ Data stewardship


- ✓ Redundant effort

- ✓ Skills required to effectively use the data

Big Data in Azure

Big Data in Azure

Compute (PaaS)



Azure
HDInsight



Azure
Databricks



Azure
Machine Learning

Compute (IaaS)



Azure Data
Science VMs



Hadoop on a cluster
of Azure virtual
machines

Storage

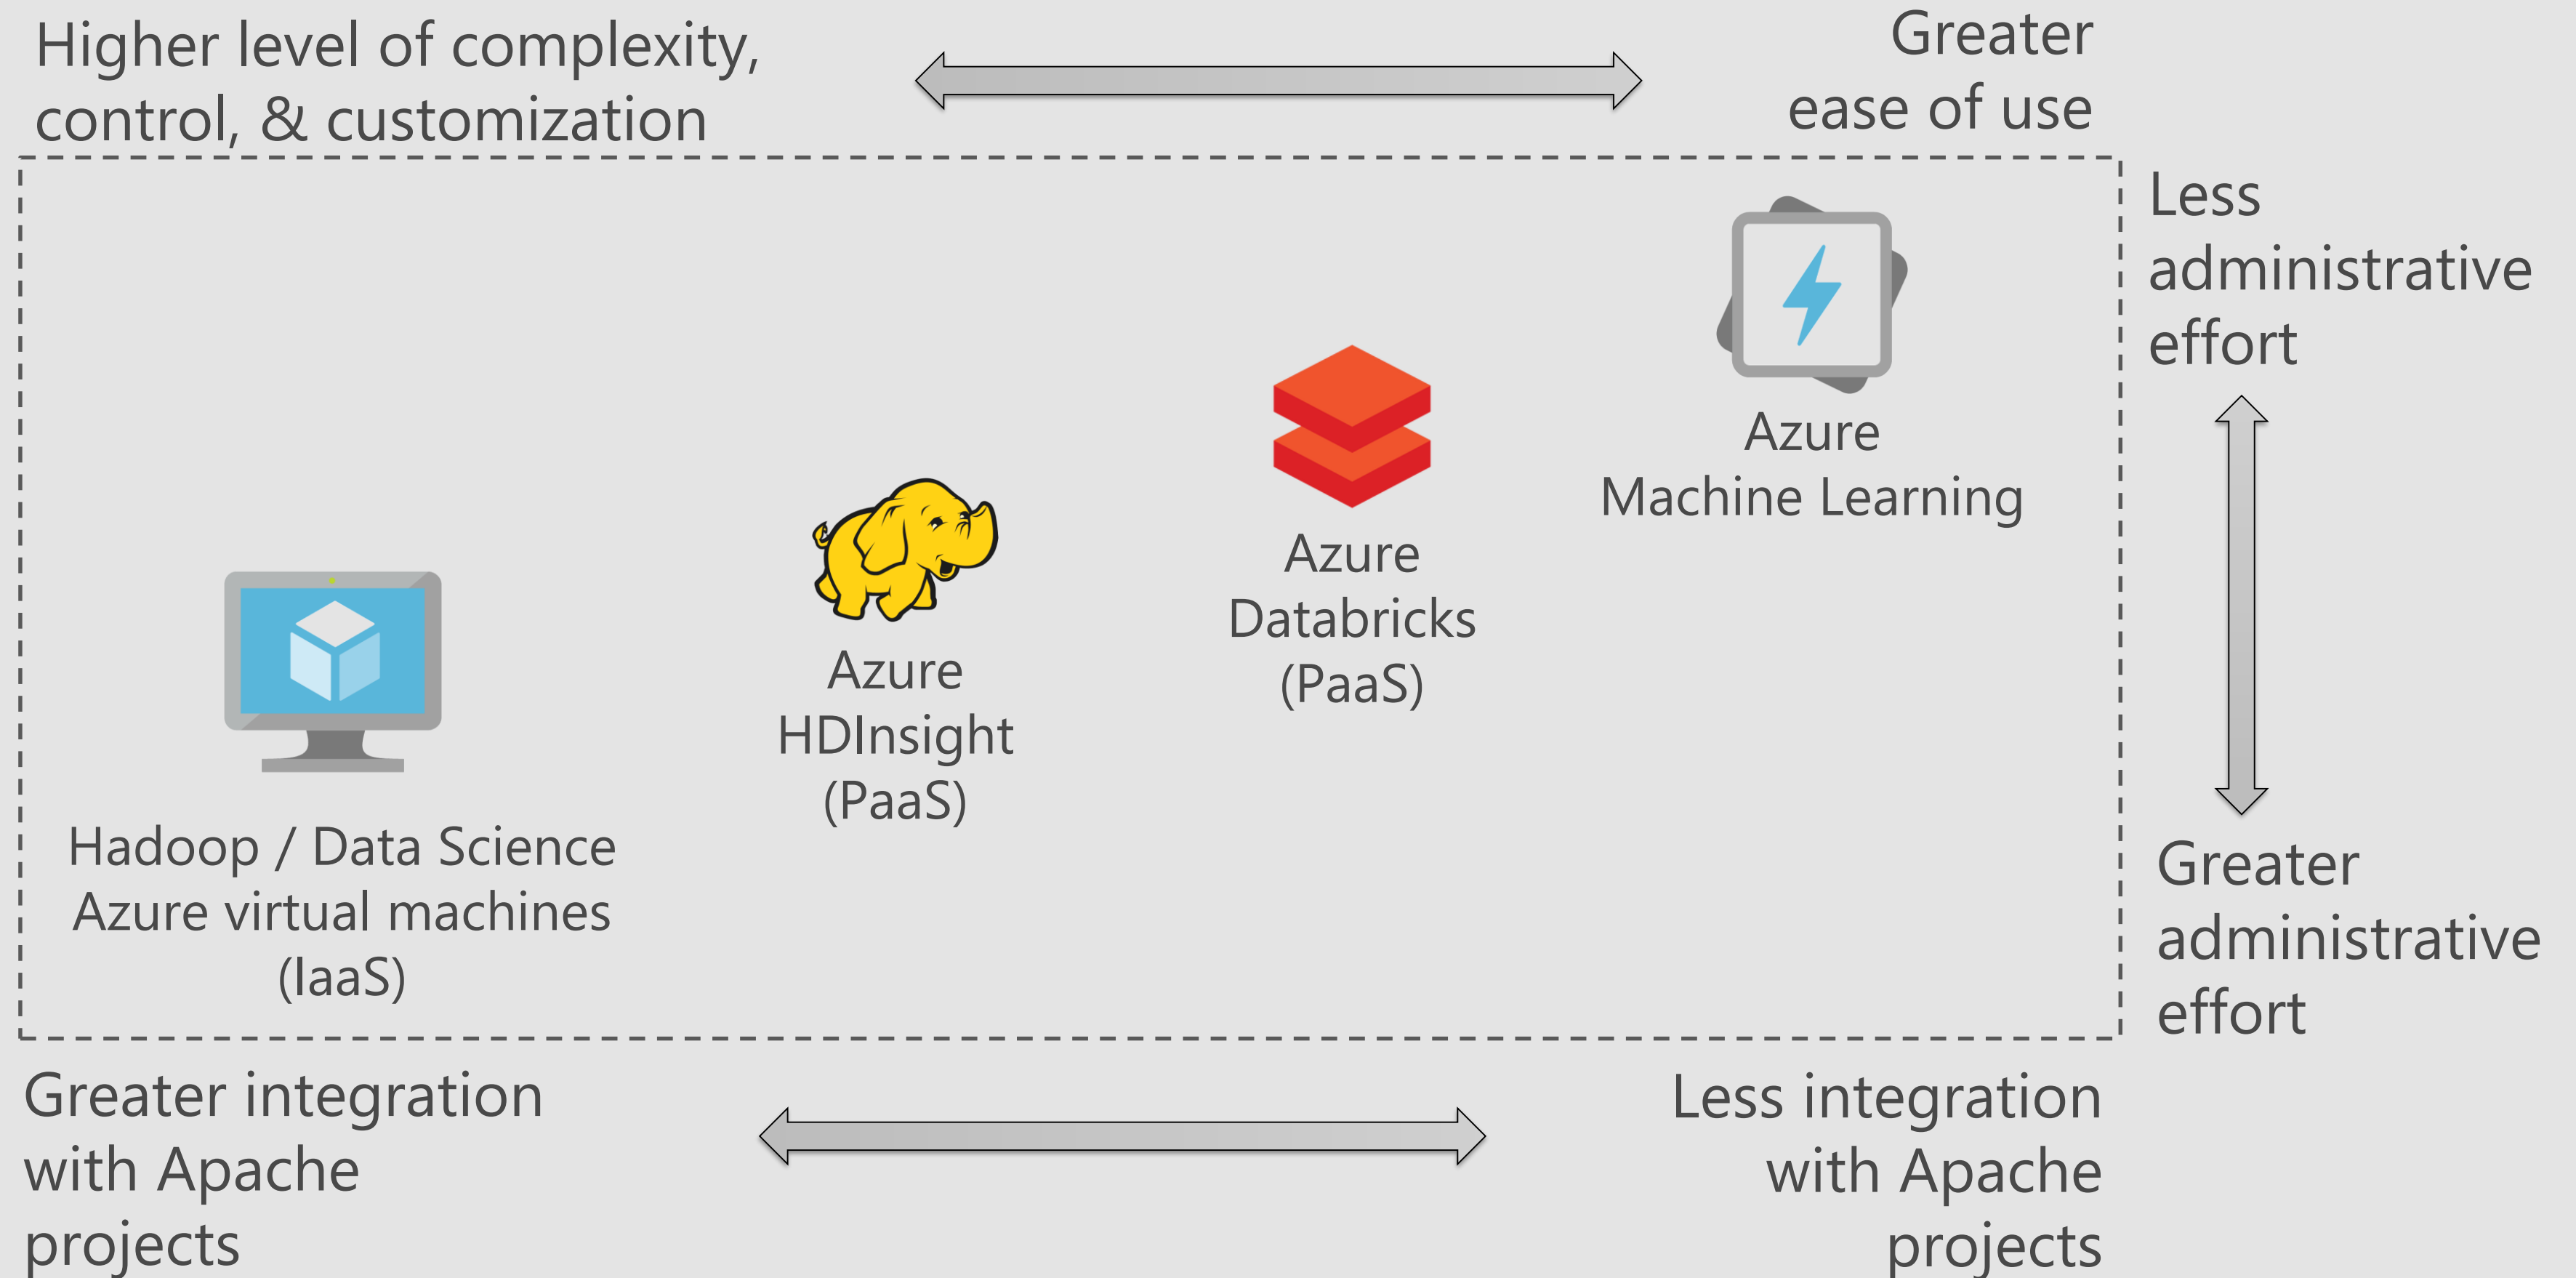


Azure Data Lake Store (Gen2)



Azure Storage

Big Data in Azure: Compute



Deciding Between Compute Services



Hadoop VM



HDInsight



Databricks



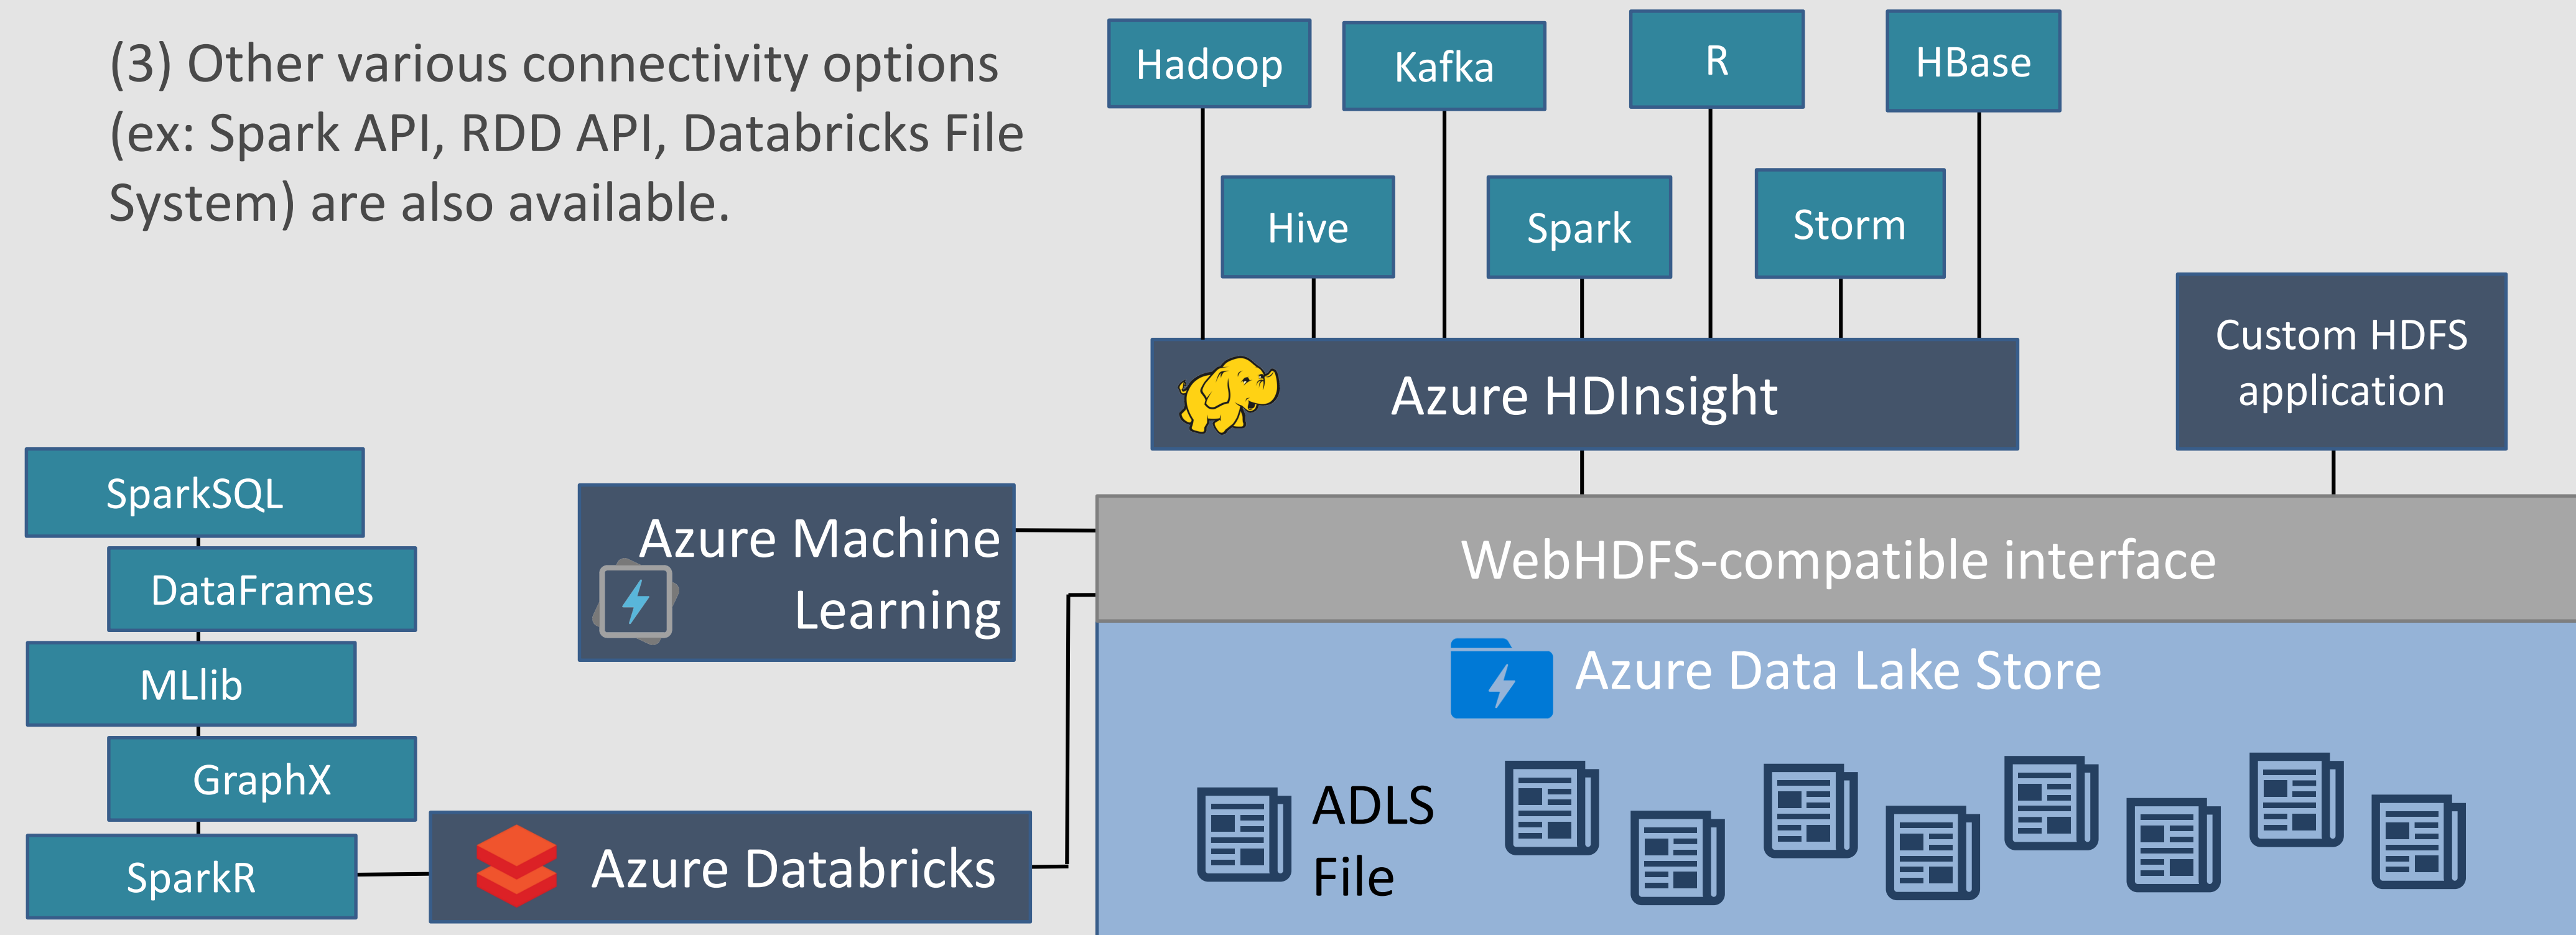
Azure ML

Type:	IaaS	PaaS	PaaS	SaaS
Purpose:	Running your own cluster of Hadoop virtual machines	Running a managed cluster	Running optimized Spark framework	Running packaged AI, R or Python Script
Suitable for:	Full control over everything; investment in distributions such as Hortonworks, Cloudera, MapR	Integration with open source Apache projects (ex: Hive, Storm, Kafka, Spark, etc)	Collaborative notebooks, easier deployments	An ideal initial entry point for sandbox experimentation

Intro to Azure Data Lake

Azure Data Lake Store - Compatibility

- (1) **WebHDFS endpoint** (https://) allows integration with open source projects.
- (2) **“AzureDataLakeFilesystem”** (adl://) provides additional performance enhancements not available in WebHDFS.
- (3) Other various connectivity options (ex: Spark API, RDD API, Databricks File System) are also available.



Azure Data Lake Store – Distributed File System

Files of any size can be stored because **ADLS is a distributed system** which file contents are divided up across backend storage nodes.

A read operation on the file is also **parallelized across the nodes**.

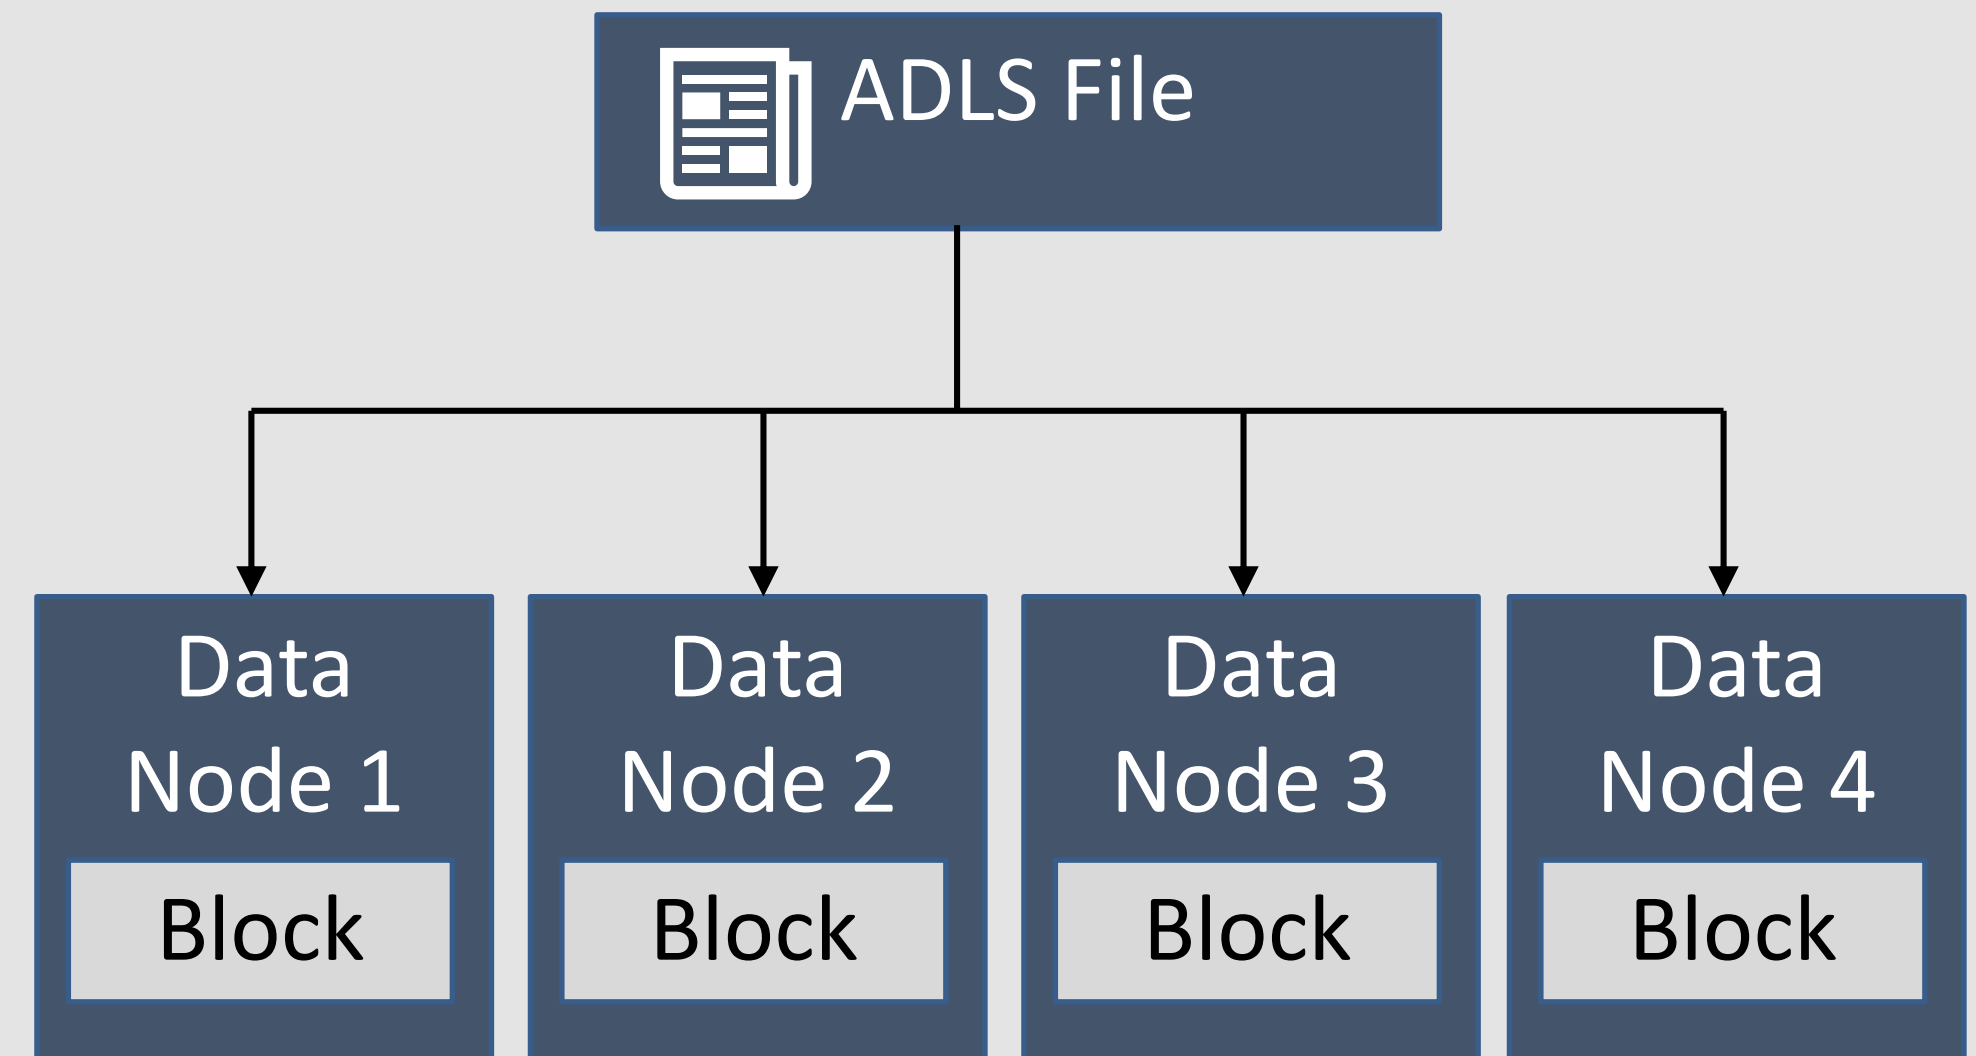
Blocks are also replicated for fault tolerance.



The **ideal file size** in ADLS is 256MB – 2GB in size.

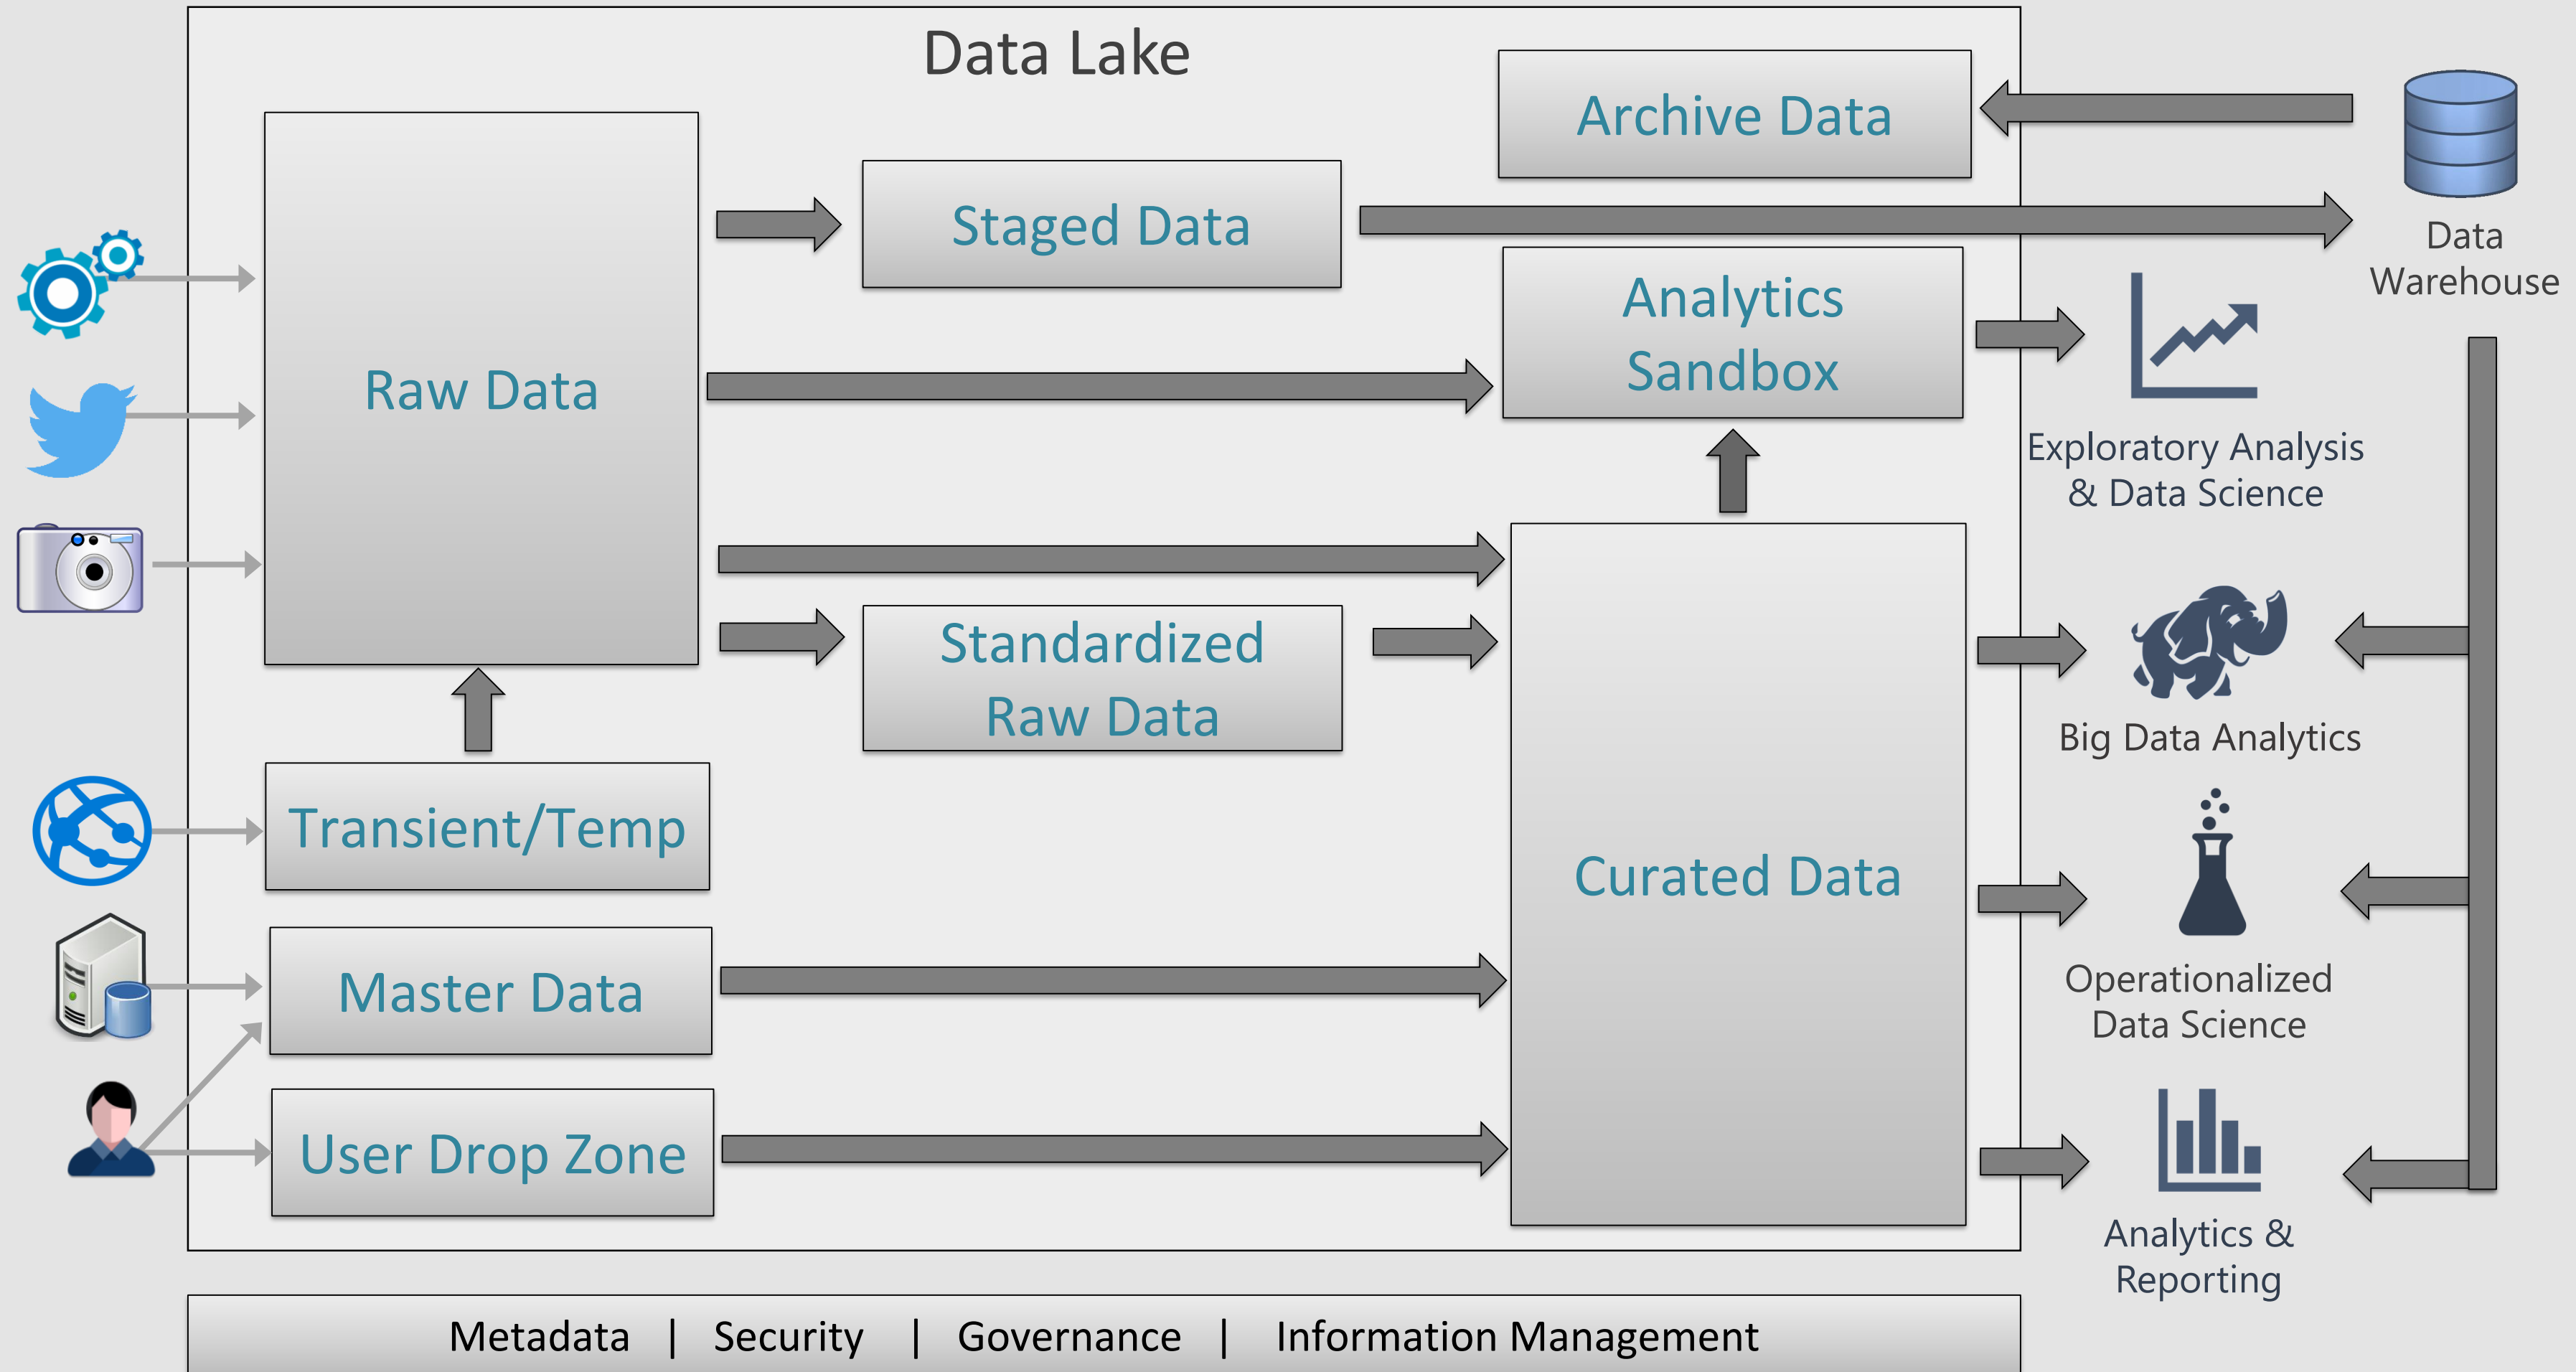
Many very tiny files introduces significant overhead which reduces performance. This is a well-known issue with storing data in HDFS. Techniques:

- Append-only data streams
- Consolidation of data into larger files

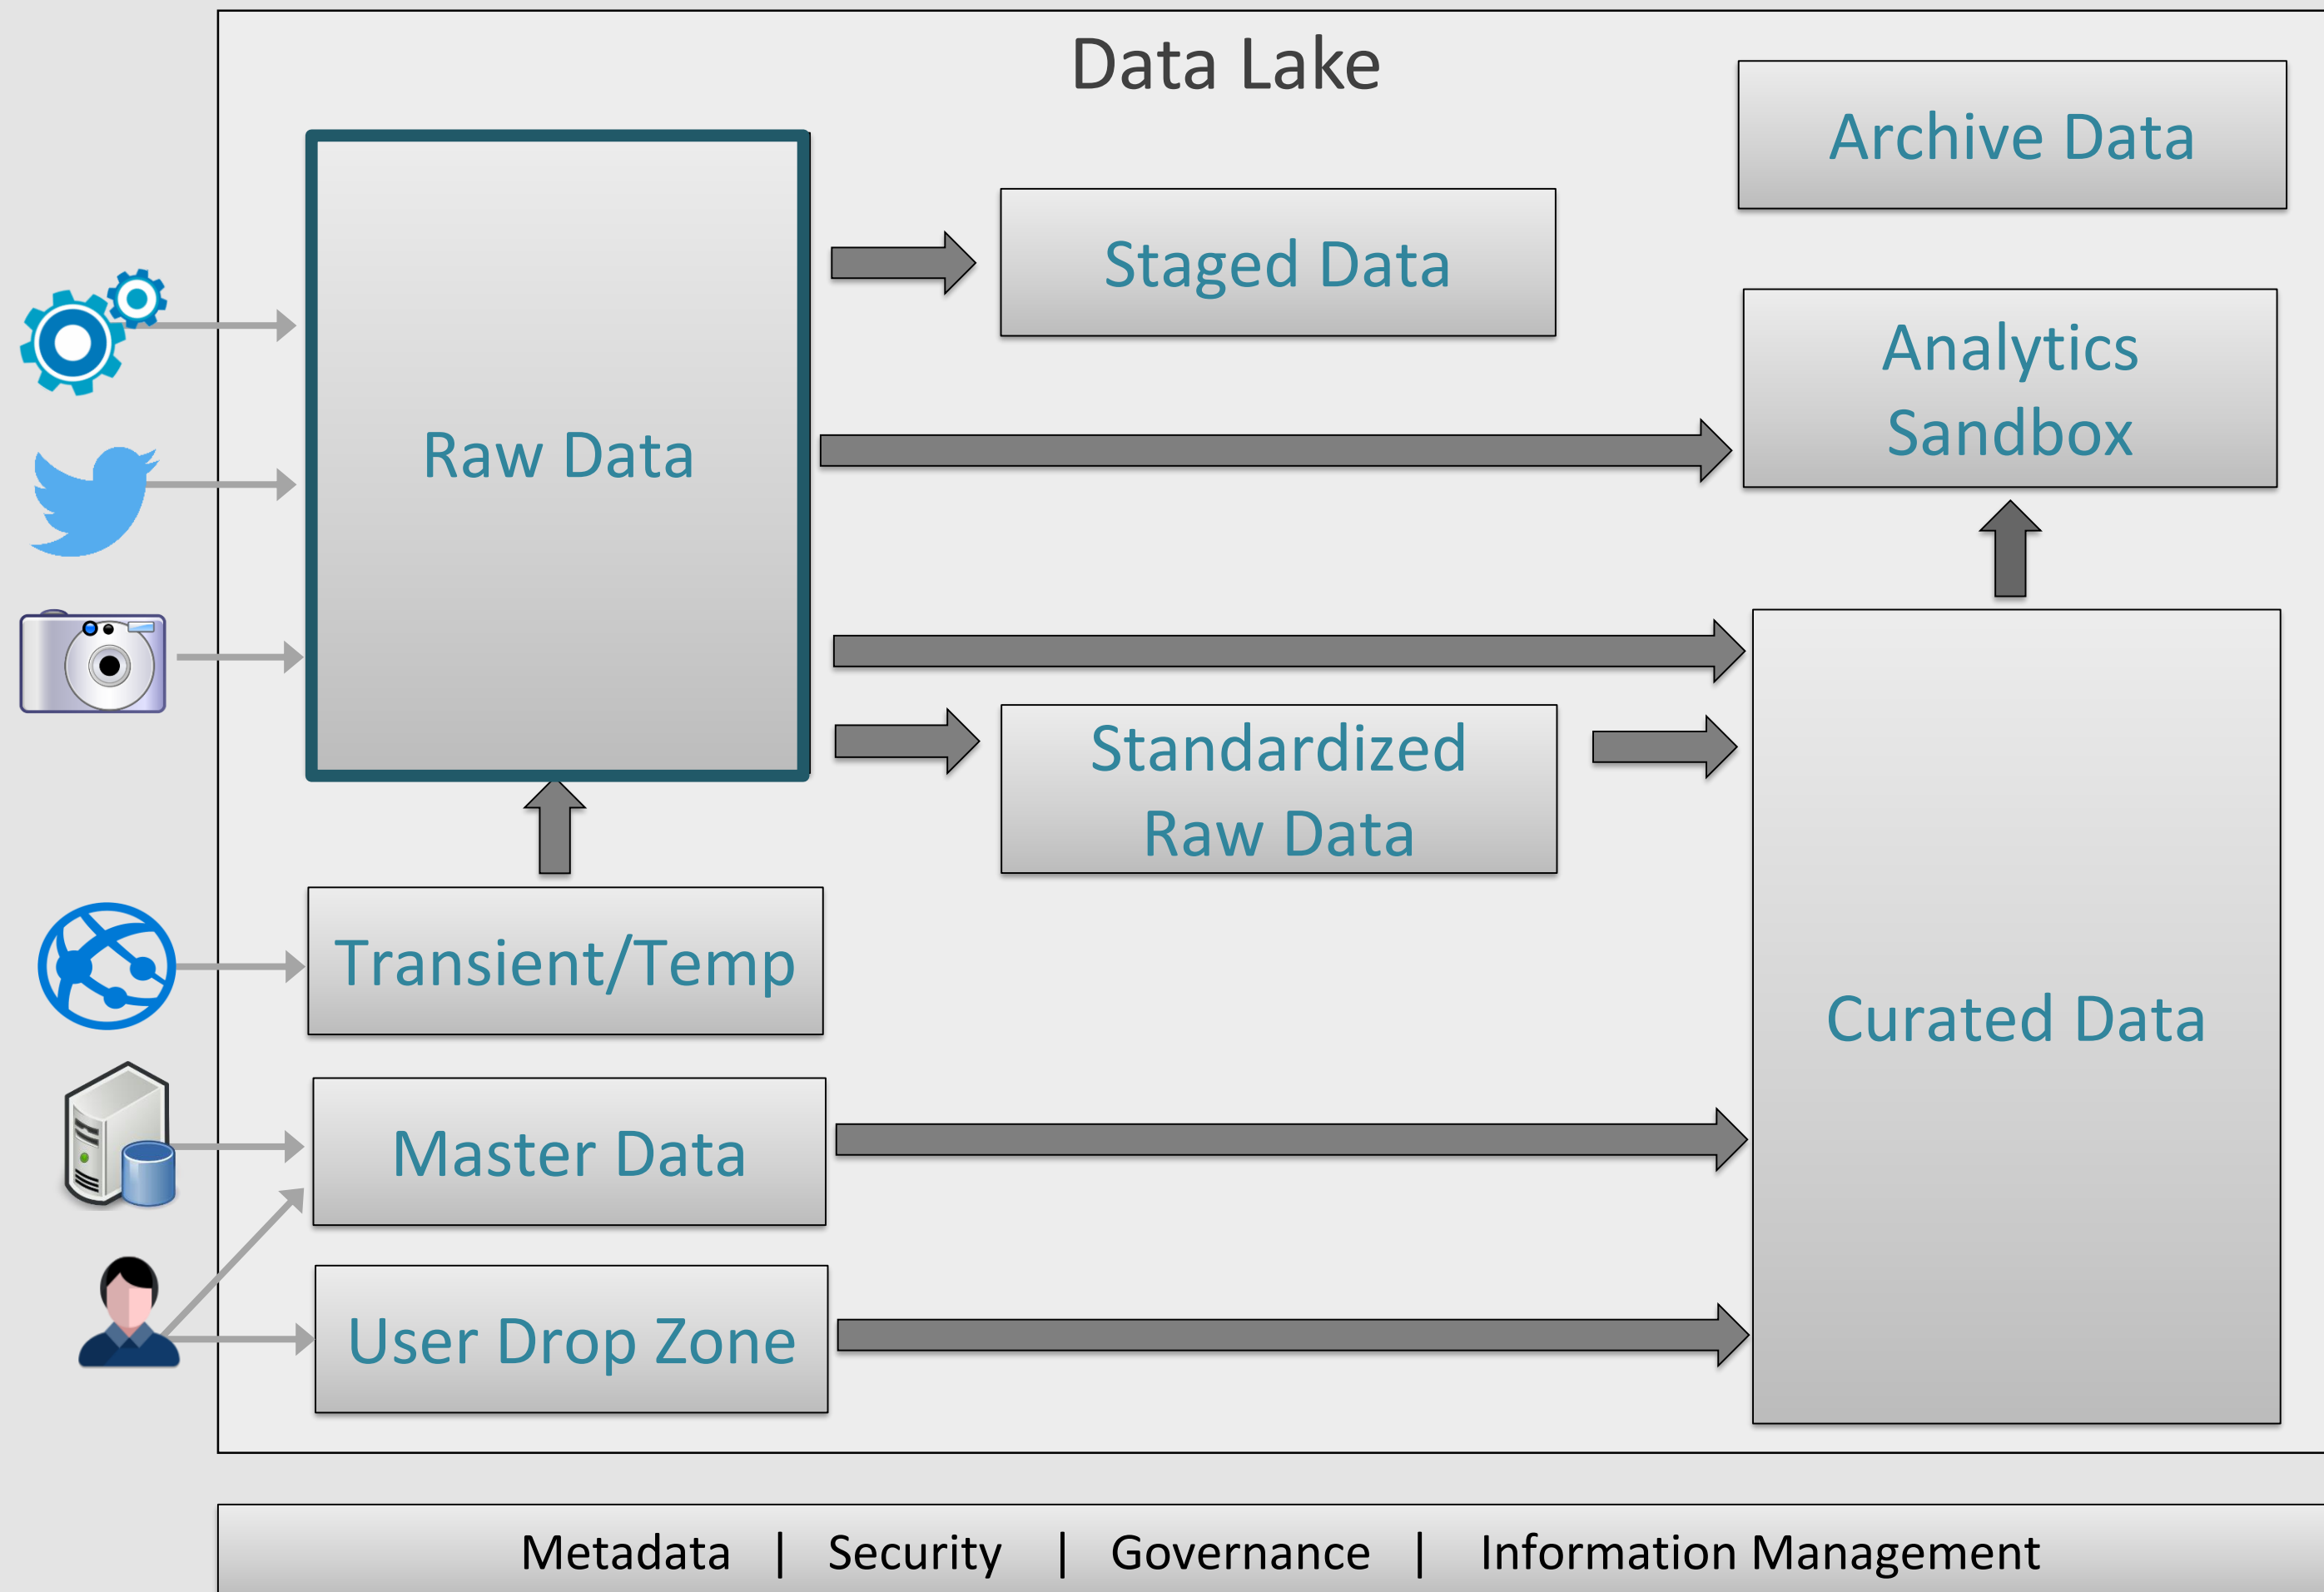


Designing the Structure of a Data Lake

Designing the Zones of a Data Lake

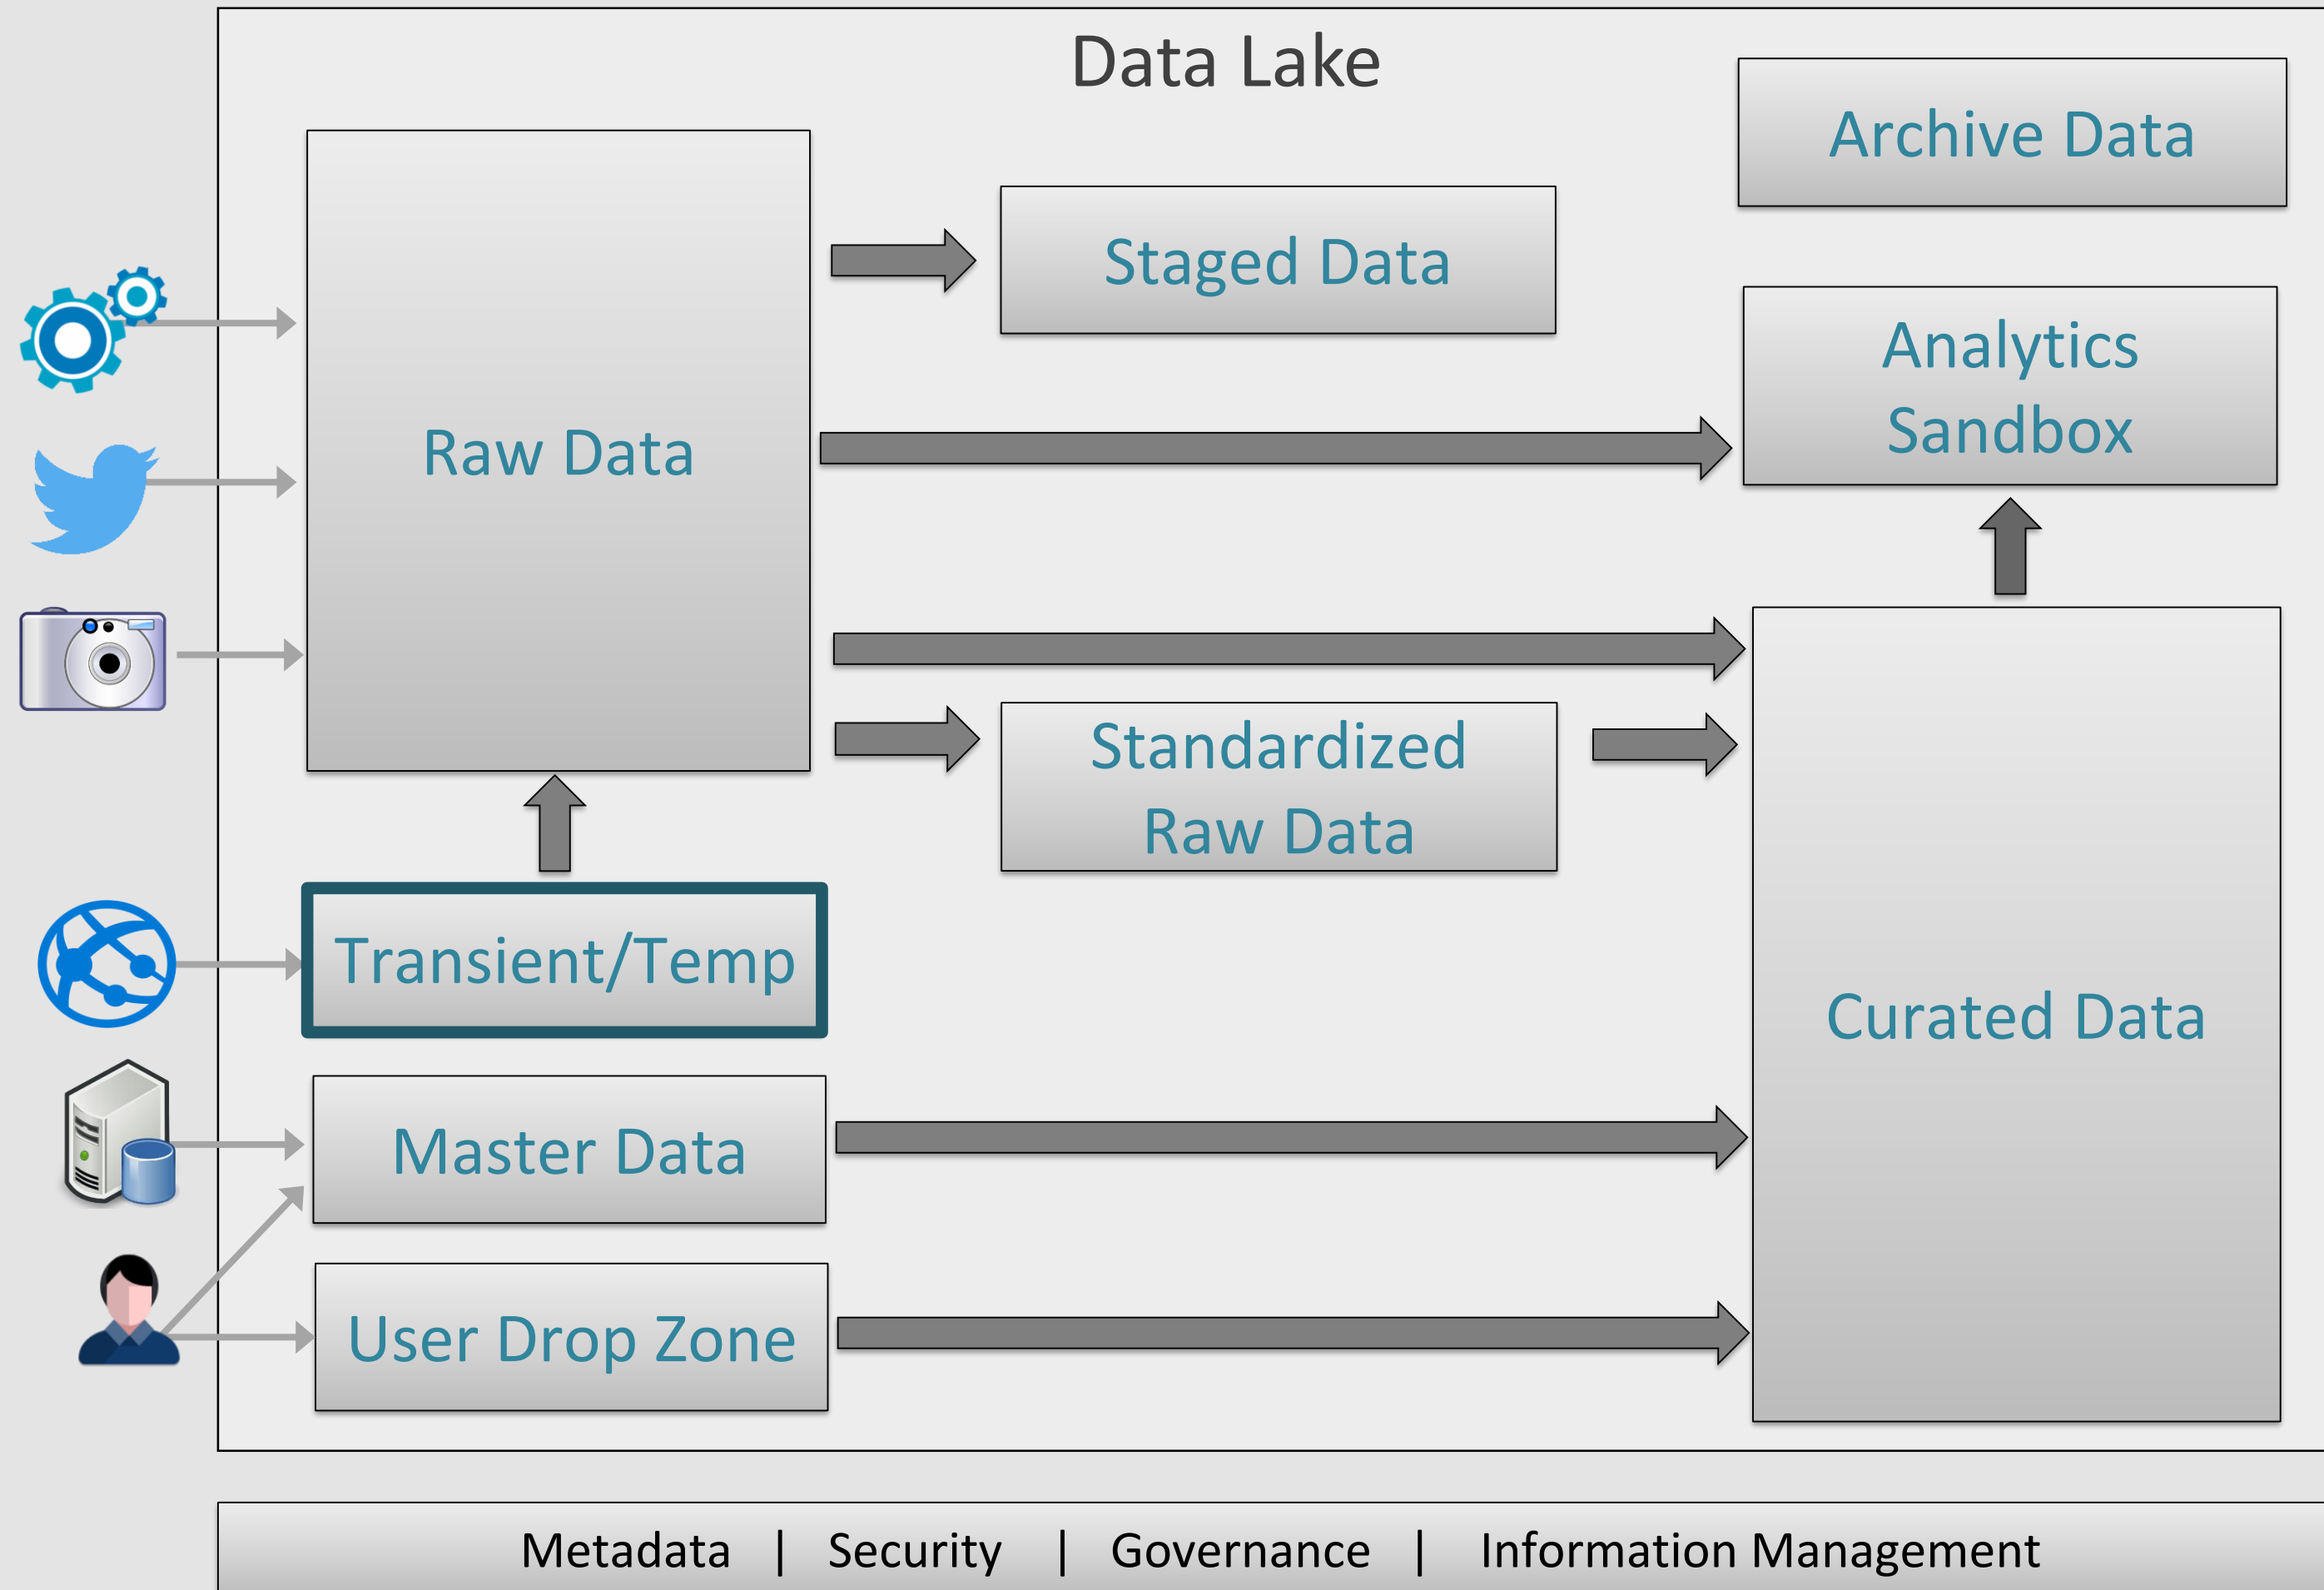


Raw Data Zone



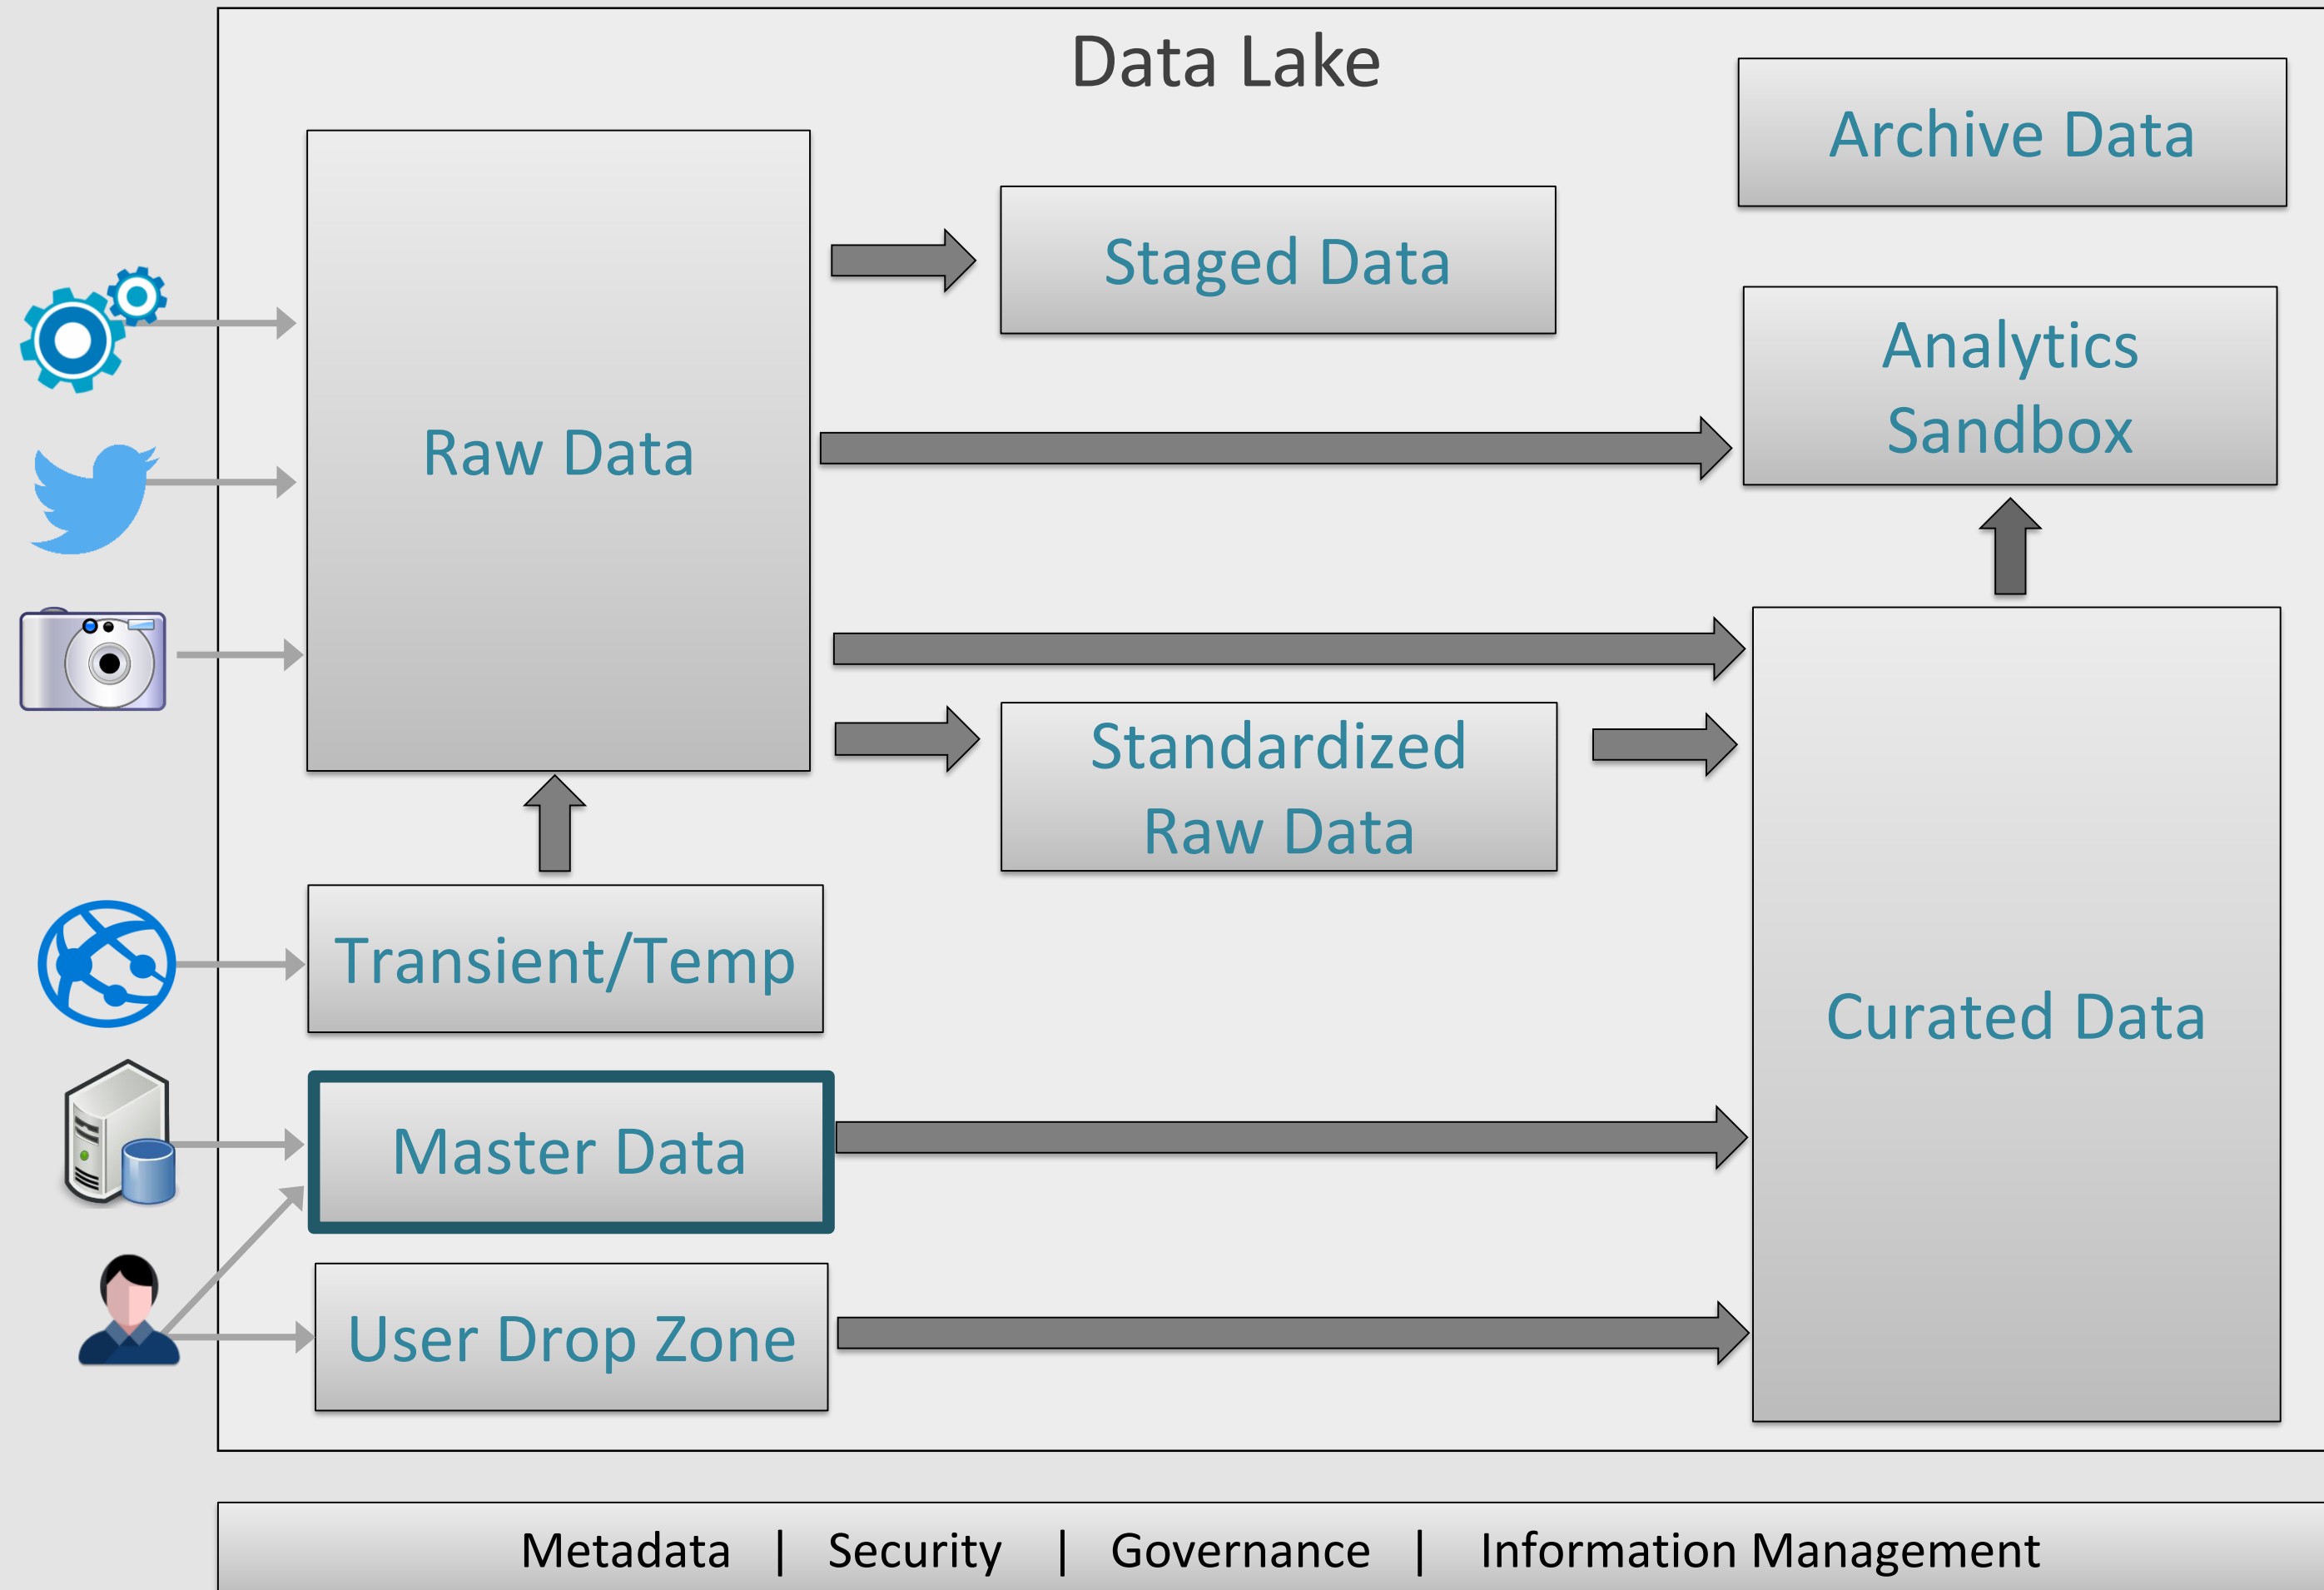
- ✓ Storage in native format for any type of data
- ✓ Exact copy from the source
- ✓ Immutable to change
- ✓ Typically append-only
- ✓ History retained indefinitely
- ✓ Extremely limited access to the Raw Data Zone – no operationalized usage
- ✓ Everything downstream from here can be regenerated from raw data

Transient/Temp Zone



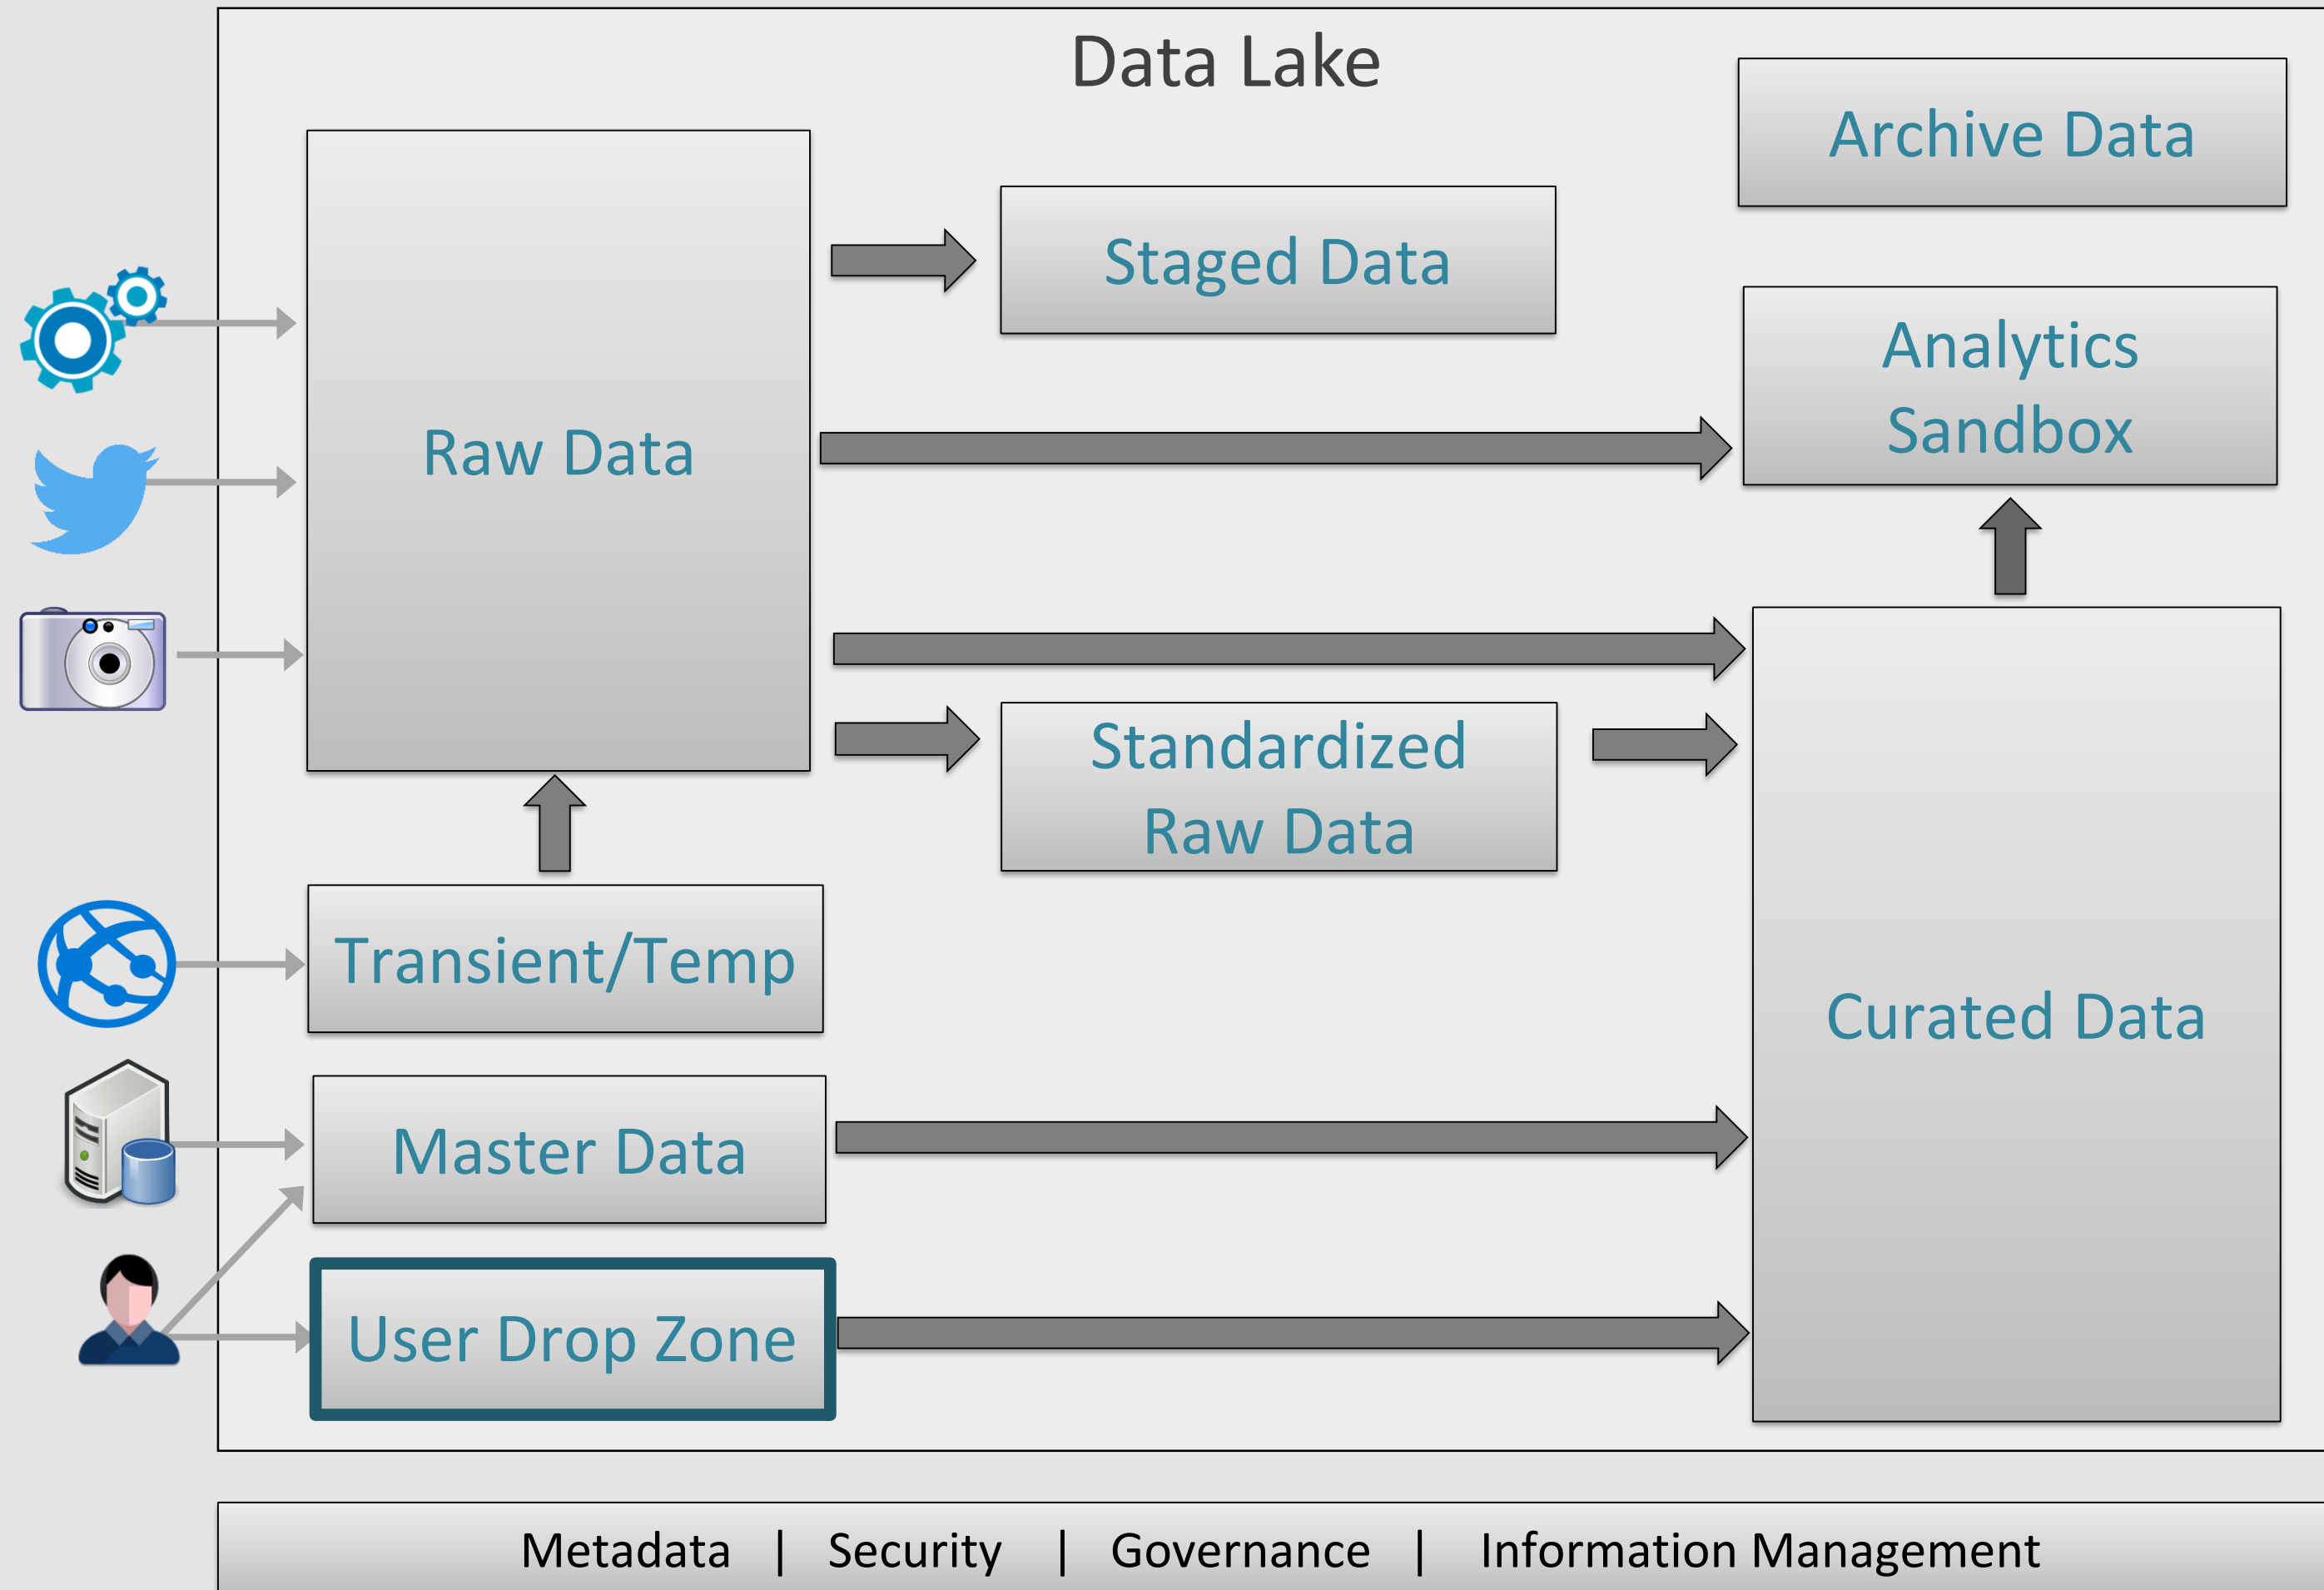
- ✓ Selectively utilized
- ✓ Useful when data quality checks or validation is required before the data is routed to the Raw Data Zone for retention
- ✓ Useful when you need a “New Data” zone separate from Raw Data Zone
(ex: to ensure that jobs pulling data from Raw receive consistent data)
- ✓ Could contain transient, low-latency data
(aka ‘speed layer’)

Master Data Zone



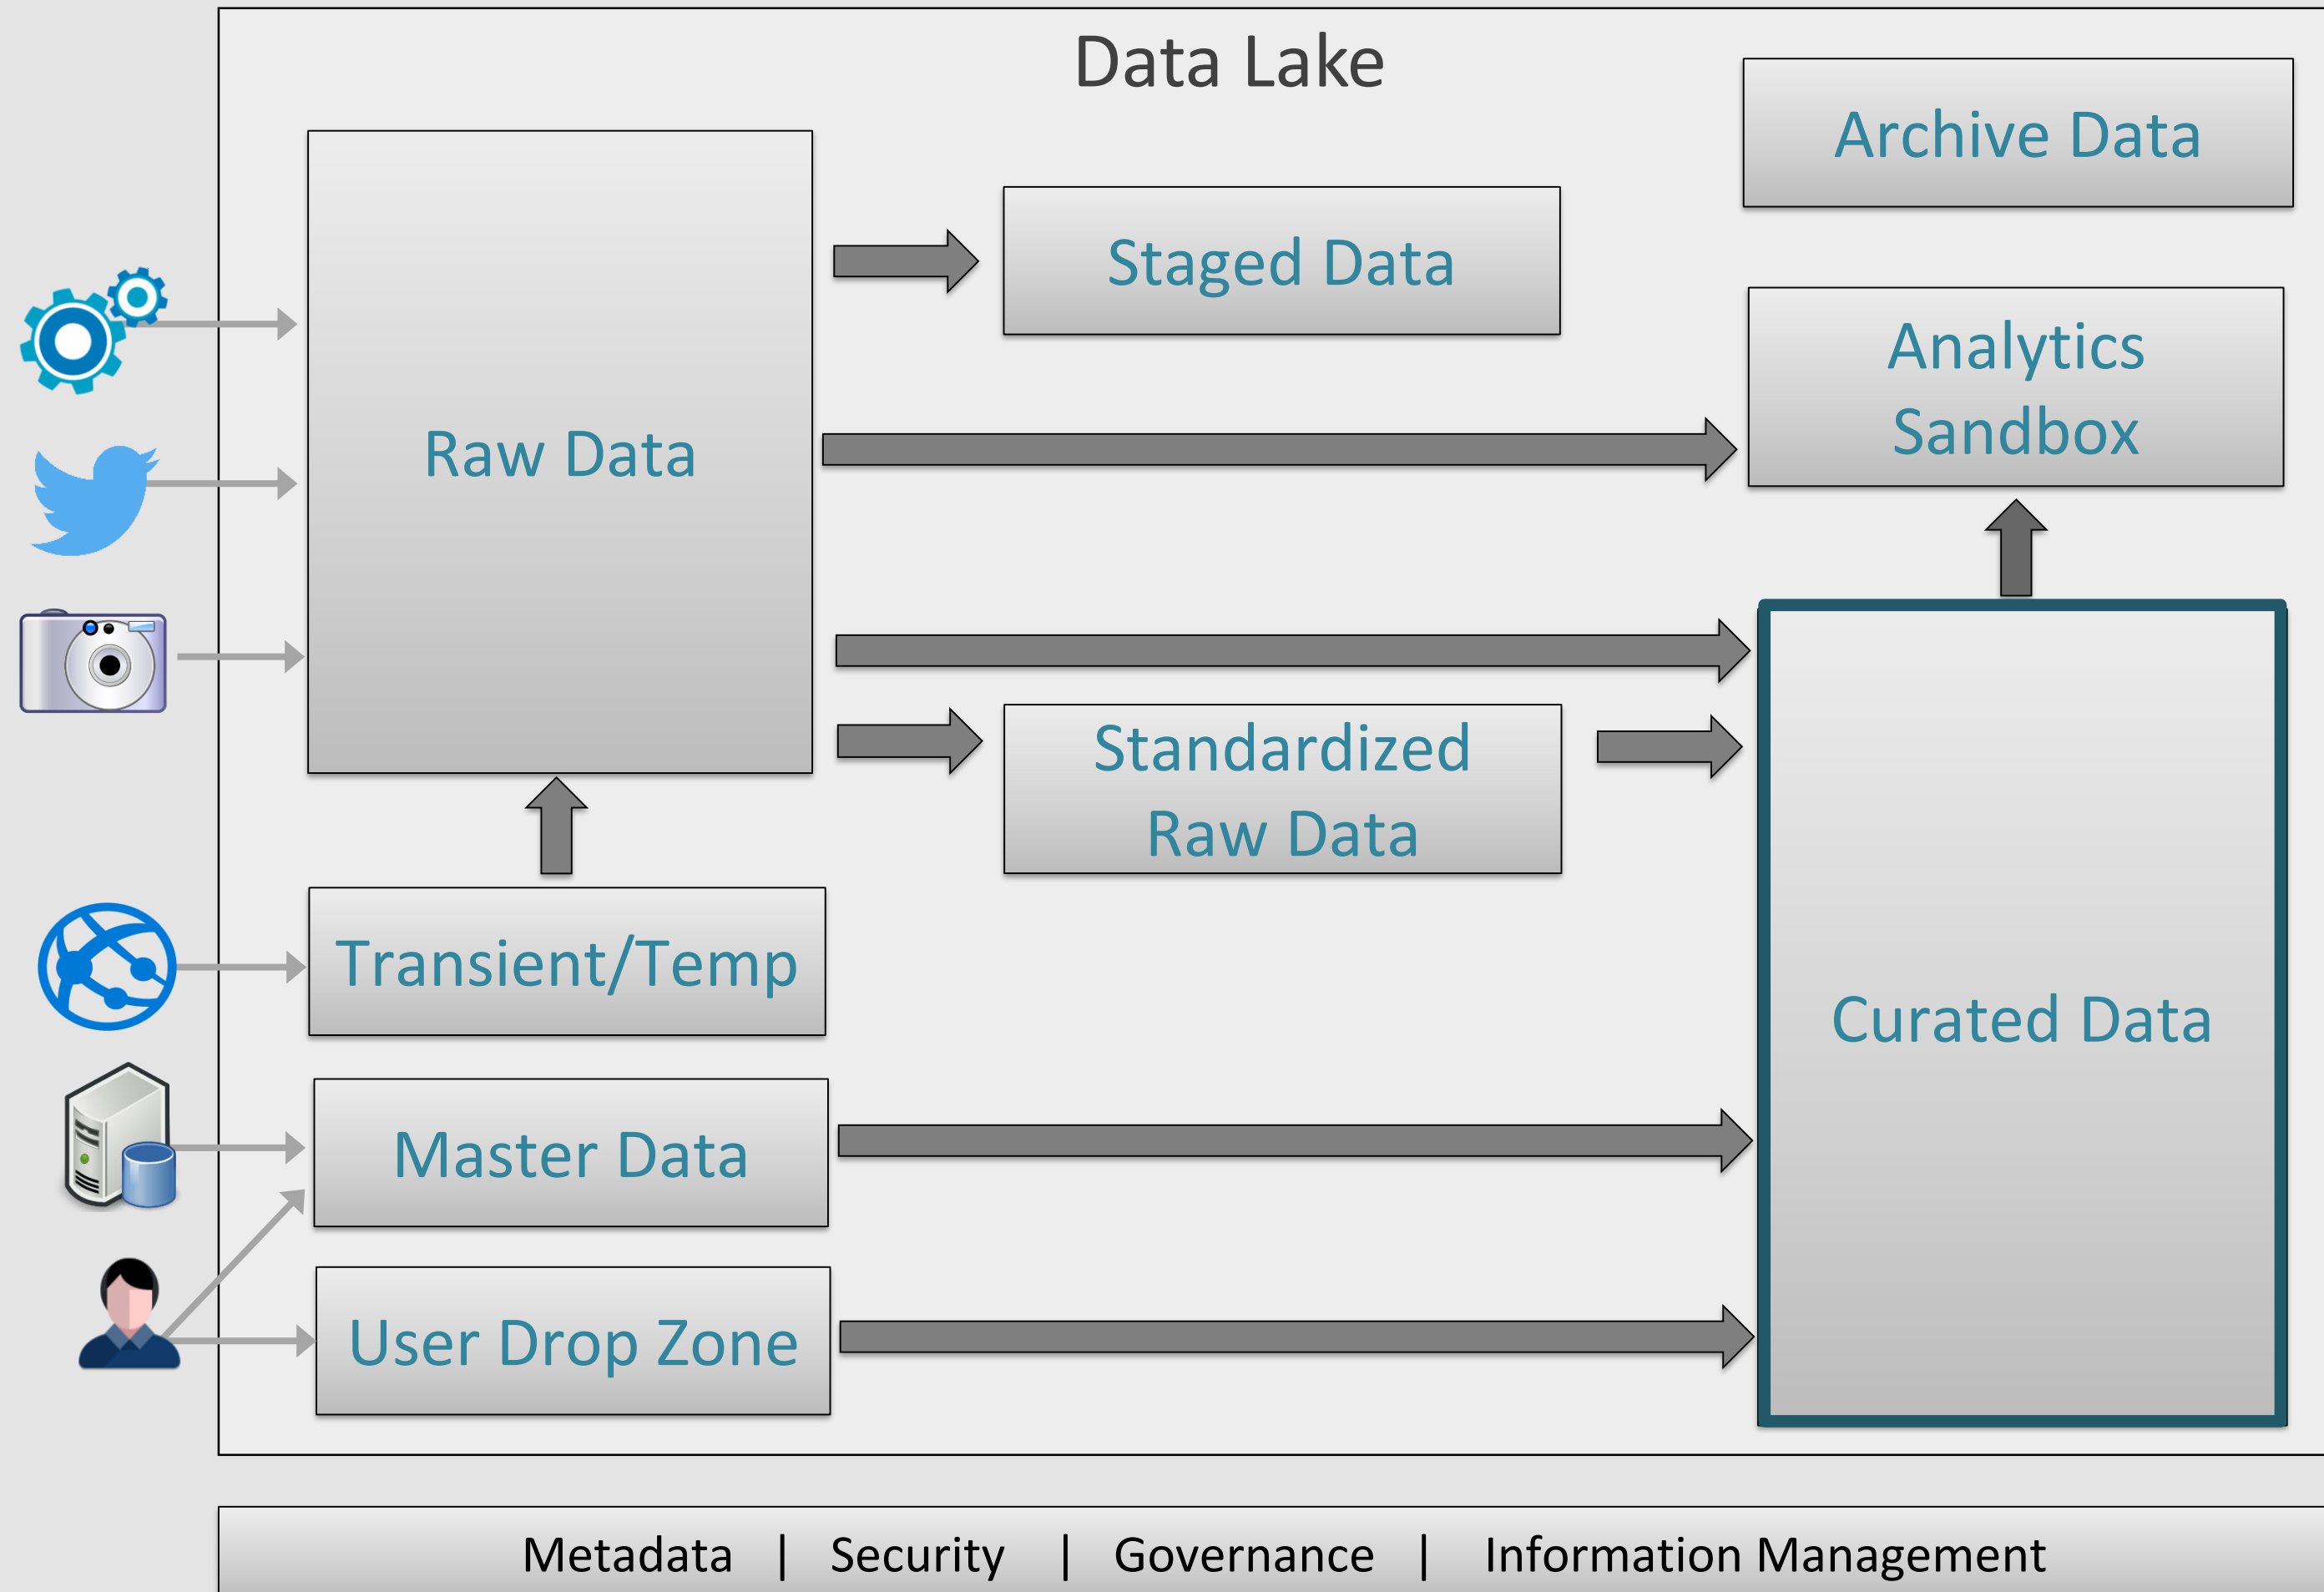
✓ Reference data to augment analysis

User Drop Zone



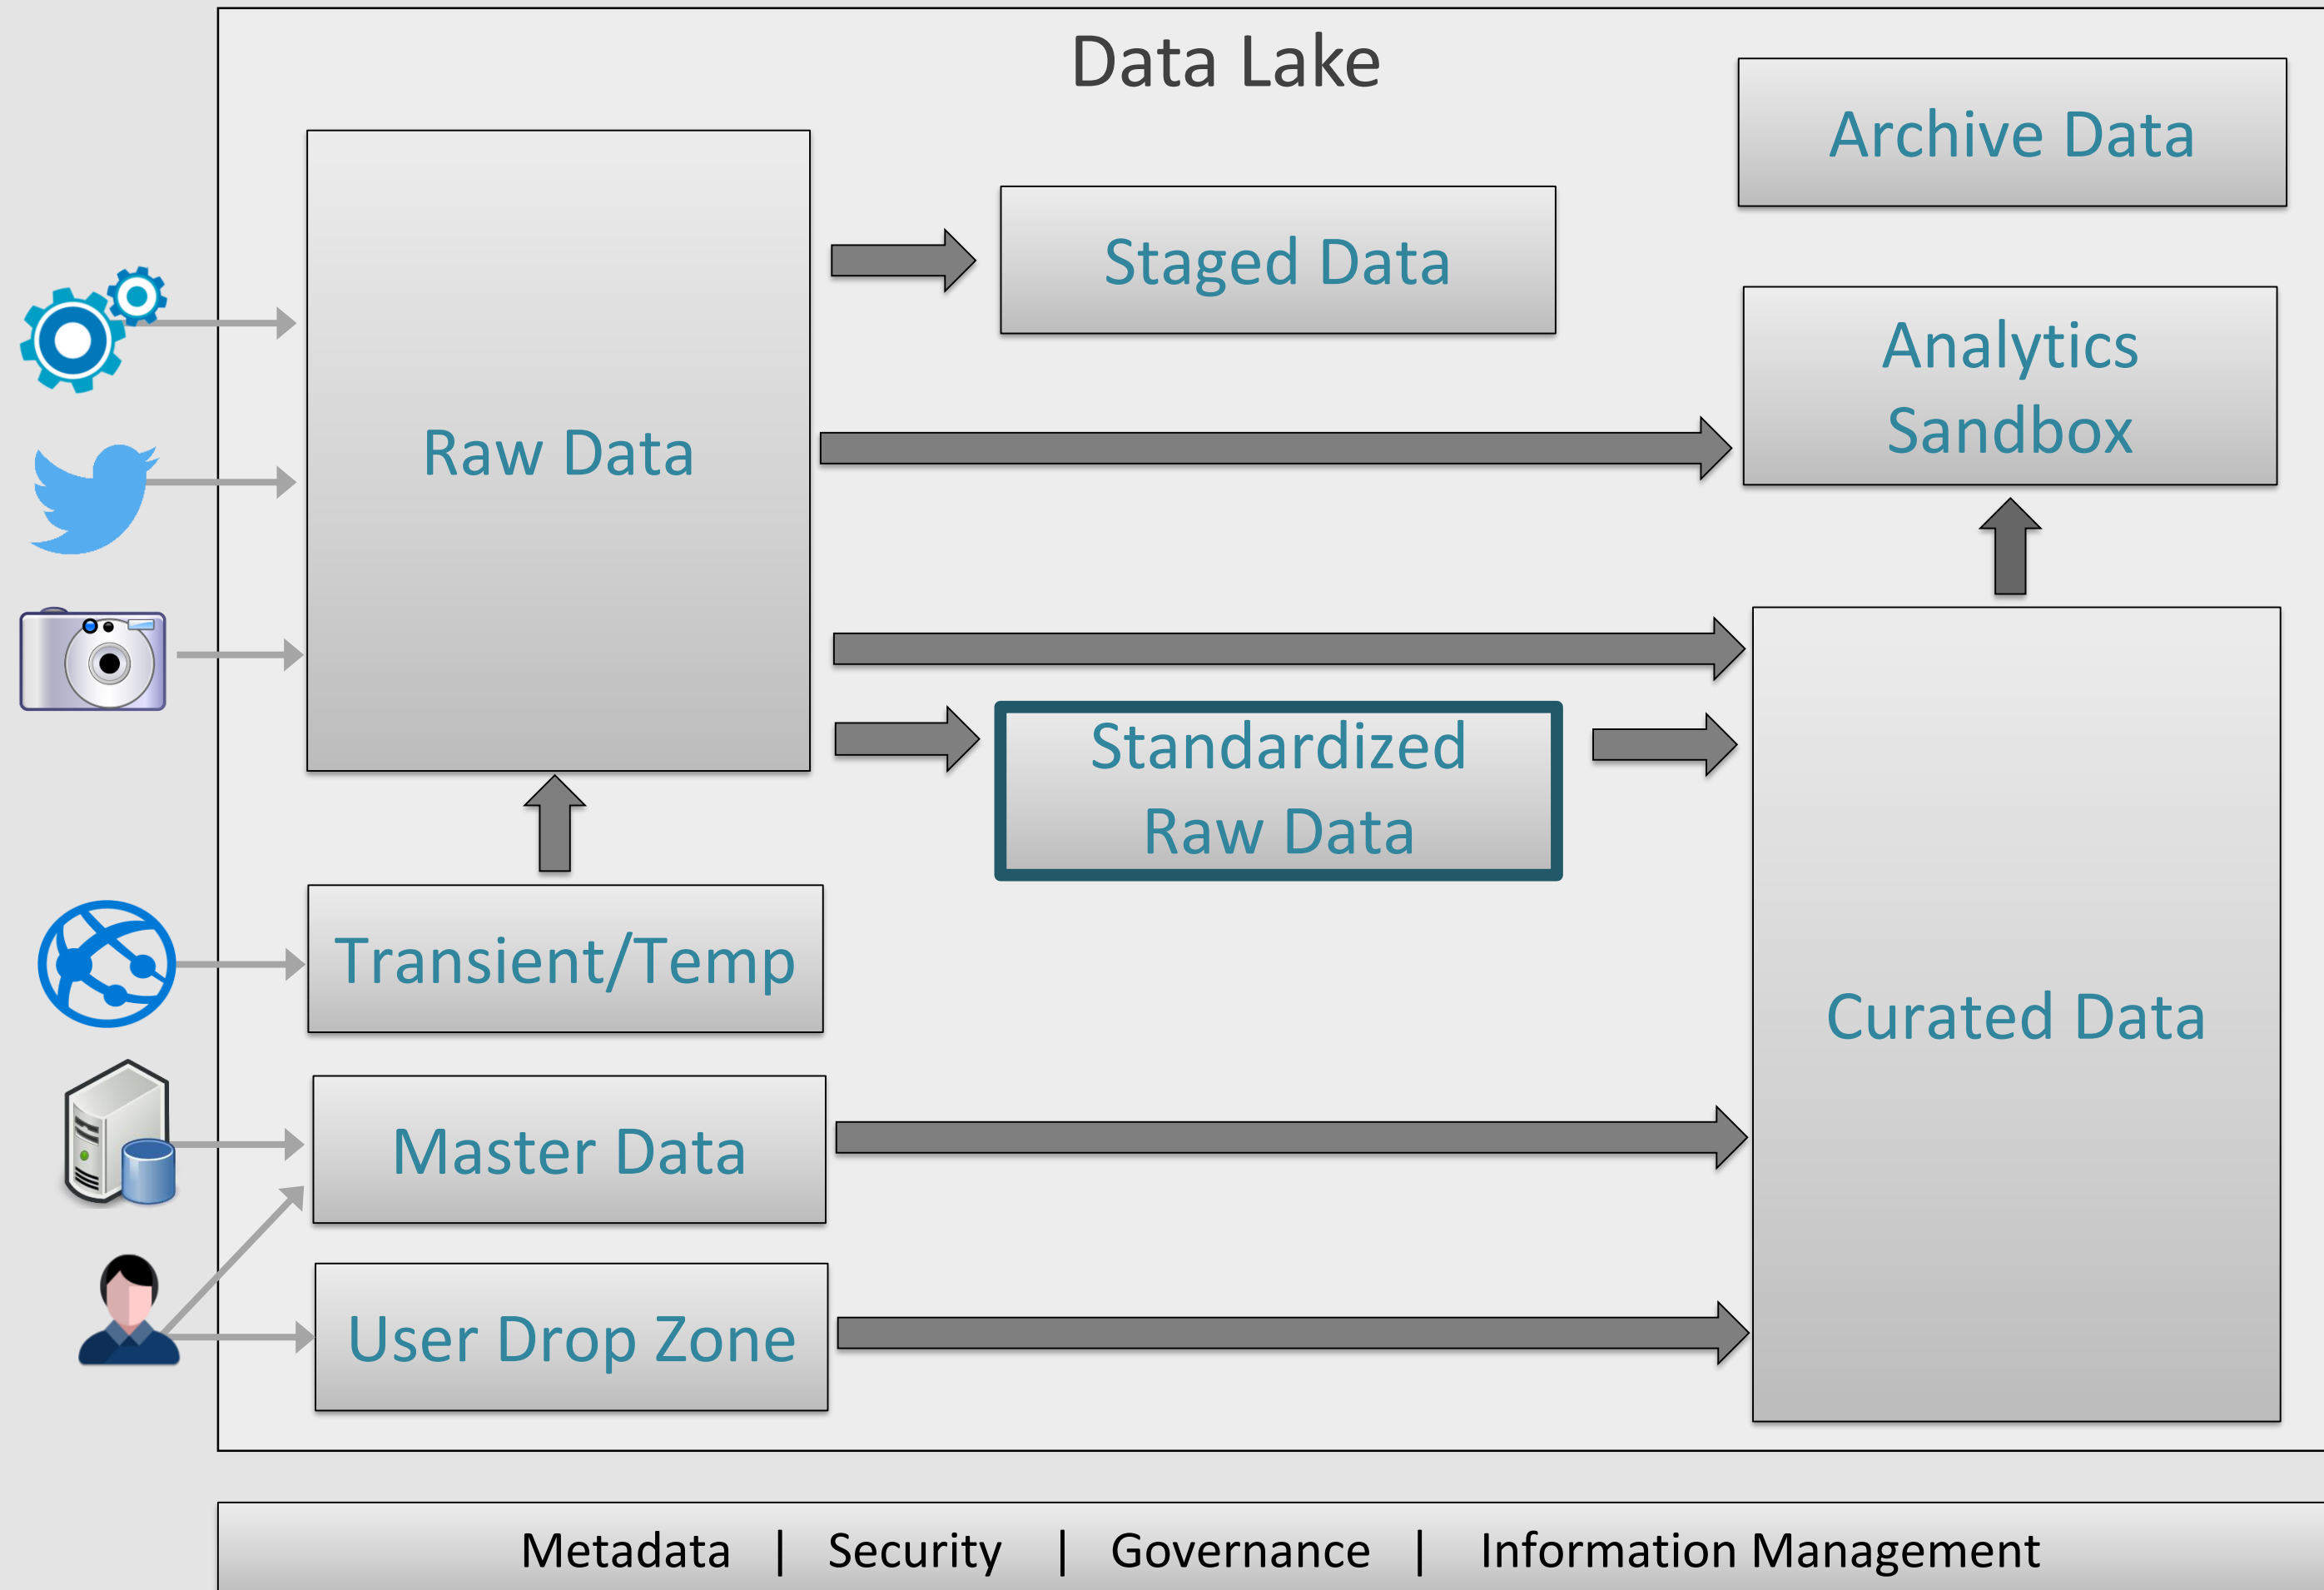
✓ Manually-generated data to augment analysis

Curated Data Zone



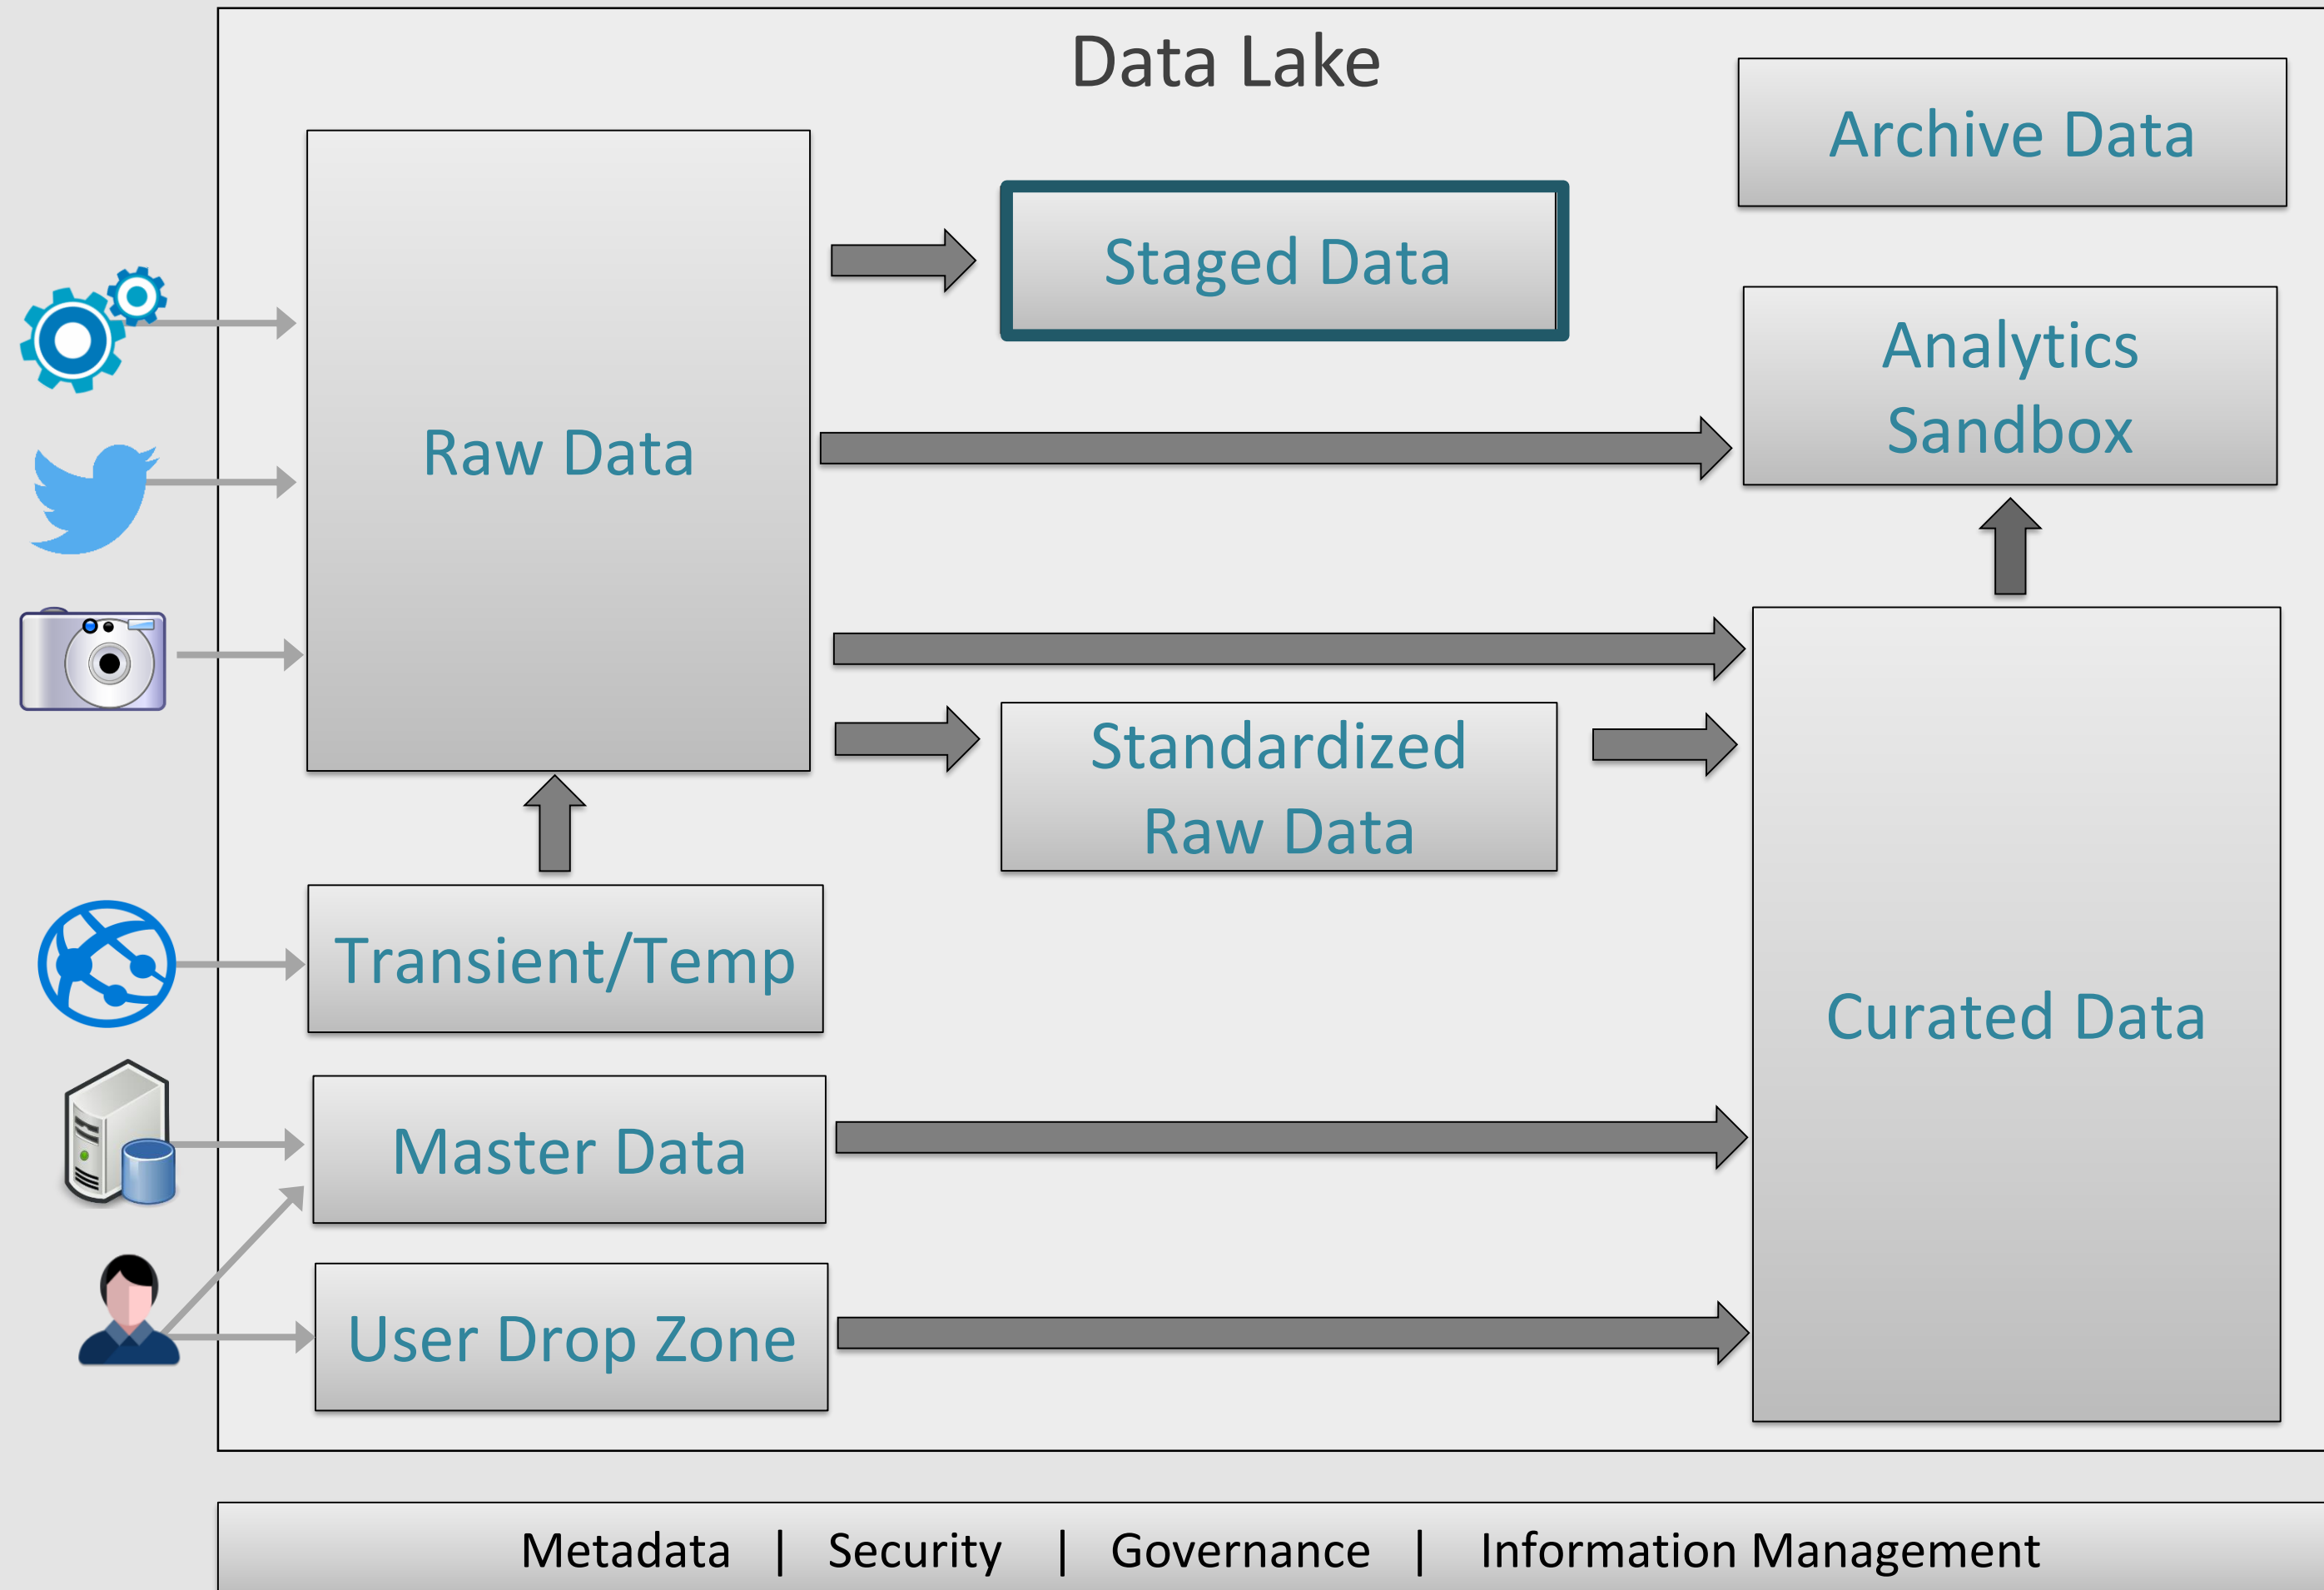
- ✓ Cleansed and transformed
- ✓ Organized for optimal data delivery (aka 'serving layer')
- ✓ Nearly all self-service data access comes from the Curated Data Zone
- ✓ Standard governance and security policies
- ✓ Standard change management principles

Standardized Data Zone



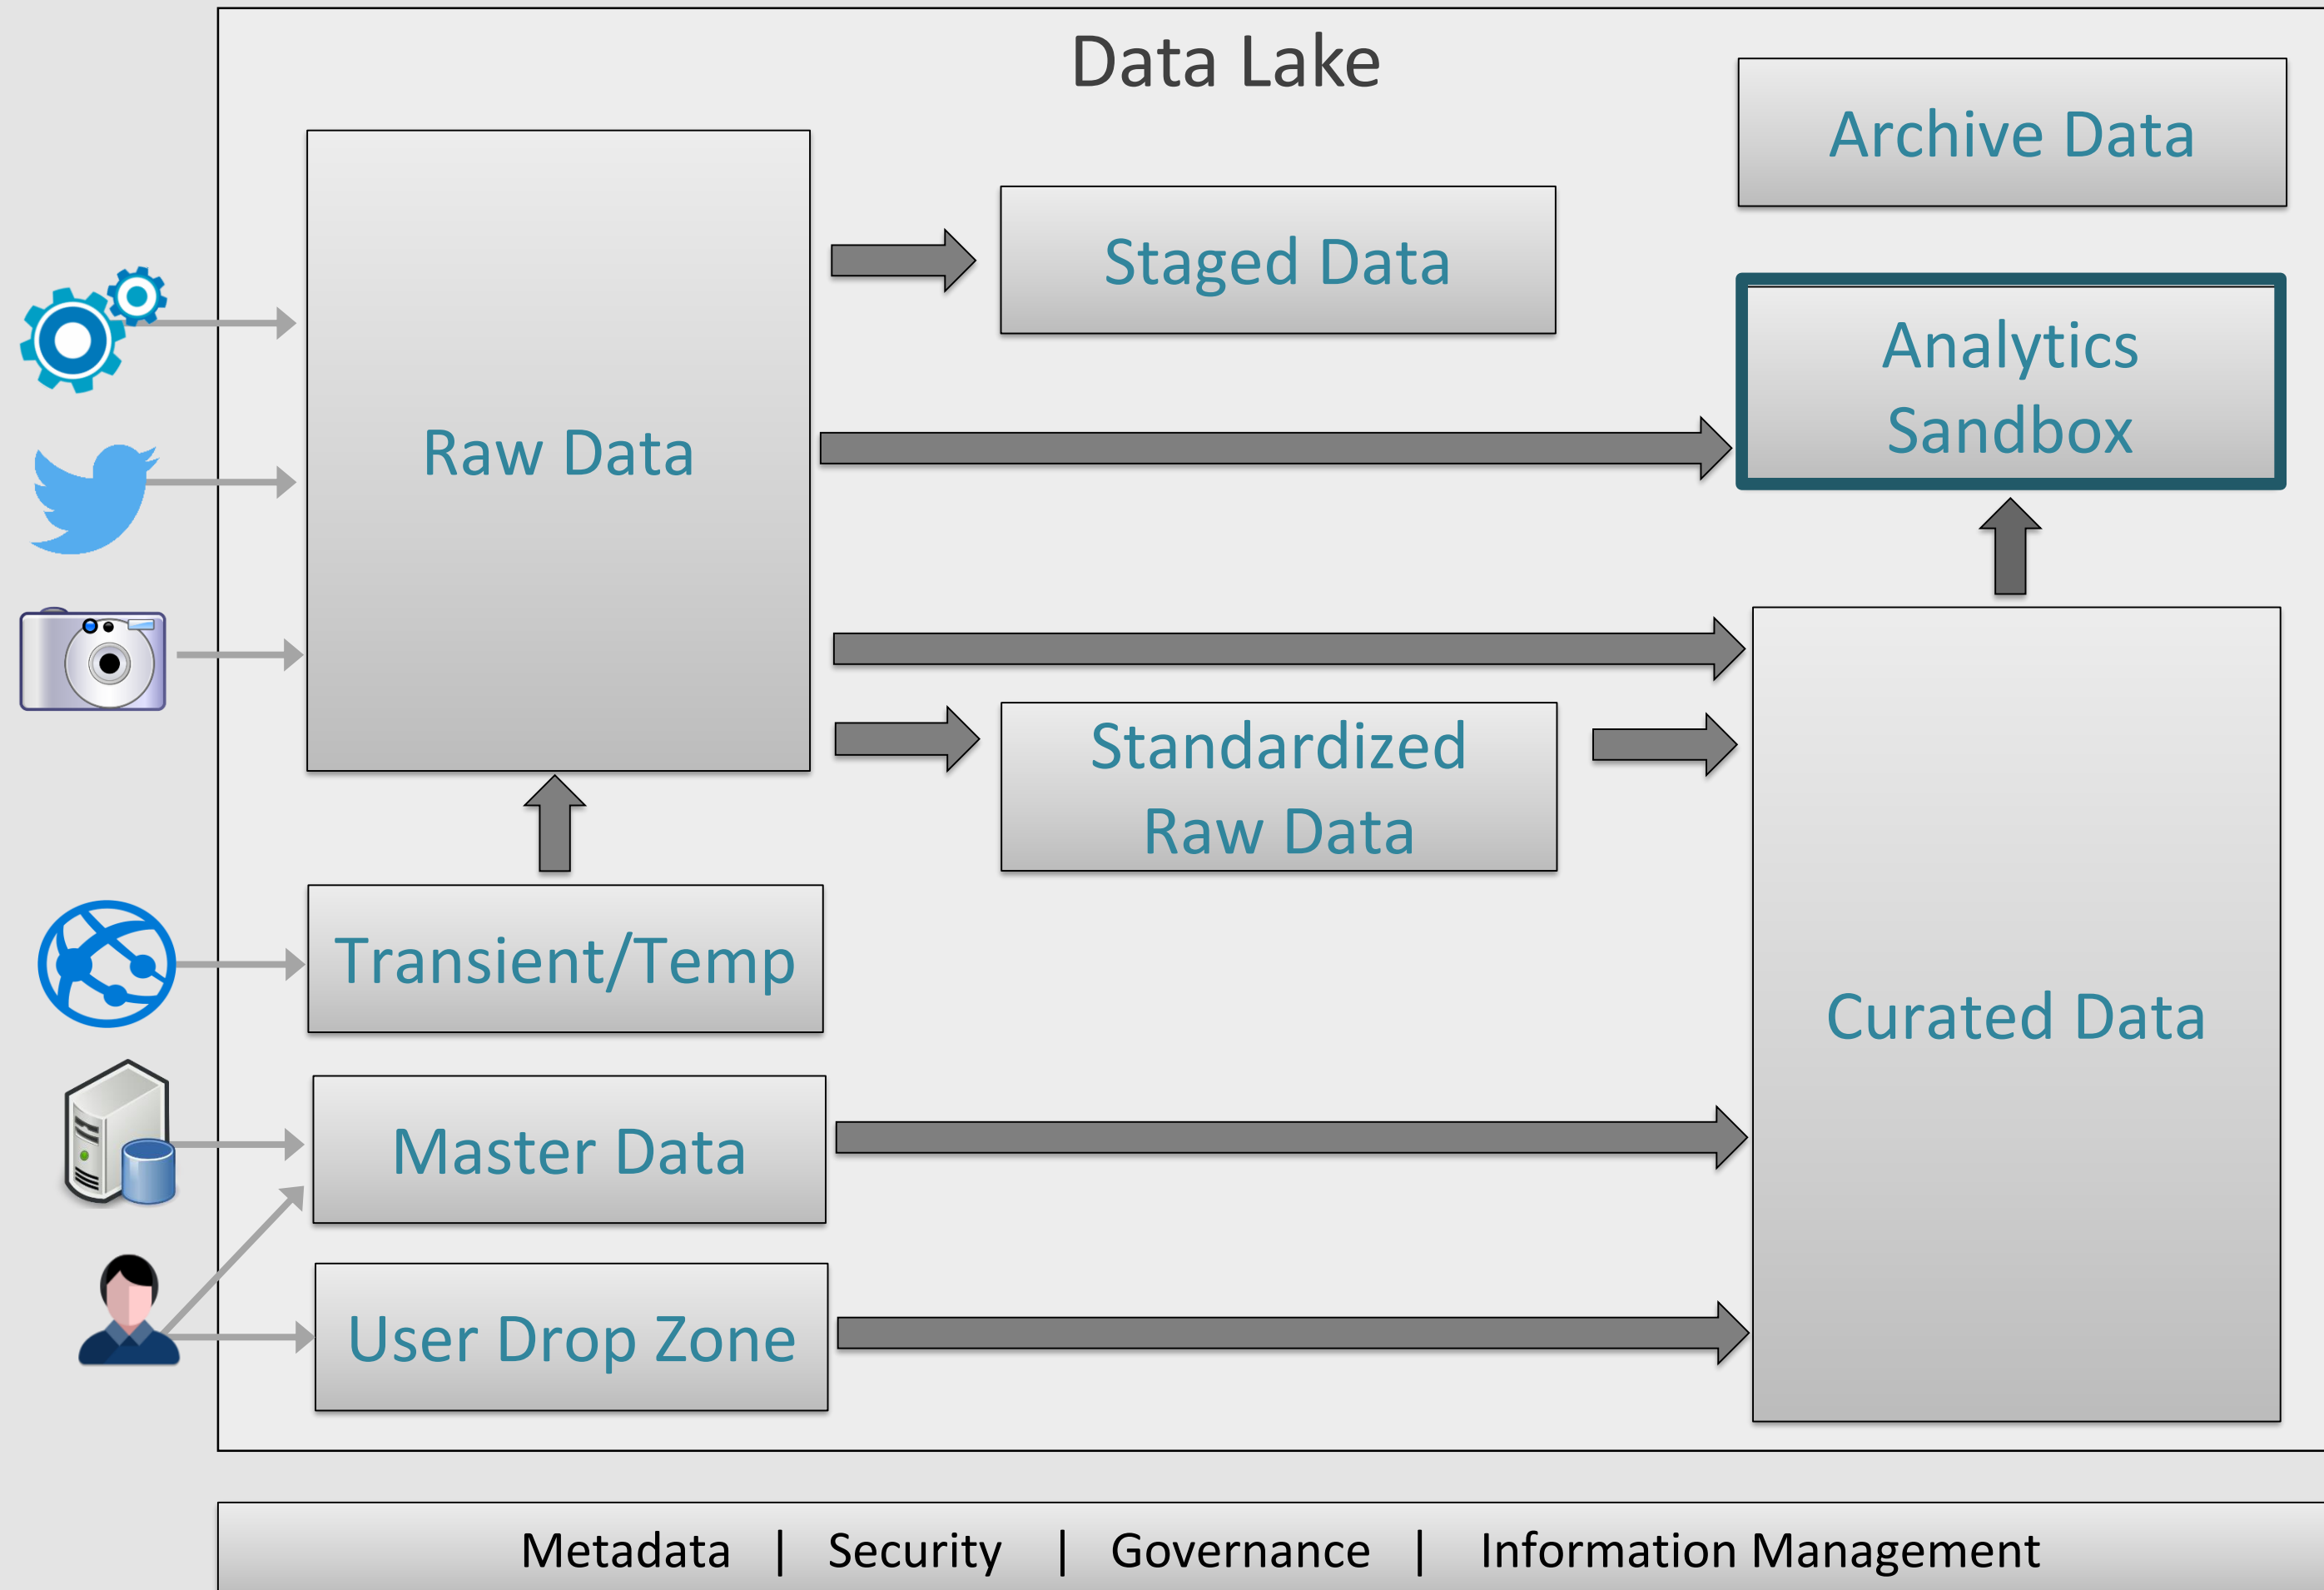
- ✓ A standardized version of the Raw Data Zone applicable to data structures which vary in format – ex: JSON which is standardized into consistent columns & rows (aka 'semantic normalization')
- ✓ No real cleansing or transformations applied
- ✓ Intermediary to assist creation of curated data
- ✓ File consolidations (ex: solve 'small files' performance issues)

Staged Data Zone



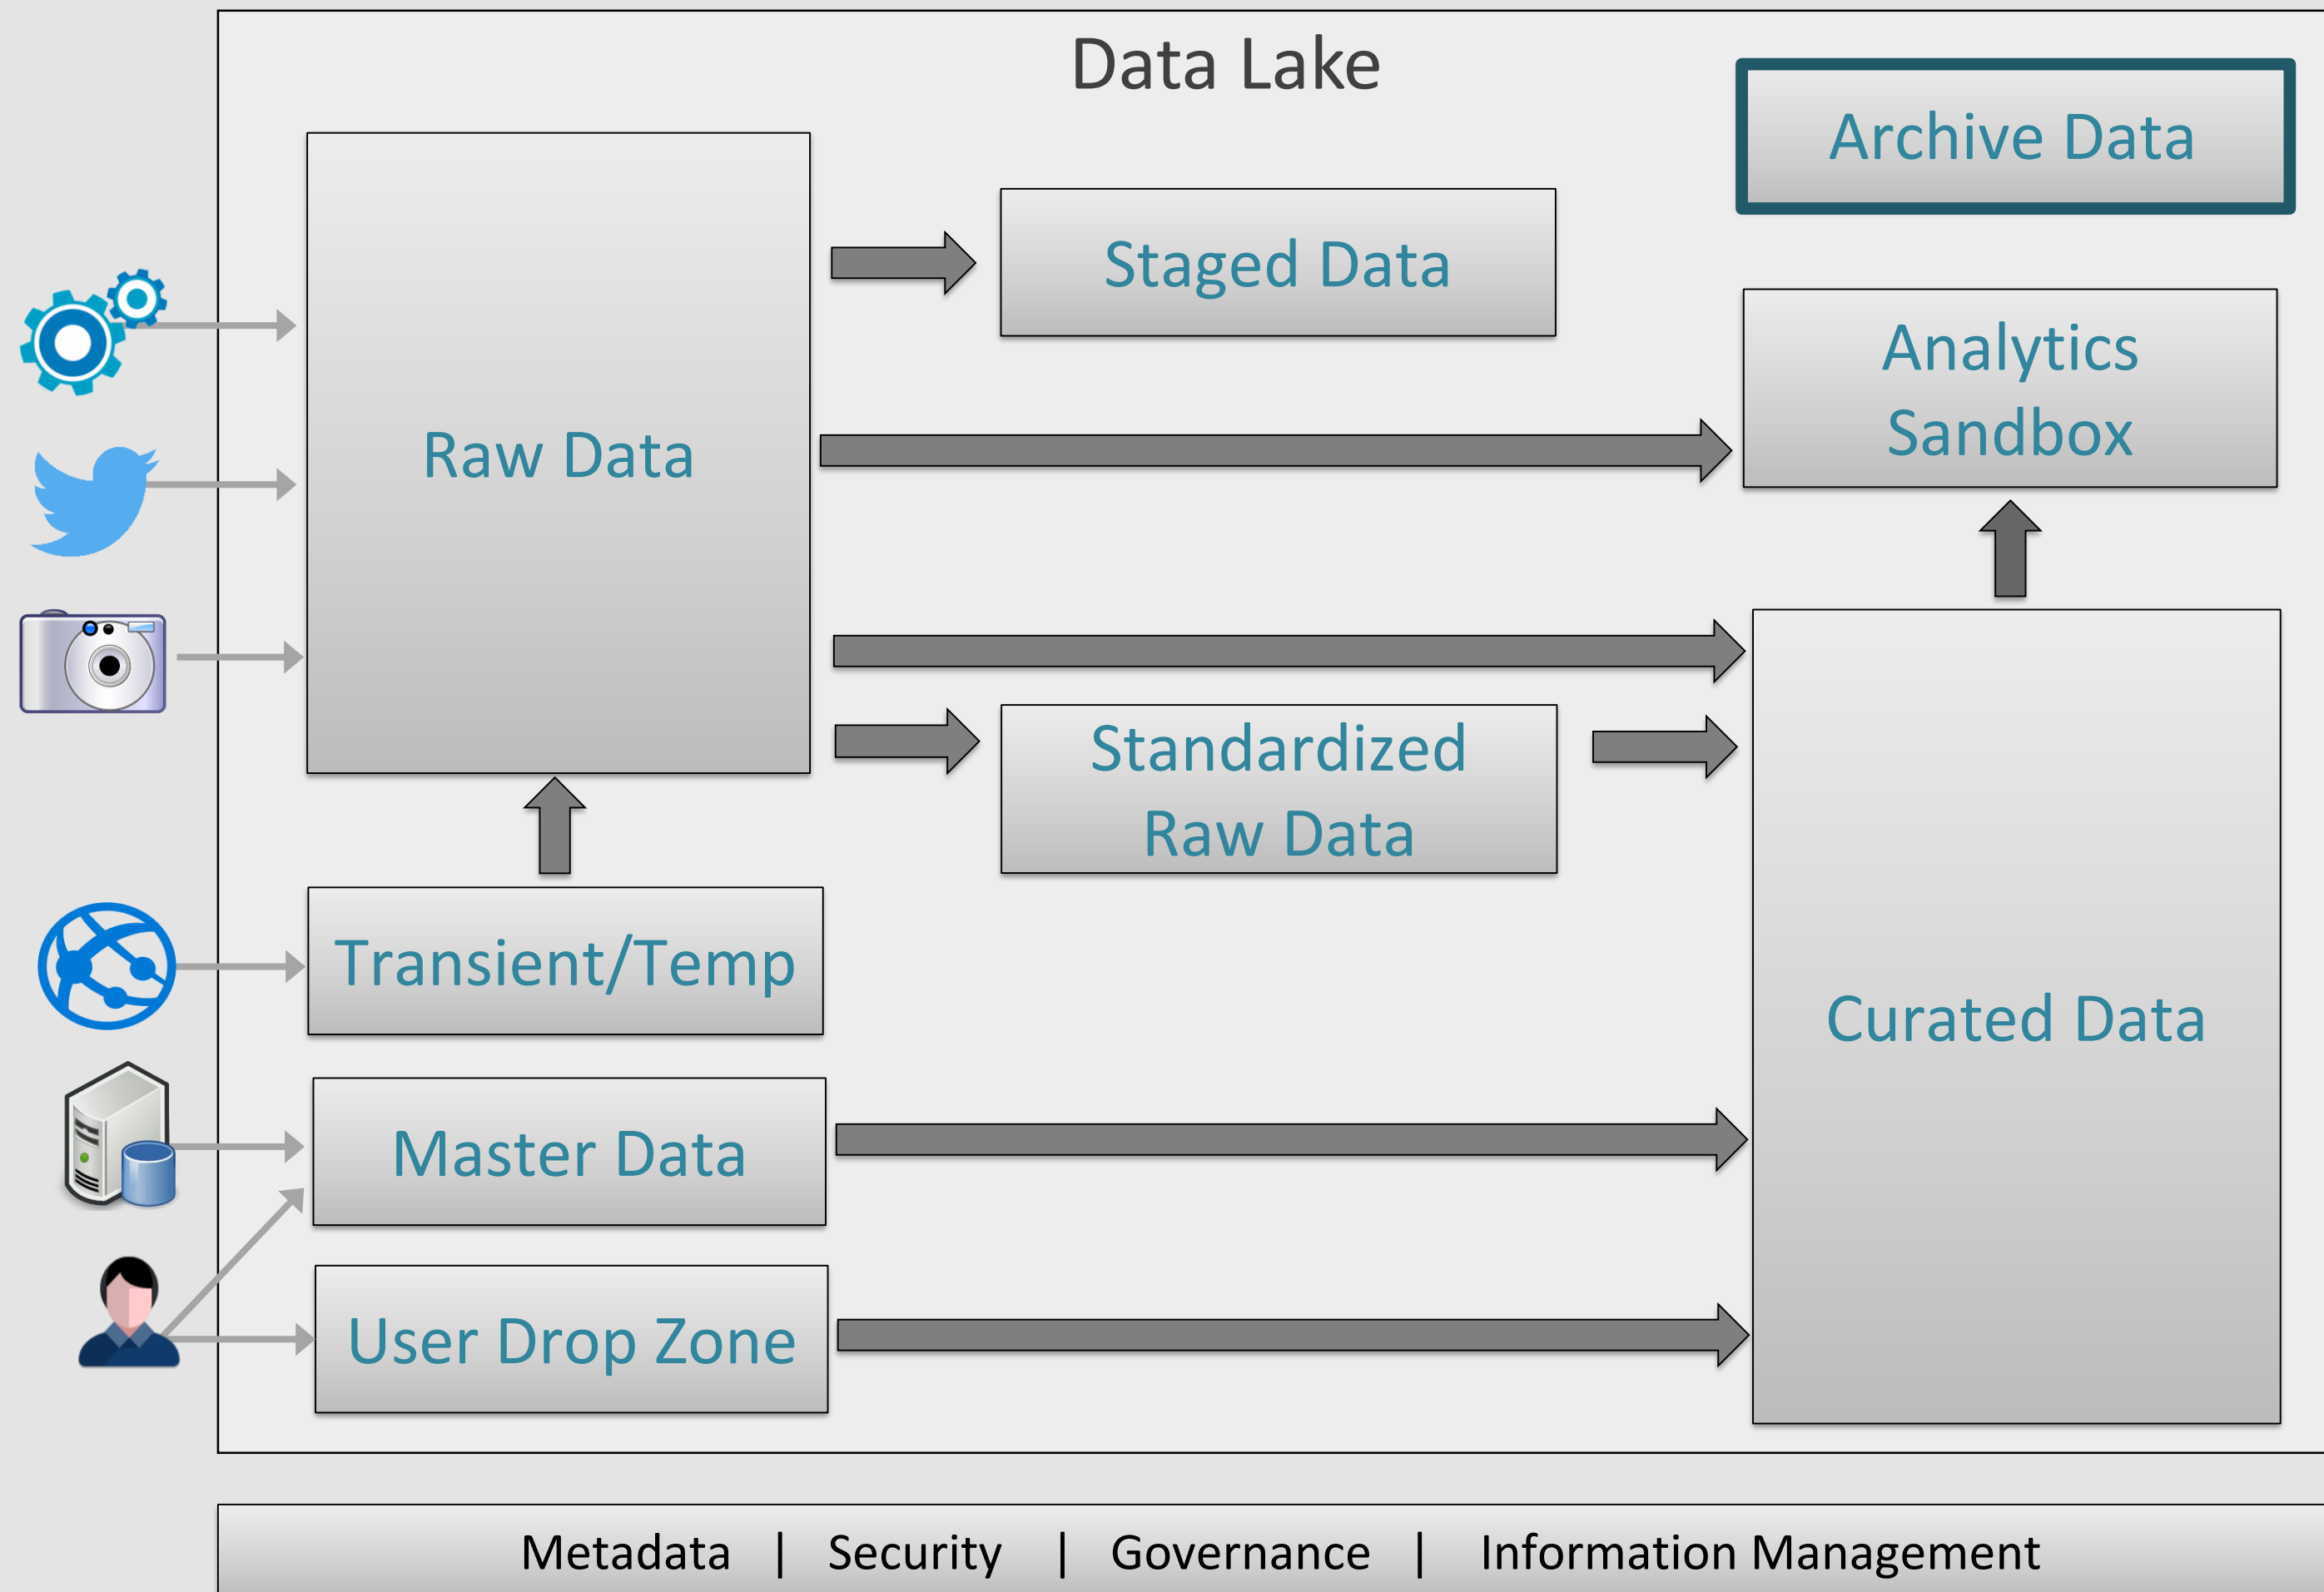
- ✓ Data which is staged for a particular purpose or application (thus has certain columns, certain formats, with or without headers, etc.)

Analytics Sandbox Zone




- ✓ Workspace for data science and exploratory activities
- ✓ Minimal, if any, governance and standards (purposely undisciplined)
- ✓ Valuable efforts are “productionized” and “operationalized” to the Curated Data Zone
- ✓ Not used for self-service, operationalized, purposes

Archive Data Zone



- ✓ An active archive
- ✓ Contains aged data offloaded from a data warehouse or other application
- ✓ Available for querying when needed (typically only occasionally)



What are some ways we
could potentially organize
data in a data lake?

Organizing a Data Lake

(1/7)

Objectives

- ✓ Plan the structure based on optimal data retrieval
- ✓ Avoid a chaotic, unorganized data swamp

Common ways to organize the data:

Time Partitioning

Year/Month/Day/Hour/Minute

Subject Area

Security Boundaries

Department
Business unit
etc...

Downstream App/Purpose

Data Retention Policy

Temporary data
Permanent data
Applicable period (ex: project lifetime)
etc...

Business Impact / Criticality

High (HBI)
Medium (MBI)
Low (LBI)
etc...

Owner / Steward / SME

Probability of Data Access

Recent/current data
Historical data
etc...

Confidential Classification

Public information
Internal use only
Supplier/partner confidential
Personally identifiable information (PII)
Sensitive – financial
Sensitive – intellectual property
etc...

Organizing a Data Lake

(2/7)

Raw Data Zone

Subject Area

Data Source

Object

Date Loaded

File(s)

Sales

Salesforce

CustomerContacts

2016

12

01

CustContact_2016_12_01.txt

Example 1

Pros: Subject area at top level, organization-wide
Partitioned by time

Cons: No obvious security or organizational boundaries

Curated Data Zone

Purpose

Type

Snapshot Date

File(s)

Sales Trending Analysis

Summarized

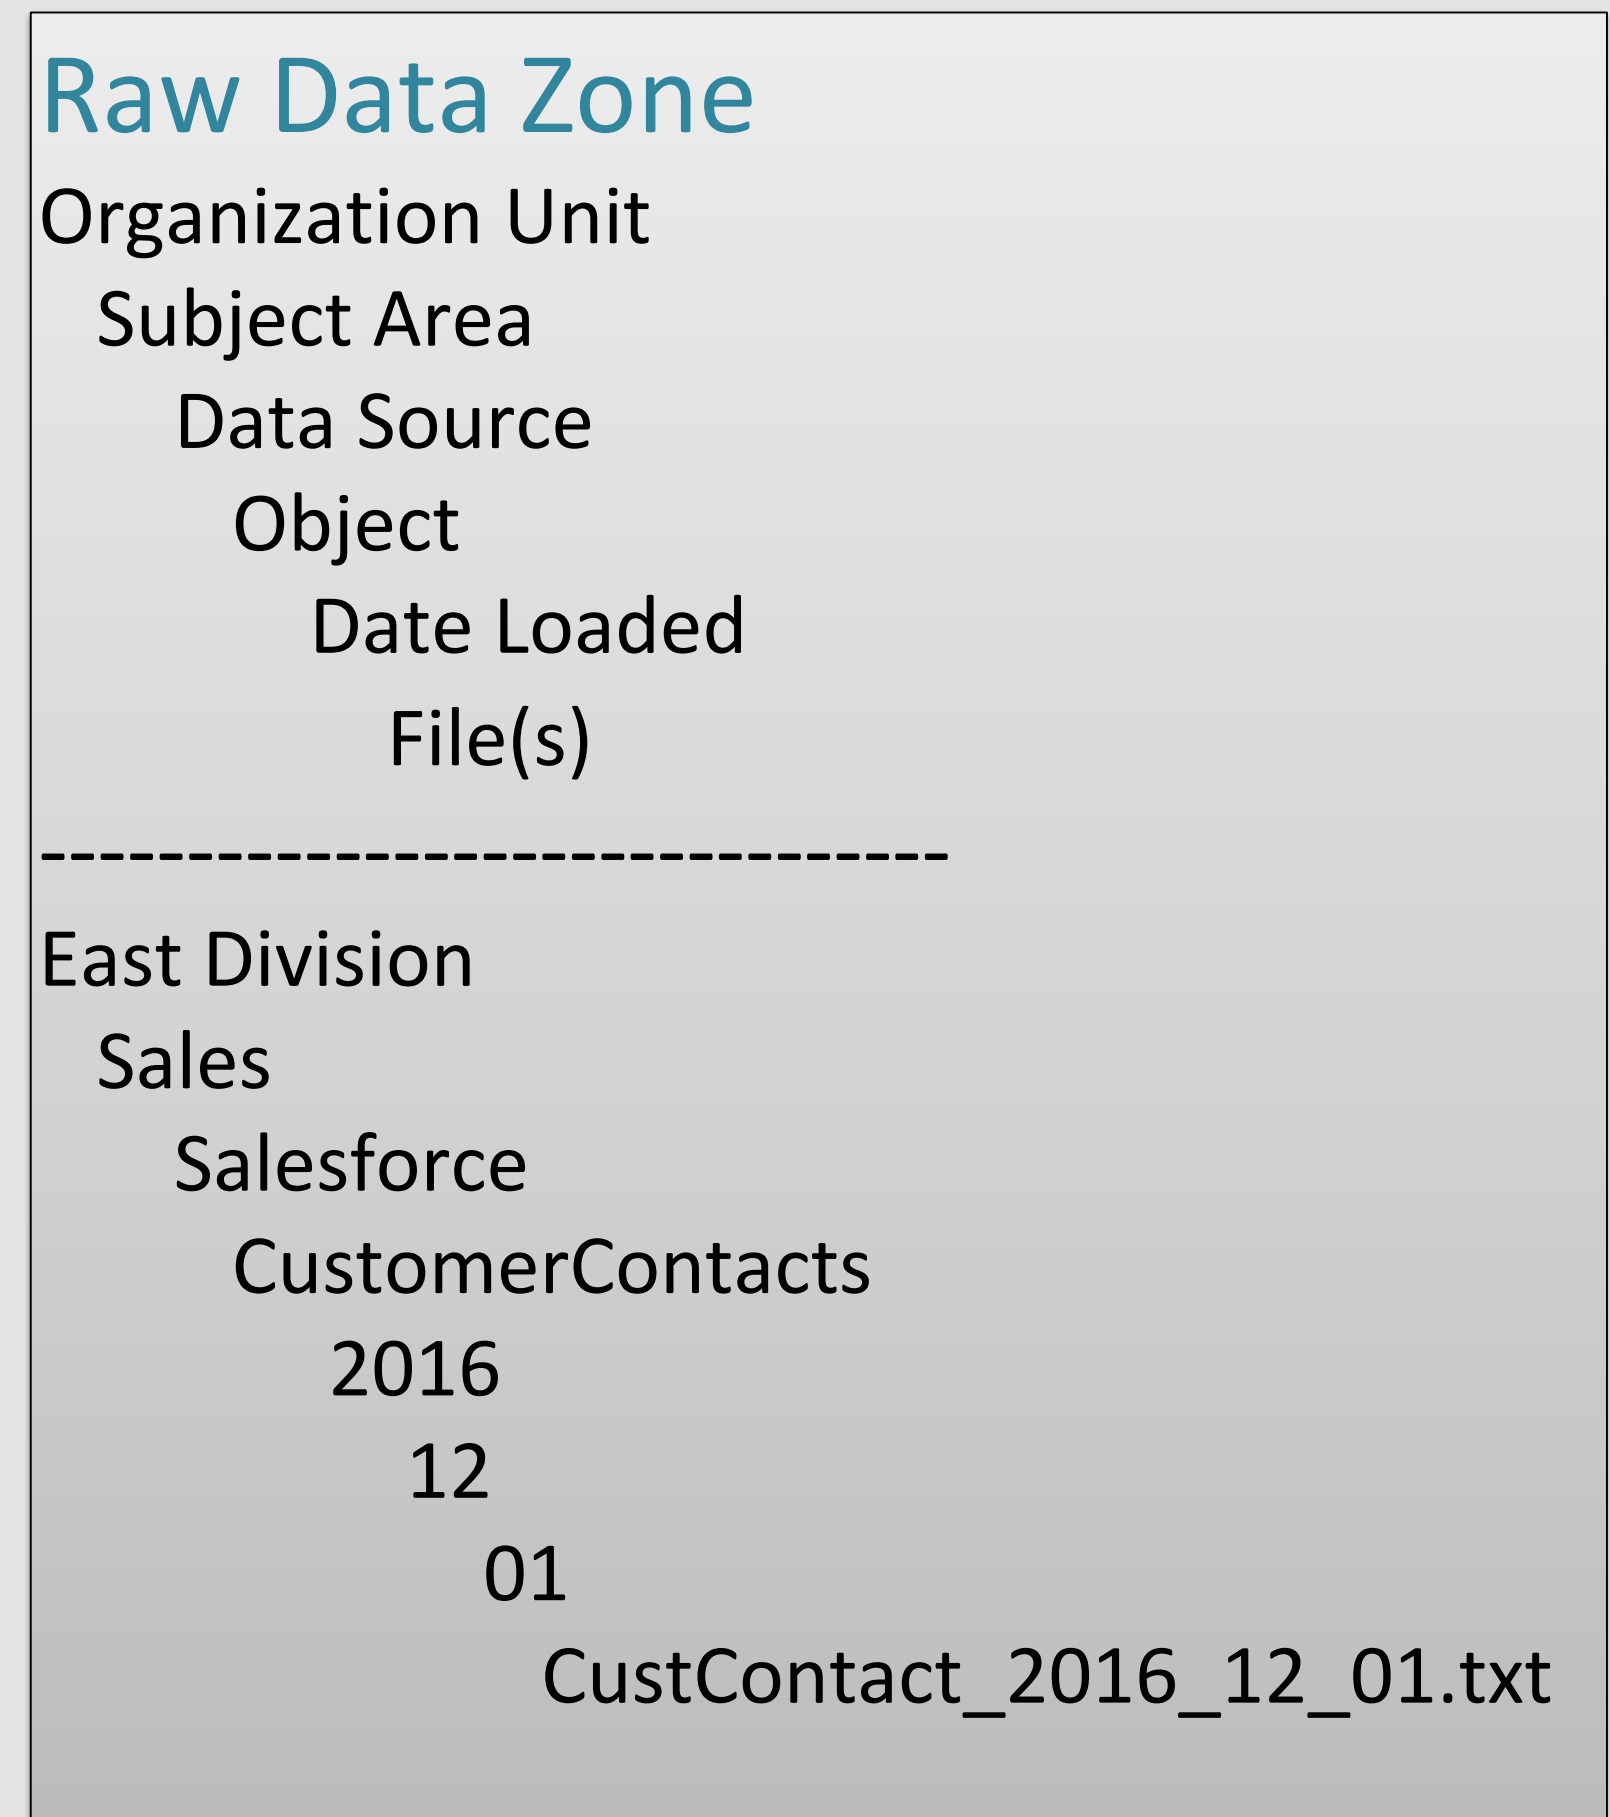
2016_12_01

SalesTrend_2016_12_01.txt

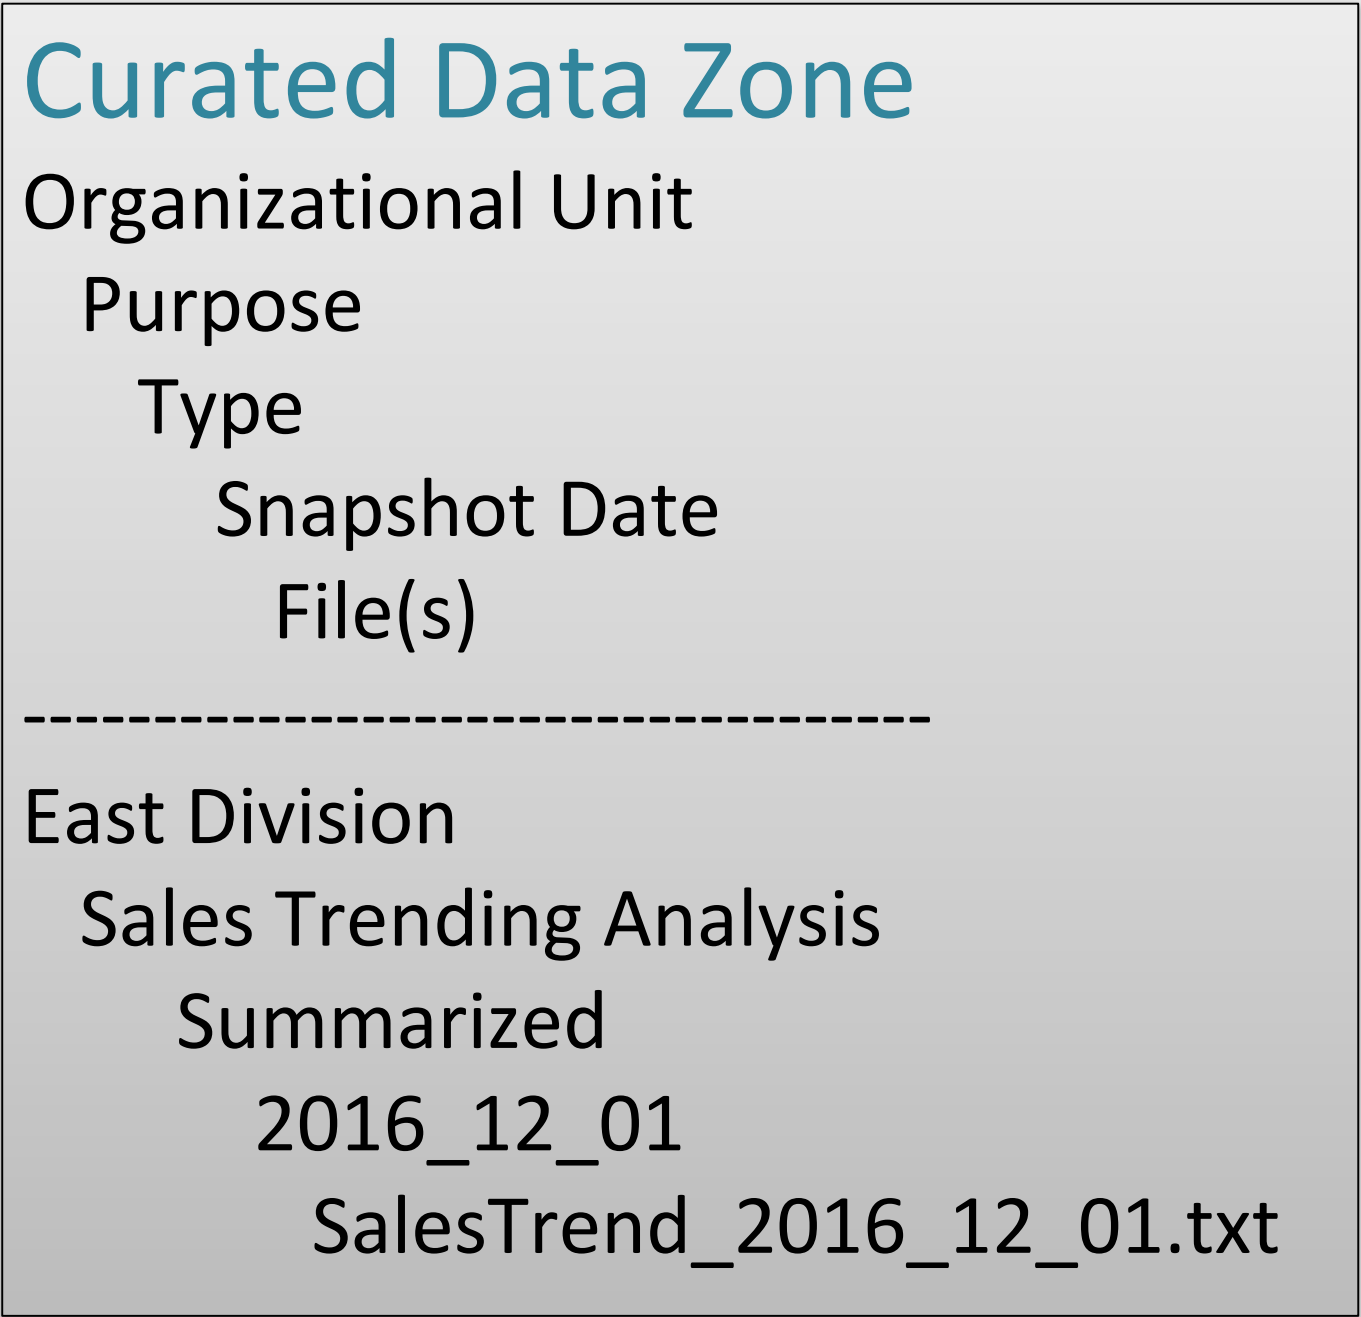


Organizing a Data Lake

(3/7)



Example 2
Pros: Security at the organizational level
Partitioned by time
Cons: Potentially siloed data, duplicated data



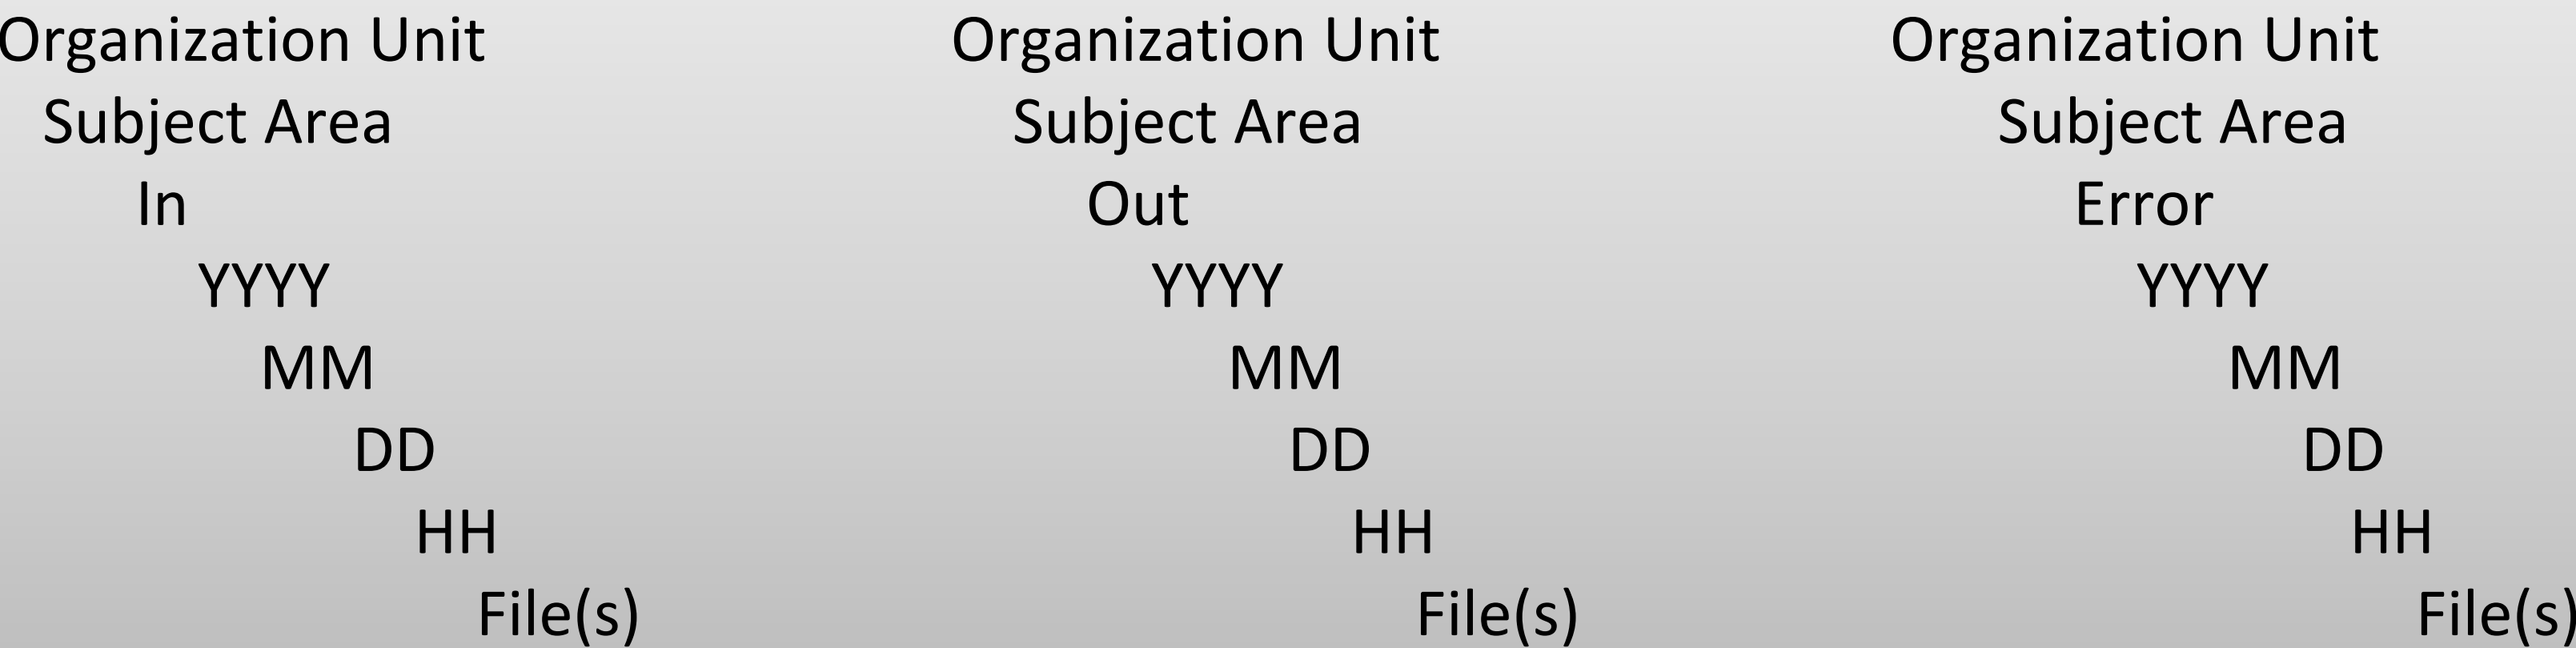
Organizing a Data Lake

(4/7)

Example 3

- Pros:** Segregates records coming in, going out, as well as error records
Time partitioning can go down to the hour, or even minute level, depending on volume (ex: IoT data)
- Cons:** Not obvious by the names what the purpose of 'out' is (which could be ok if numerous downstream applications utilize the same 'out' data)

Raw Data Zone



Organizing a Data Lake

(5/7)

Subject Area 1

RawData

YYYY

MM

CuratedData

MasterData

StagedData

Subject Area 2

RawData

YYYY

MM

CuratedData

MasterData

StagedData

Example 4

Zones are a logical need, but they don't necessarily have to be at the top of the structure

Pros: Security by subject area

Cons: All raw data is not centralized

Organizing a Data Lake

(6/7)

Do:

- ✓ Hyper-focus on **ease of data discovery** & retrieval – will one type of structure make more sense?
- ✓ Focus on **security implications** early – what data redundancy is allowed in exchange for security
- ✓ Include **data lineage** & relevant **metadata** with the data file itself whenever possible (ex: columns indicating source system where the data originated, source date, processed date, etc)
- ✓ Include the **time element** in **both** the folder structure & the file name
- ✓ Be **liberal yet disciplined** with folder structure (lots of nests are ok)
- ✓ Clearly separate out the zones so **governance & policies** can be applied separately
- ✓ Register the curated data with a catalog (ex: **Azure Data Catalog**) to document the metadata—a data catalog is even more important with a data lake
- ✓ Implement **change management** for migrating from a sandbox zone (discourage production use from the sandbox)
- ✓ Assign a **data owner** & **data archival** policies as part of the structure, or part of the metadata

Organizing a Data Lake

(7/7)

Don't:

- × **Do not combine mixed formats** in a single folder structure
 - ✓ If it's looping through all files in a folder schema-on-read will fail if it finds a different format
 - ✓ Files in one folder should all be able to be traversed with the same script

- × **Do not put your date partitions at the beginning of the file path** -- it's much easier to organize & secure by subject area/department/etc if dates are the lowest folder level

Optimal for top level security:

\SubjectArea\YYYY\MM\DD\FileData_YYYY_MM_DD.txt

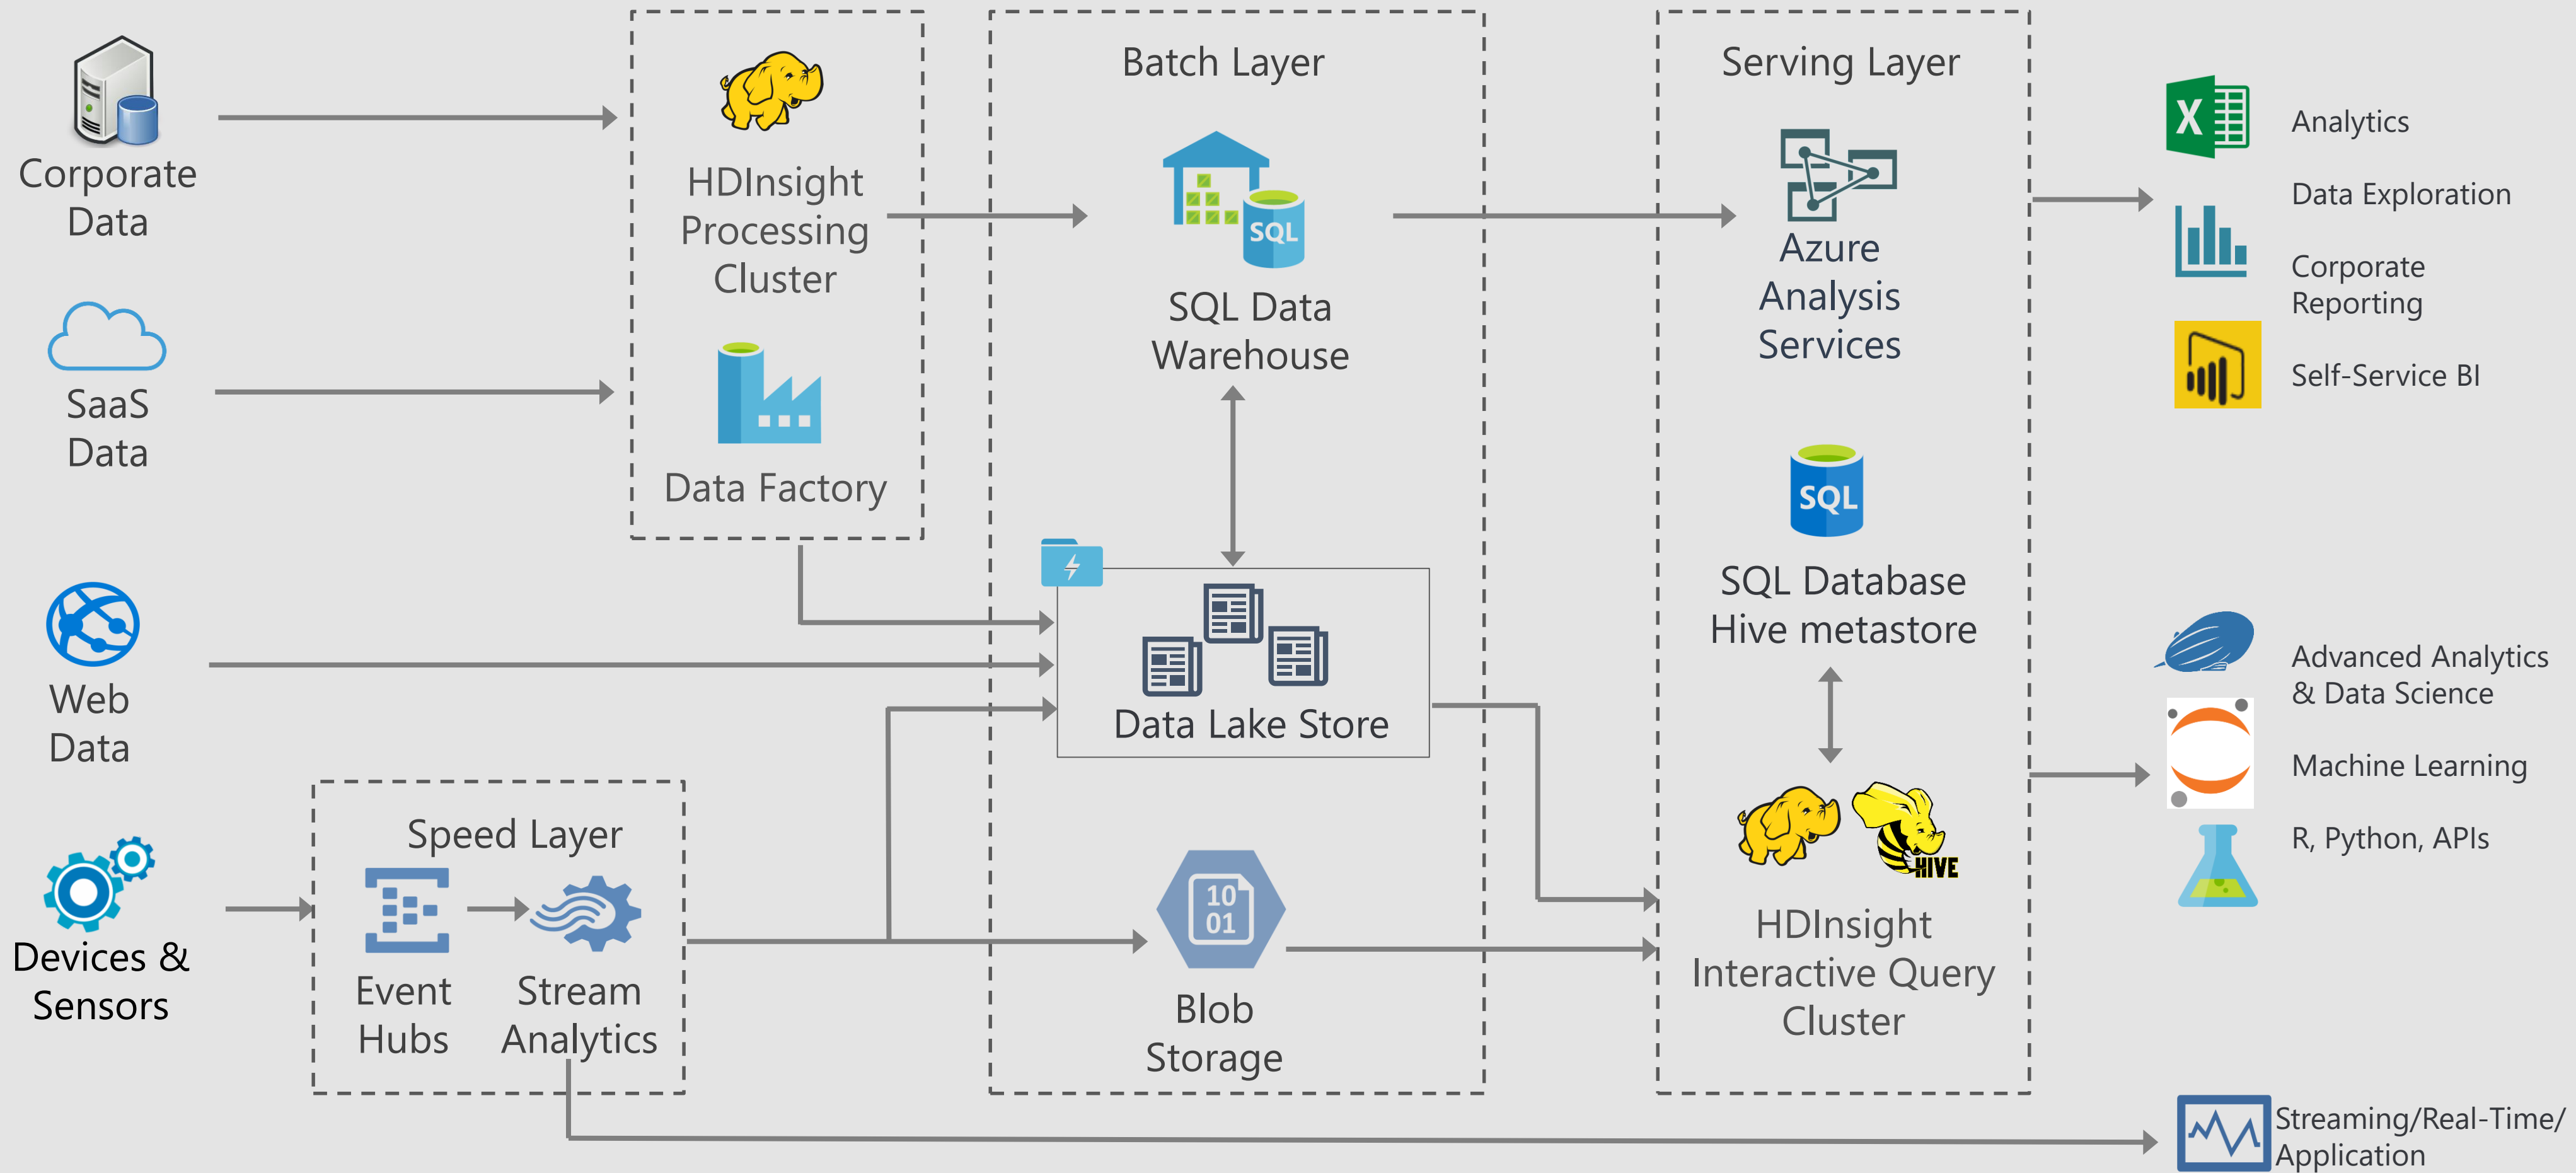
Tedious for enforcing security:

\YYYY\MM\DD\SubjectArea\FileData_YYYY_MM_DD.txt

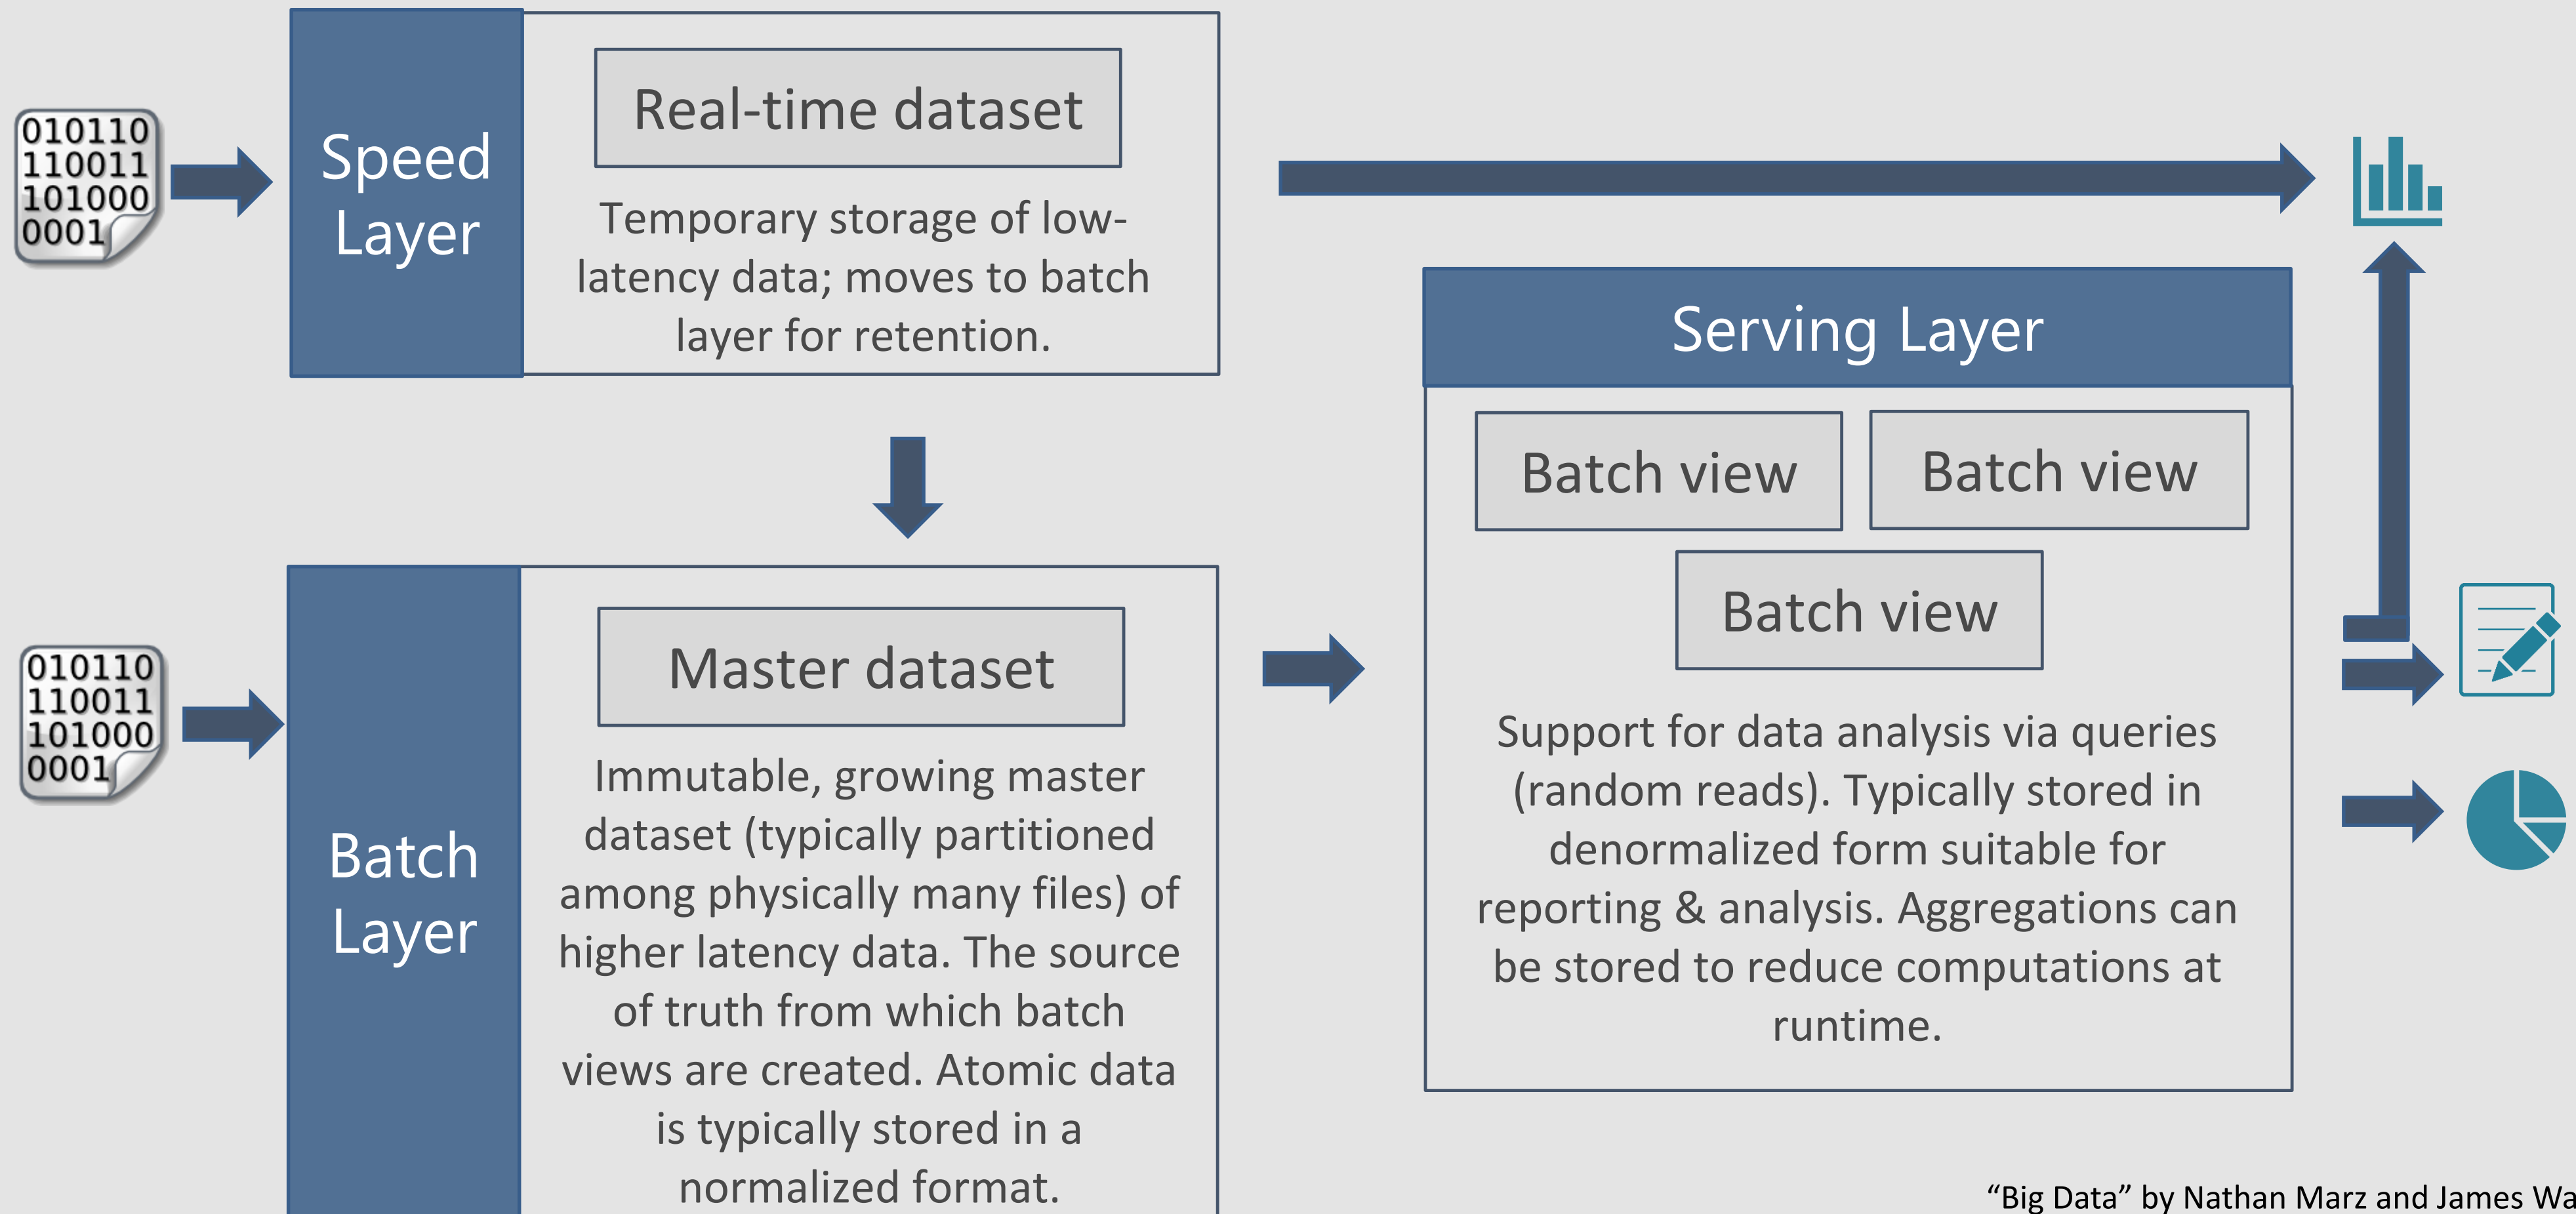
- × **Do not neglect naming conventions.** You might use camel case, or you might just go with all lower case – either is ok, as long as you're consistent because some languages are **case-sensitive**


Following Big Data Principles When Designing A Data Lake

Lambda Architecture



Lambda Architecture





What principles might you
expect to follow
in a big data project?

Big Data Principles to Follow in a Data Lake Project

Immutable Raw Data

- Raw data is append-only & unchanging
- Continually growing
- No summarizations or deletions
- Bad data can be deleted, but it's rare
- Immutable data is resilient to human error

Recreatable

- Everything downstream from the raw data can be regenerated (error tolerant)
- Schema changes can be handled
- Unstructured data can always be re-structured (“semantic normalization”)
- Speed layer may use approximations, corrected in the batch layer (eventual consistency)

Identifiable Data

- Timestamped
- Unique (tolerant of duplicates from retries)

Rawness of Data

- Obtain the rawest, most atomic, data available

Separate Layers

- Redundant data in both the batch & serving layers allows normalized & denormalized data

Schema Changes Over Time

(1/2)

Schema changes include:

Addition of new columns

Removal of columns

Renaming of columns

Two options:

(1) Schema enforcement upon the ingestion of data

(2) Schema flexibility for the developers; deal with “standardizing” the data after ingestion

Schema Changes Over Time (1/2)

Raw Data:

```
{
  "ID": "a4791906-a9a7-4d47-956b-9df12b5f7296",
  "SID": "a93c45cb-9420-4efc-919e-ca5338660bbc",
  "TID": null,
  "CID": "98FD9198EA4A47C41AB8CDC31DD53077",
  "LID": "87F4BE0AEB0CD4CF56C280F7660F0281",
  "PID": "S1",
  "PV": "11.2.0.0",
  "TS": "2017-09-18T00:01:17.7781017Z",
  "EC": "47623b79-f290-45fe-a23a-2e840ff5f63f",
  "ED": {
    {
      "Name": "UsageStopped",
      "Path": "WindowsPerformanceConsole",
      "D": "00:02:01.0644843"
    },
    "EventProcessedUtcTime": "2017-09-18T00:01:17.7781017Z",
    "PartitionId": 1,
    "EventEnqueuedUtcTime": "2017-09-18T00:01:17.7781017Z"
  }
}
```

```
{
  "ID": "28439047-b4a7-4aa6-9a4b-ebcd9cc1e028",
  "SID": "69f947db-52de-4f19-8d4a-3b16ed11c27d",
  "TID": null,
  "CID": "BD5145B8D4CDBFEC72A23793F309ECF",
  "LID": "577C0A72744D3A13F9A40663C3D34CC0",
  "PID": "S1",
  "PV": "11.2.0.0",
  "TS": "2017-09-18T14:14:38.4585269Z",
  "EC": "c0a0d257-752a-4899-a0eb-5680a323bd19",
  "ED": {
    {
      "Val": 74.0,
      "PrevVal": 0.0,
      "Name": "PropertyValueChanged",
      "Path": "EH\\ScoreChange"
    },
    "EventProcessedUtcTime": "2017-09-18T14:14:38.4585269Z",
    "PartitionId": 1,
    "EventEnqueuedUtcTime": "2017-09-18T14:14:38.4585269Z"
  }
}
```

```
{
  "ID": "fa27bf3f-b101-4c92-87e3-dc63f20614b2",
  "SID": "a73215bb-b831-43f2-a3cd-fed140d2c1eb",
  "TID": null,
  "CID": "71021EFAAF18DF534D0F1E986E687B04",
  "LID": "DAA3C04DF132B74027475ADE1D1788D8",
  "PID": "S1",
  "PV": "11.2.0.0",
  "TS": "2017-09-18T03:29:49.7600422Z",
  "EC": "e95861bb-cb49-4579-804c-6d0fafbed33c",
  "ED": {
    {
      "Name": "Navigation",
      "Path": "EM"
    },
    "EventProcessedUtcTime": "2017-09-18T03:29:49.7600422Z",
    "PartitionId": 1,
    "EventEnqueuedUtcTime": "2017-09-18T03:29:49.7600422Z"
  }
}
```



Standardized

Raw Data:

InstanceID	SessionID	TransactionID	ClientID	LocationID	ProductID	ProductVersion
a4791906-a9a7-4d47-956b-9df12b5f7296	a93c45cb-9420-4efc-919e-ca5338660bbc	0	98FD9198EA4A47C41AB8CDC31DD53077	87F4BE0AEB0CD4CF56C280F7660F0281	S1	11.2.0.0
fa27bf3f-b101-4c92-87e3-dc63f20614b2	a73215bb-b831-43f2-a3cd-fed140d2c1eb	0	71021EFAAF18DF534D0F1E986E687B04	DAA3C04DF132B74027475ADE1D1788D8	S1	11.2.0.0
28439047-b4a7-4aa6-9a4b-ebcd9cc1e028	69f947db-52de-4f19-8d4a-3b16ed11c27d	0	BD5145B8D4CDBFEC72A23793F309ECF	577C0A72744D3A13F9A40663C3D34CC0	S1	11.2.0.0
daf30e65-c7ab-4bed-b651-4e122e63cebf	d5fa951e-5ea9-4230-b2ef-9b48ef0e622f	0	BD5145B8D4CDBFEC72A23793F309ECF	577C0A72744D3A13F9A40663C3D34CC0	S1	11.2.0.0
Row continues						

TimestampUtc	EventClassID	EventClassName	EventPath	EventDuration	EventValue	EventPreviousValue
2017-09-18T00:01:17.7781017Z	47623b79-f290-45fe-a23a-2e840ff5f63f	UsageStopped	WindowsPerformanceConsole	02:01.1	0	0
2017-09-18T03:29:49.7600422Z	e95861bb-cb49-4579-804c-6d0fafbed33c	Navigation	EM	00:00.0	0	0
2017-09-18T14:14:38.4585269Z	c0a0d257-752a-4899-a0eb-5680a323bd19	PropertyValueChanged	EH\\ScoreChange	00:00.0	74	0
2017-09-18T14:16:29.8271102Z	4989e98a-fedb-4156-9eba-ee7f6961a48a	PropertyValue	MonitoringService\\Count	00:00.0	1	0

(semantic normalization)

Data Formats & Data Compression

CSV

Commonly used. Human-readable. Not compressed. Typically not the best choice for large datasets.

JSON

Commonly used. Human-readable. Self-describing schema.

Parquet

Columnar format; highly compressed.

Avro

Row-based format. Supports compression. Schema encoded on the file.

ORC (optimized row columnar)

Columnar format with collections of rows. Light indexing and statistics.

Deciding on a format

- Supported formats by key systems
- Integration with other systems
- File sizes
- Schema changes over time
- If a self-describing schema is desired
- Data type support
- Data format compatibility
- Performance of workload (read vs. write)
- Convenience & ease of use

Techniques to Recompute the Serving Layer

Full recomputation

The entire master dataset is used to recompute the batch views in the serving layer.

Pros: Simplicity

- Better human fault-tolerance

- Ability to continually reap benefits of improved algorithms or calculations

- Easier to keep wide datasets which contain redundant data synchronized/consistent

Cons: Performance; speed of updates

- CPU and I/O heavy

- Not practical for extremely large datasets


Incremental recomputation

Only new data from the master dataset is involved in recomputations.

Pros: Better performance

Cons: Significantly more complex

- Still need a way to do a full recomputation in the event of errors or significant changes



What is the state of data modeling
for files stored in
a data lake?

Data Modeling for Files in a Data Lake

Wide datasets, with all data needed in one file, are commonly used

Pros: Easy to do analysis.

- Data can be co-located on the nodes as the data gets distributed (depending on the tool).

- The desired format frequently for data scientists & the tools they use.

- Usually well-suited to in-memory, columnar, data formats.

Cons: Data is repeated (particularly dimensional data) across lots of files.

- Keeping data updated across many files can take time.

- Data of different granularities can get tricky.

- Immutable, append-only data means everything acts like a slowly changing dimension.

Recap, Suggestions, Q&A